





CSE 330: Operating Systems

Adil Ahmad

Lecture #2: Privileges, design, and general concepts

A modern OS wears many hats!



What makes the OS so special that it can do its tasks?

First, let's take a closer look at CPUs

- Let's think of CPUs in *very naïve* terms as having the following:
 - **Instructions** with hardcoded operations (e.g., add)
 - **Registers** where operations (e.g., add) are performed
- Code: $x = 10 + 20$. CPU executes the following “instructions”:
 - LOAD 10 from memory to a register (rax)
 - LOAD 20 from memory to a register (rbx)
 - ADD rbx to rax
 - STORE rax to memory
- LOAD, ADD, and STORE are examples of **memory and arithmetic instructions** implemented by your CPU

Some other instructions exposed by modern CPUs

- **Memory and arithmetic instructions**

- LOAD (read data from memory into register)
- ADD (add values in two registers)

- **Device access instructions**

- IN → read data from device (e.g., SSD)
- OUT → write data to device

- **CPU shutdown instructions**

- HLT → stop the CPU from spinning
- ... (many more)



Which of these instructions look “sensitive” to you?

- **Memory and arithmetic instructions**

- LOAD (read data from memory into register)
- ADD (add values in two registers)

- **Device access instructions**

- IN → read data from device (e.g., SSD)
- OUT → write data to device

Program can overwrite important files in SSD

- **CPU shutdown instructions**

- HLT → stop the CPU from spinning

Program can shut down your computer on a “whim” or mistake

- ... (many more)

CPU privilege modes **restrict** sensitive instructions

User mode
(U-mode)

Supervisor mode
(S-mode)

Only **non-sensitive (unprivileged)** instructions allowed:

- Memory access (e.g., load, store)
- Arithmetic (e.g., add, subtract)
- ...

Both **non-sensitive** and **sensitive (privileged)** instructions allowed:

- Memory access (e.g., load, store)
- Arithmetic (e.g., add, subtract)
- Device access (e.g., in, out)
- Shutdown (hlt),



What makes the OS so special that it can do its tasks?

Only OS executes at the S-mode and controls “privileged” functionality


The effective impact of user-mode on programs

- Cannot directly read/write to storage
- Cannot access devices like hardware timers, USB, GPU, etc.
- Cannot arbitrarily change the memory regions allocated to program
- Cannot read/write other process' memory regions
- Cannot read/write kernel regions, etc...


Does not impact the program itself → **limits changes a program can make outside of its own scope**

Privilege modes can be even *more* expressive!

Each **higher privilege mode** allows **more control** over system functions




A software should only have as much privilege as it needs.
(The principle of least privilege)




First topic: designing an OS

The three “hats” of the OS require lots of code



How should we combine the different code together?

Let's start with one large software: a **monolithic kernel**




The OS is one “big” program that runs within the same (S-mode) privilege

Easier for kernel *sub-systems* to collaborate
– no irritating boundaries

All kernel code runs within S-mode –
no communication slowdown

Why not to keep a monolithic kernel?

A big kernel means **complexity**, which reduces innovation and changes




More than 31 million
code lines today!

Changes must not break
prior functionality → hard



Why not to keep a monolithic kernel?

No internal enforcement means **minor bugs** → crashes or full compromise




Why not to keep a monolithic kernel?

Over-general and emphasis on reusing abstractions, which **can be inefficient**



Let's look at the alternative: a *micro-kernel*



Move most of the OS' functionality to user-space servers – *Small kernel*

Independent servers can use well-defined communication APIs – *Easy to update servers*

Minor bugs in *user-space servers* cannot compromise full system – *Better security*

Cardinal rule of designing a micro-kernel


“A concept is **tolerated inside the microkernel** only if **moving it outside** the kernel, i.e., permitting competing implementations, **would prevent the implementation** of the system's required functionality.”

(Jochen Liedtke, one of the inventors of the L4 micro-kernel)

Why not to keep a micro-kernel?

IPC makes **communication between components slow**

- One of the biggest challenges older micro-kernels faced (e.g., pre-L4)



Protected procedural call (PPC)

PPC entails:

- (a) Switching privilege
- (b) Copying args/results
- (c) Blocking

These aspects make communication slow

Why not to keep a micro-kernel?

- **Compatibility is also a problem**, since micro-kernels do not need to follow the widely-used APIs (communication protocols)
 - E.g., POSIX standards used by monolithic kernels
- Each version of a “user space” server can have its own API which can be evolved somewhat arbitrarily
- Programs must follow that API for execution

OS design (and life) is all about **trade-offs!**

- Monolithic has better performance and compatibility
- Micro-kernel has better modularity and security

Consumer OSs adopt a monolithic approach *with elements of micro-kernels*

- The kernel is kept as *small as possible*
- But most servers (e.g., device drivers) are in the privileged kernel space

Small security-critical systems (e.g., secure co-processors) use micro-kernels



My favorite design: a *nested kernel*


Small monitor (*the nested kernel*) that interposes HW and monolithic kernel

Benefits


- Still a simple-ish design
- Nested kernel can enforce security policies on buggy monolithic kernels

Drawbacks

- Hard to enable isolation within same privilege level (*software tricks?*)
- Slower than a monolithic kernel



Nested Kernel (Dautenhahn et. Al, ASPLOS '15)



General OS concepts to get us started

What is the difference between the “OS” and “kernel”?

- Kernel is “core” part of the operating system that interacts with the hardware directly (e.g., for memory management, scheduling, etc.)
- The ”OS” is essentially a collection of the ”kernel” and pre-packaged user-level software. These software basically provide:
 - Graphical user input (GUI)
 - Window management
 - File explorers, etc.
- In this class, whenever I say OS, I really mean the “kernel” :p


Integral concept of a **process** in OS

- *Programs* are a set of codes (static entity) or executable files
 - e.g., compile a “Hello World” C program → a.out
- *Processes* are the *running version of programs*
 - We can run the same program in multiple processes
 - One host can run multiple processes
- Let’s check a C++ analogy
 - Definition of a class → “program”
 - Object instance of a class → “process”

What events result in process creation?

- Principle events that initiate the creation of processes
 - System initialization (e.g., *init* process)
 - Execution of a process creation system call by a running process
 - User request to create a process

An illustration of typical process creation



Each process has an (isolated) **address space**

- A process requires hardware resources to execute, including CPU and main memory (also called DRAM)
 - DRAM → Dynamic Random Access Memory
- OS does not show all system memory to a process, rather abstracts a portion of memory to each process
- View of memory shown to each process is called its “address space”
 - Typically, each process’ memory view is different from others

(More details about address spaces in lecture #8 onwards.)

Process can spawn other instances of itself

- Ask OS to create a completely different clone of the same program.
 - Clone will be a new “process” with isolated memory
 - E.g., running two instances of a web browser
- Ask OS to execute it *simultaneously* on multiple CPUs
 - Each context is what we call a **thread**
 - Threads share the same memory but can execute different functions independently of each other
 - Useful for information sharing
 - E.g., two tabs within a web browser (*not always true but for illustrative purpose!*)
- We will learn all about “multi-processing” and “multi-threading”!

Recall that processes (and their threads) are *unprivileged*


Will the process ever require access to “privileged” functionality?

- Yes!
 - Reading a file from the SSD
 - Writing to a file in USB storage
 - Sending data through the network
 - Reading a hardware provided timer, etc.
- We **need a mechanism to ask the operating system kernel** to help us with privileged tasks
 - OS kernels can check if the process is “authorized” or not

Process interacts with the kernel using system calls

Requests made to the OS from a user process for a certain operation (e.g., file system access, device access, etc.)

Made using special instructions in modern CPUs (e.g., SYSCALL in x86)



System calls require:

- (a) Switching privilege/address spaces
- (b) Copying arguments/results

(More details about transitions in lecture #9 onwards.)

Some examples of famous system calls in Linux

`fd = open(file, permissions, ...)`

- *Open a file for reading, writing, or both*

`s = close(file)`

- *Close an open file*

`n = read(fd, buf, nbytes)`

- *Read data from a file into a buffer*

`n = write(fd, buf, nbytes)`


- *Write data from a buffer into a file*

`pos = lseek(fd, offset, whence)`

- *Move the file pointer*

`s = stat(name, &buf)`

- *Get a file's status info (e.g., size, etc.)*



That's all, folks. See you in the next class!