




CSE 330: Operating Systems

Adil Ahmad

Lecture #14: Device communication and protocols

Recall the OS does the **plumbing** in modern computers

Provides an abstraction for the messy hardware



Manipulating hardware requires: (a) programming knowledge and (b) understanding of the hardware

Plumbing analogy:

With the OS' abstraction, users do not care about pipes, rather they can simply turn the taps on or off.

What kind of hardware plumbing have we discussed thus far?


Provided below are some examples:

- **Memory management:** OS abstracted away aspects of
 - Memory allocation and deallocation from a user process
- **CPU management:** OS abstracted away aspects of
 - Translation Lookaside Buffer (TLB) management
 - CPU cache management

External hardware device management

- This is another core operating system **plumbing** duty
- Provide a clean, easy-to-use way for external “connected” devices and processes in your computer to communicate
- Device management is a highly-complicated task for several reasons:
 - Devices are very complex and diverse in their functions
 - Devices can be connected to your computer “dynamically”
 - Many processes want to access devices at the same time
 - Many devices want to communicate with your computer at the same time

Remember when I showed you this graph of Linux code




- A large chunk of the new lines (**and bugs!**) come from the ever-increasing list of devices that Linux is trying to support! ☺


Let's dive into *Input/Output (I/O)* devices

- **Definition:** a set of external connected interfaces that send and receive data (or information) to your computer
- **What are example of I/O devices connected to your computer?**
 - Hard disk drives (HDDs)
 - Solid state drives (SSDs)
 - Keyboard
 - Mouse
 - Graphics Processing Units (GPUs)
 - USB drives, etc.


Let's see how devices are connected to a computer



Bus interconnect allows two-way communication




Let's see what a typical device looks like to the OS



- **Status:** Tells the OS whether the device is ready, busy, initialized, etc.
- **Command:** OS will write what it needs here, and if device is ready, it will perform the corresponding operation
- **Data:** OS will write any information the device needs to complete its task at this location


How does the OS know how to talk to device registers?



A2	A1	A0	READ MODE	WRITE MODE
0	0	0	• Receive Holding Register	• Transmit Holding Register
0	0	0	N/A	• LSB of Divisor Latch when Enabled
0	0	1	N/A	• Interrupt Enable Register
0	0	1	N/A	• MSB of Divisor Latch when Enabled
0	1	0	• Interrupt Status Register	• FIFO control Register
0	1	1	N/A	• Line Control Register
1	0	0	N/A	• Modem Control Register
1	0	1	• Line Status Register	N/A
1	1	0	• Modem Status Register	N/A
1	1	1	• Scratchpad Register Read	• Scratchpad Register Write


- Hardware manuals provided by hardware vendors explain the “*language*”!
- If you want to do OS research, get ready to read some badly written manuals! 😊


The other “hidden” half of a typical device



- Device needs a **tiny CPU** (micro-controller) to perform computations
- Device (like CPU) will need some **tiny DRAM** for computation
- Some example of other chips:
 - Capacitor for power-backup
 - Hardware security chips, etc.

An example of the hard drive disk PCB





Understanding I/O communication

Acknowledgement

- Most of the slide content in this section are taken directly from Prof. Zhicao Cao and Prof. Ming Zhao's CSE 330 classes
- These slides are fantastic!!
- Hence, I will present them mostly as is (with their permission, of course!)

Basic I/O protocol from the OS perspective


```
while (STATUS == BUSY)
    ; // spin
```

Write data to DATA register

Write command to COMMAND register


```
while (STATUS == BUSY)
    ; // spin
```

Examining the workings of the basic I/O protocol




```
while (STATUS == BUSY)          //1
    ;
    ; // spin
Write data to DATA register    //2
Write command to COMMAND register //3
while (STATUS == BUSY)          //4
    ;
    ; // spin
```

Examining the workings of the basic I/O protocol




```
while (STATUS == BUSY)          //1
    ;
    // spin
Write data to DATA register    //2
Write command to COMMAND register //3
while (STATUS == BUSY)          //4
    ;
    // spin
```

Examining the workings of the basic I/O protocol




```
while (STATUS == BUSY) //1
    ; // spin
Write data to DATA register //2
Write command to COMMAND register //3
while (STATUS == BUSY) //4
    ; // spin
```

Examining the workings of the basic I/O protocol




```
while (STATUS == BUSY)          //1  
    ; // spin  
Write data to DATA register    //2  
Write command to COMMAND register //3  
while (STATUS == BUSY)          //4  
    ; // spin
```

Examining the workings of the basic I/O protocol




```
while (STATUS == BUSY) //1  
    ; // spin  
Write data to DATA register //2  
Write command to COMMAND register //3  
while (STATUS == BUSY) //4  
    ; // spin
```

Is there some inefficiency (resource waste) problem in this example?




```
while (STATUS == BUSY)          //1
    ; // spin
Write data to DATA register    //2
Write command to COMMAND register //3
while (STATUS == BUSY)          //4
    ; // spin
```

All this waiting (called *polling*) is not very efficient




```
while (STATUS == BUSY) //1  
    ; // spin  
Write data to DATA register //2  
Write command to COMMAND register //3  
while (STATUS == BUSY) //4  
    ; // spin
```

How would you solve this resource waste?




```
while (STATUS == BUSY)          //1  
    ; // spin  
Write data to DATA register    //2  
Write command to COMMAND register //3  
while (STATUS == BUSY)          //4  
    ; // spin
```

Solving the inefficiency problem using interrupts



```
while (STATUS == BUSY) //1  
    wait for interrupt;  
Write data to DATA register //2  
Write command to COMMAND register //3  
while (STATUS == BUSY) //4  
    wait for interrupt;
```

Service other processes while the device is handling data




```
while (STATUS == BUSY) //1  
    wait for interrupt;  
Write data to DATA register //2  
Write command to COMMAND register //3  
while (STATUS == BUSY) //4  
    wait for interrupt;
```

Is there any disadvantage to using interrupts vs polling?


- Can lead to **more frequent context switches** (e.g., context switch to a different process and back)
- In practice today, OS uses a hybrid approach:
 - short device computations → polling and wait
 - long device computations → interrupts
- Another hardware optimization: interrupt batching
 - CPU queues a sequence of interrupts, and sends them all at once

Is there something else that could be expensive in this protocol?



```
while (STATUS == BUSY) //1
    wait for interrupt;
    Write data to DATA register //2
    Write command to COMMAND register //3
    while (STATUS == BUSY) //4
        wait for interrupt;
```

Is there something else that could be expensive in this protocol?



```
while (STATUS == BUSY) //1  
    wait for interrupt;  
    Write data to DATA register //2
```

- A basic, menial operation that **peasants** can do (*copy data from one place to another*)
- Wasting your **expensive, precious, royal** CPU for this task is not ideal


```
    wait for interrupt;
```

Programmed I/O versus Direct Memory Access

- **Programmed I/O (PIO)**
 - This is the version we've discussed so far, and it is also called memory-mapped I/O (MMIO)
 - CPU directly tells device what data is needed
 - CPU is involved in the data transfer (e.g., writing to DATA register)
- **Direct Memory Access (DMA)**
 - CPU leaves (prepares) data in main memory
 - DMA hardware performs the data copy
 - CPU is free to complete other tasks in the background


Understanding DMA: memory access and PIO

- CPUs access memory by sending requests to a system bus
- Consider PIO: `writeReg(CMD, 0x100)`
 - Copy data at 0x100 to CMD register
 - System bus forwards data to device
- Depending on size to send, the CPU must:
 - Prepare data in DRAM
 - Write to MMIO address to copy data (e.g., byte-by-byte)



Understanding DMA: using controllers for speed-up


- **DMA controllers:**
 - Specialized hardware units that can only be programmed to read/write memory
- OS can then do the following:
 - Programs DMA controller to send data to the device in the background
 - Set up interrupts for when DMA transfers are completed
- OS can resume execution of other tasks



Is this approach still (somewhat) inefficient?

DMA version 2.0: removing the controller

- Allow device to directly use the system bus (also called *bus mastering*)
- OS can then do the following:
 - Prepares memory region (e.g., 0x100) with a batch of data
 - Tell device (e.g., using PIO) where the data is and let device handle the rest
 - Interrupts enabled like before
- No need for synch. between the controller and device; hence, faster!



Quick comparison: *controllers* and *bus mastering*


Controller approach (also called *third-party DMA*)

- Leaves complex handling of system bus requests to the system (e.g., motherboard/chipset)
- Suitable for simpler devices
- Slower since it requires interactions b/w CPU-controller-device

Bus mastering approach (also called *first-party DMA*)


- Requires complex tasks be done by the device directly
- Suitable for more advanced devices (e.g., GPUs)
- Faster since it avoids having the middle-man

Recap of code-level illustration of Programmed I/O (before)



```
while (STATUS == BUSY) //1  
    wait for interrupt;  
    Write data to DATA register //2  
    Write command to COMMAND register //3  
    while (STATUS == BUSY) //4  
        wait for interrupt;
```

Code-level illustration for Direct Memory Access (after)



```
while (STATUS == BUSY) //1
    wait for interrupt;
Initiate DMA transfer //2a
Wait for interrupt //2b
Write command to COMMAND register //3
while (STATUS == BUSY) //4
    wait for interrupt;
```

Is there a security threat with bus mastering DMA?


- Sadly, yes! Can directly access physical memory region
- **Malicious devices can send DMA requests for any regions**
 - e.g., belonging to the OS or other programs etc.
 - Google the **BadUSB** attack! 😊
- **OSs can also be malicious and leverage DMA to steal data**
 - **Any example scenario?**
 - Consider the OS running inside your VirtualBox setups
 - If it is malicious, it can try to use a device (e.g., keyboard) to access data from your host OS using DMA

Is there a similarity to this problem with processes?

- **Yes!** If processes directly access physical memory, they could also access memory regions belonging to other processes (lecture 4 – 7)
- **How did we solve that problem back then?**
- We leveraged a translation layer (aka *page tables*) to define and restrict regions of memory accessible to processes

Adding *another* translation layer for DMA protection

- Enabled by a hardware unit inside the CPU called the **I/O memory management unit (IOMMU)**
- The OS sets *translation tables* for DMA and all access to the DRAM from devices occurs through I/O virtual address
- Like normal page tables, the OS can set permissions on mapped regions (e.g., **read-only**, **read-write**, etc.).



Questions? Otherwise, see you next class!