

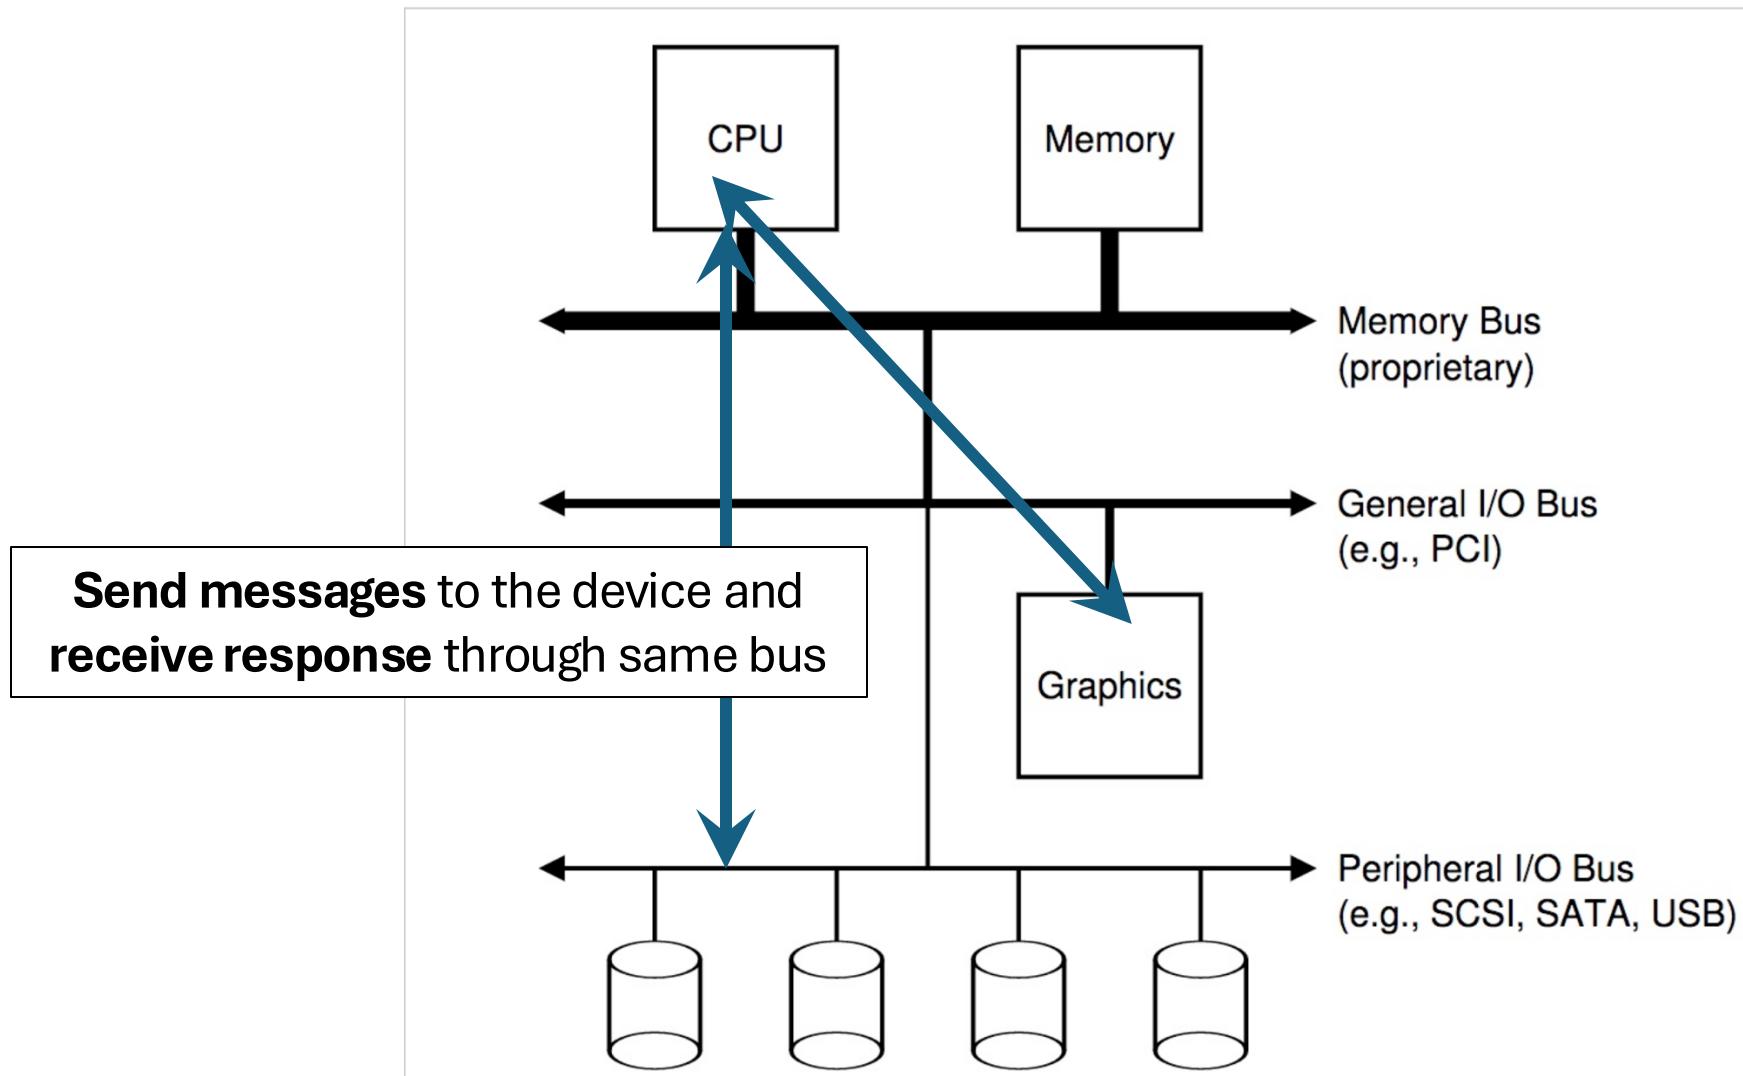


# CSE 330: Operating Systems

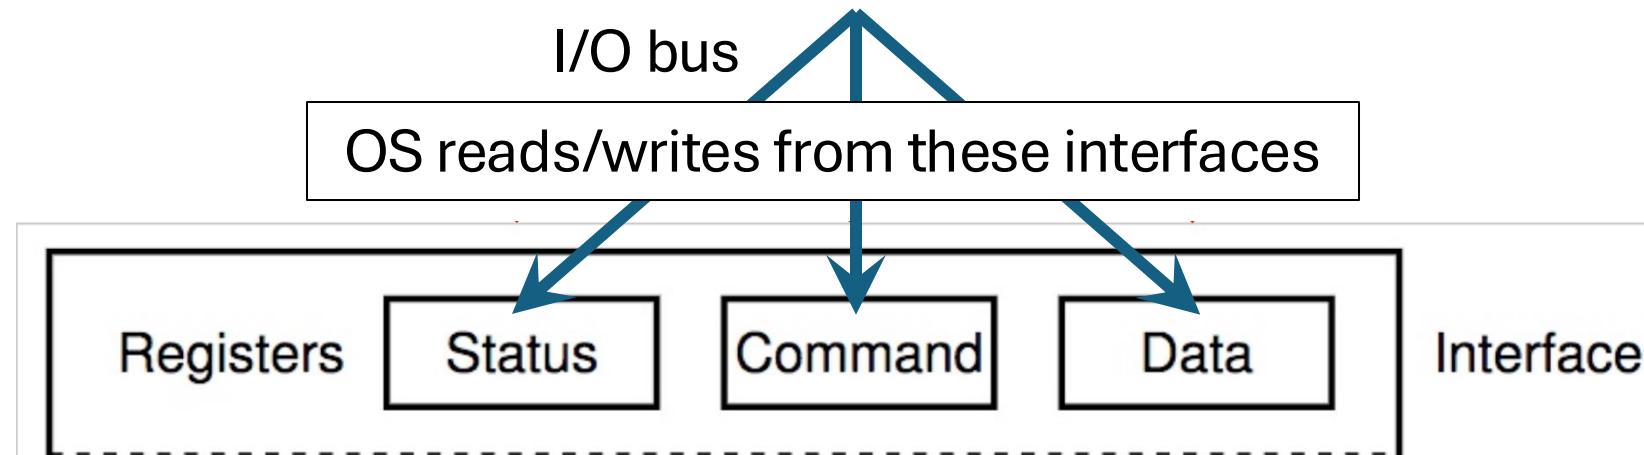
Adil Ahmad

**Lecture #15:** Device examples, HDDs and SSDs

# Bus interconnect allows two-way communication



# Let's see what a typical device looks like to the OS



- **Status:** Tells the OS whether the device is ready, busy, initialized, etc.
- **Command:** OS will write what it needs here, and if device is ready, it will perform the corresponding operation
- **Data:** OS will write any information the device needs to complete its task at this location

# Basic I/O protocol from the OS perspective

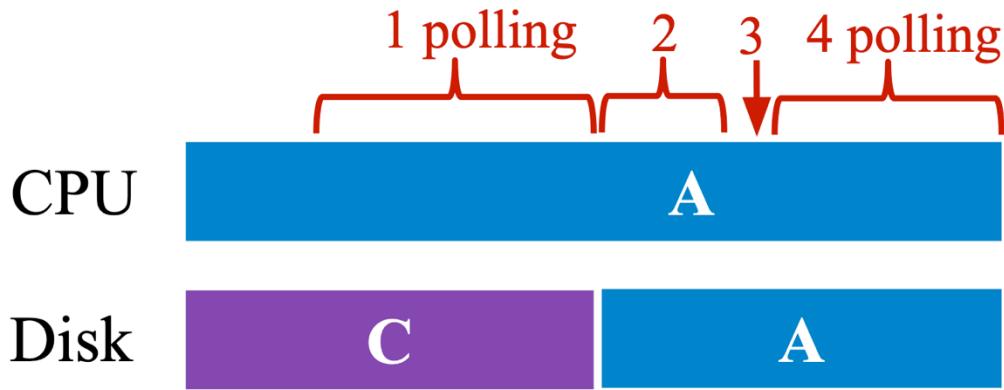
```
while (STATUS == BUSY)
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

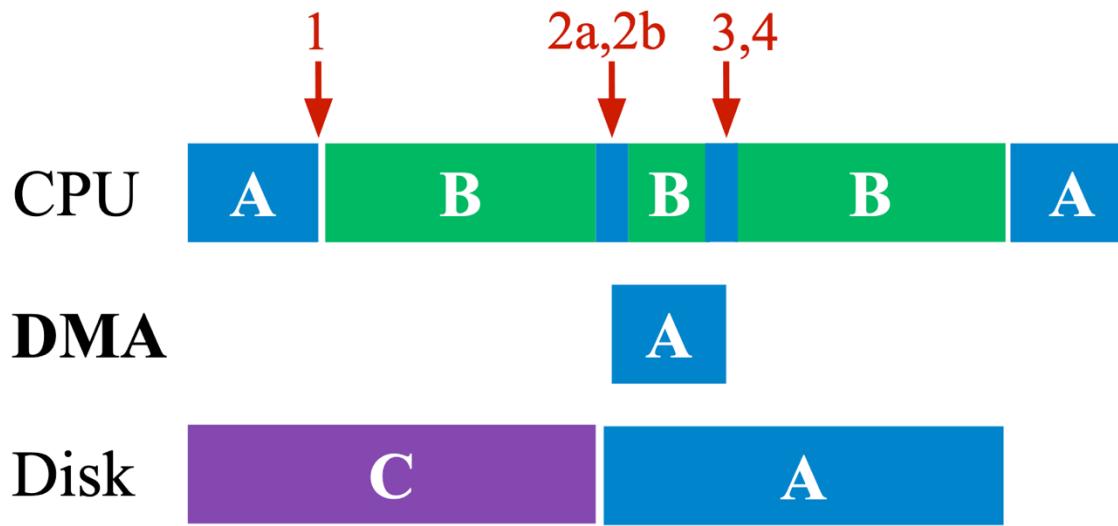
```
while (STATUS == BUSY)
    ; // spin
```

# Solving the inefficiency problem using interrupts



```
while (STATUS == BUSY) //1  
    wait for interrupt;  
Write data to DATA register //2  
Write command to COMMAND register //3  
while (STATUS == BUSY) //4  
    wait for interrupt;
```

# Code-level illustration for Direct Memory Access (after)



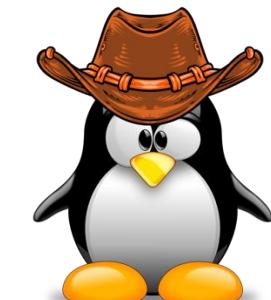
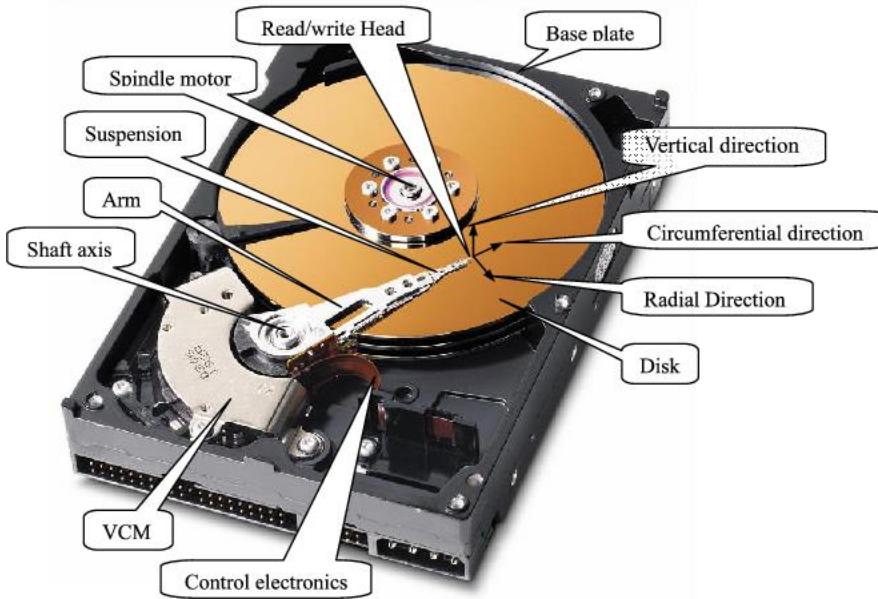
```
while (STATUS == BUSY) //1
    wait for interrupt;
Initiate DMA transfer //2a
Wait for interrupt //2b
Write command to COMMAND register //3
while (STATUS == BUSY) //4
    wait for interrupt;
```

# Is there a security threat with DMA?

- Sadly, yes! Can directly access physical memory region
- Malicious devices can send DMA requests for any regions
  - e.g., belonging to the OS, secret keys, etc.
  - Check out the BadUSB attack! ☺
- OSs can also be malicious and leverage DMA to steal data
  - **Any example scenario?**
  - Consider the OS running inside your VirtualBox setups
  - If it is malicious, it can try to use a device (e.g., keyboard) to access data from your host OS using DMA

# Adding translation tables for DMA protection

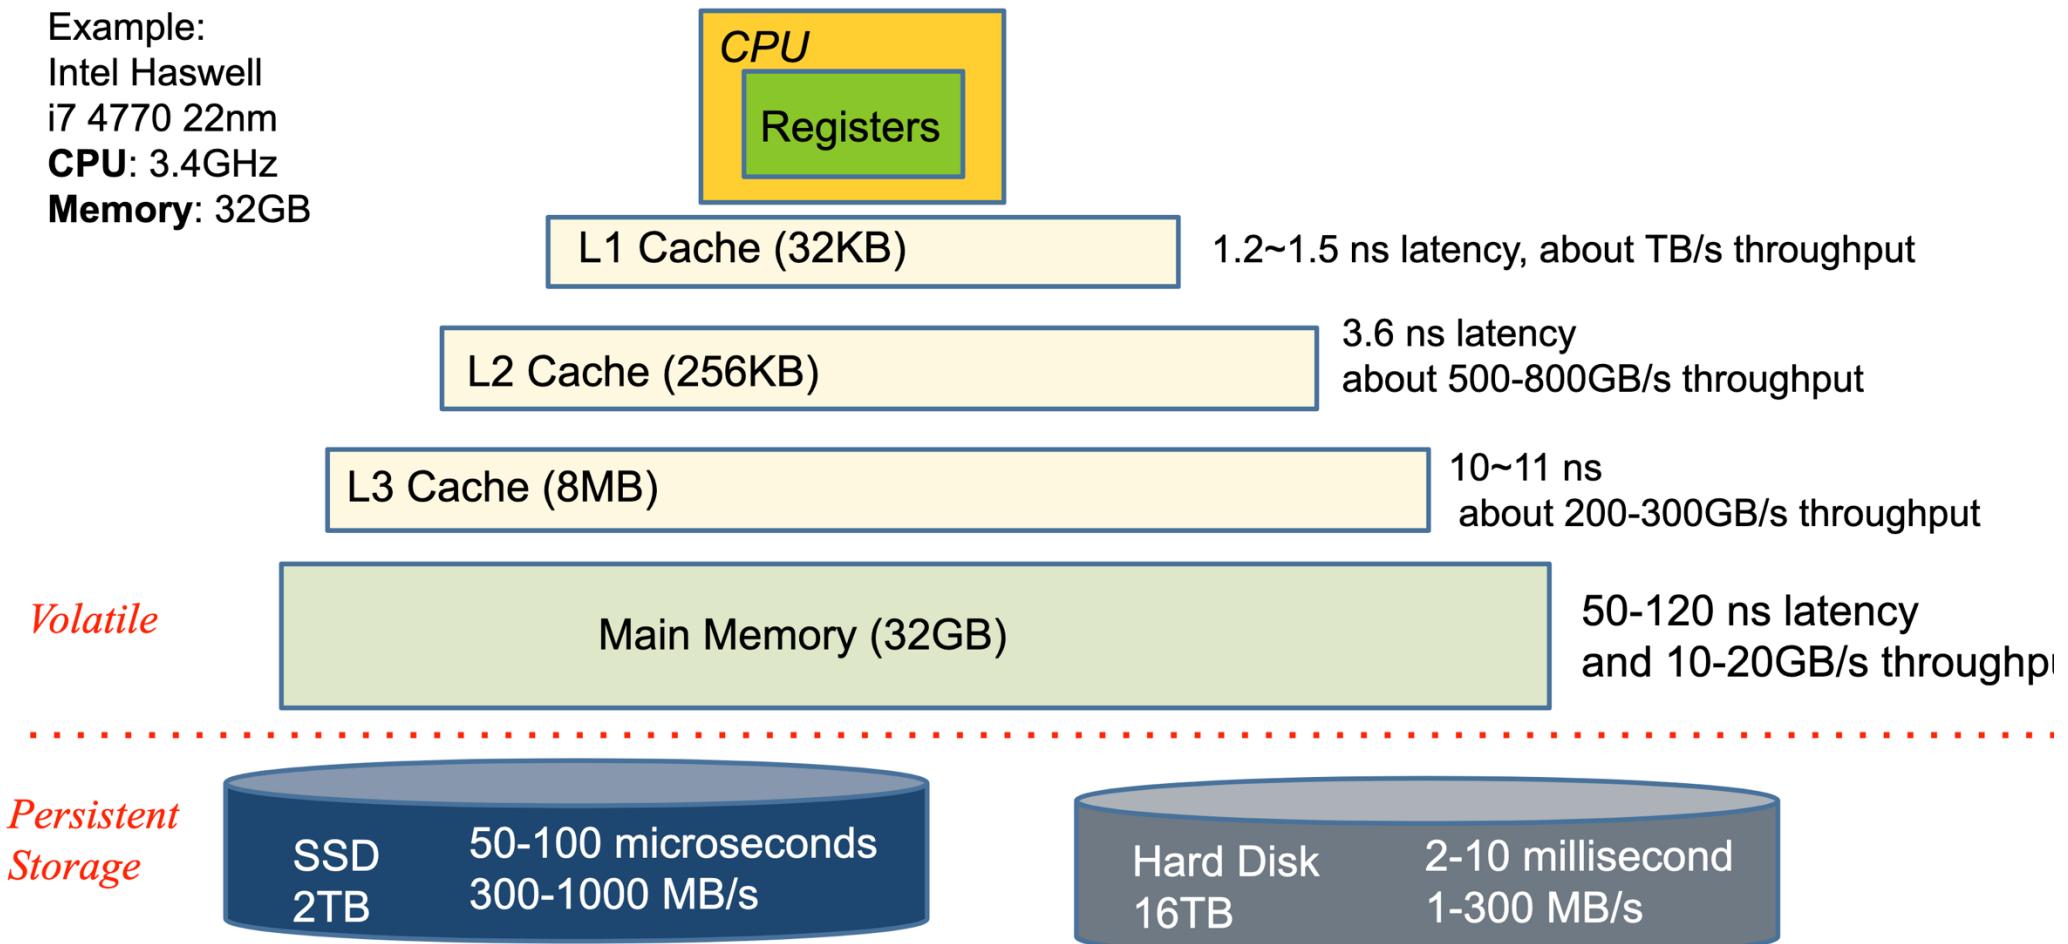
- Enabled by a hardware unit inside the CPU called the *I/O memory management unit (IOMMU)*
  - Like the memory management unit (MMU)
- The OS sets **I/O page tables** for DMA and all access to the DRAM from devices occurs through **I/O virtual address**
- Like normal page tables, the OS can define some regions as read-protected, write-protected, etc.



## Device example: Hard disk drives (HDDs)

# The memory-storage hierarchy of modern computers

Example:  
Intel Haswell  
i7 4770 22nm  
**CPU:** 3.4GHz  
**Memory:** 32GB



# Basic interface of an HDD

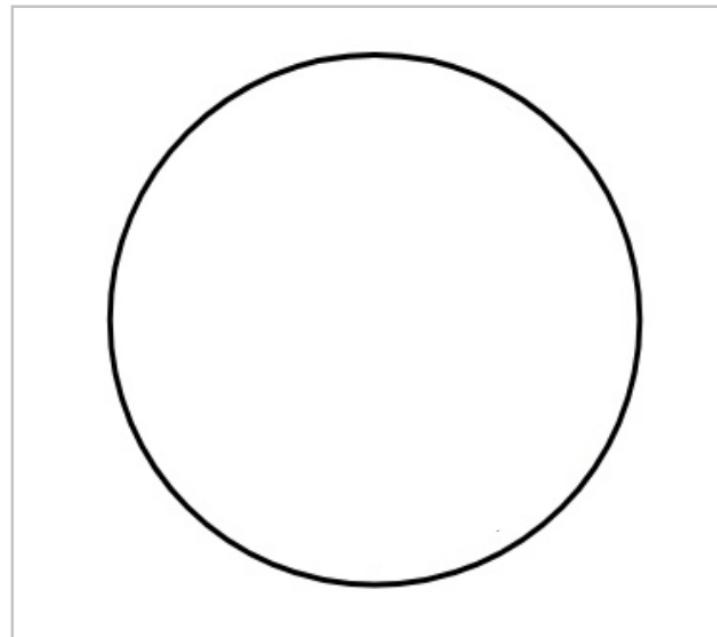
- A magnetic HDD has a **sector-addressable** space
  - You can think of a disk as an array of sectors
  - Each sector (also called “logical block”) is the smallest unit of transfer
- Sectors are typically 512 or 4096 bytes
  - 100 GB disk of 512-byte sectors = **209715200** total sectors
- HDDs support two general I/O commands:
  - Read from sectors (blocks)
  - Write to sectors (blocks)

# High-level description of an HDD's internal structure

- The 1-dimensional array of logical blocks is mapped into the sectors of the disk sequentially
  - Sector 0 is the first sector of the first track on the outermost cylinder
  - Mapping proceeds in order through that track, then the rest of the tracks in that cylinder, and then through the rest of the cylinders from outermost to innermost

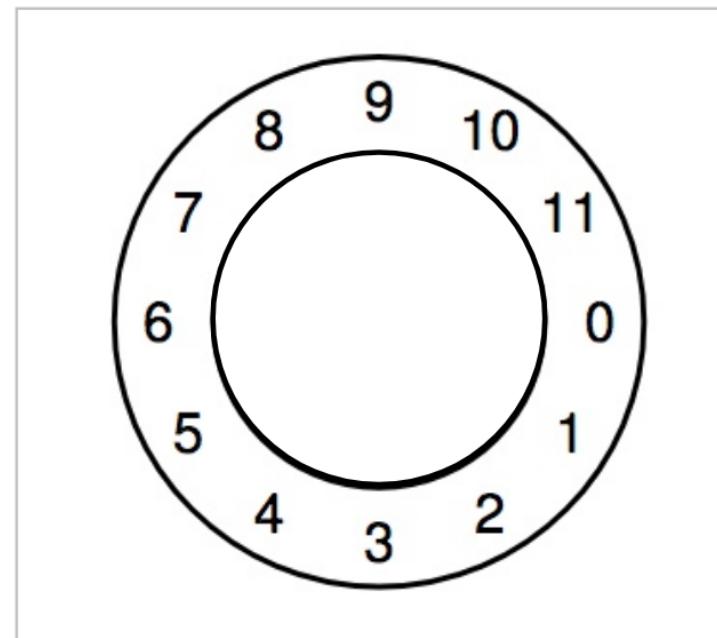
# Internal workings of a Hard Disk Drive (HDD)

Platter  
Covered with a magnetic  
film



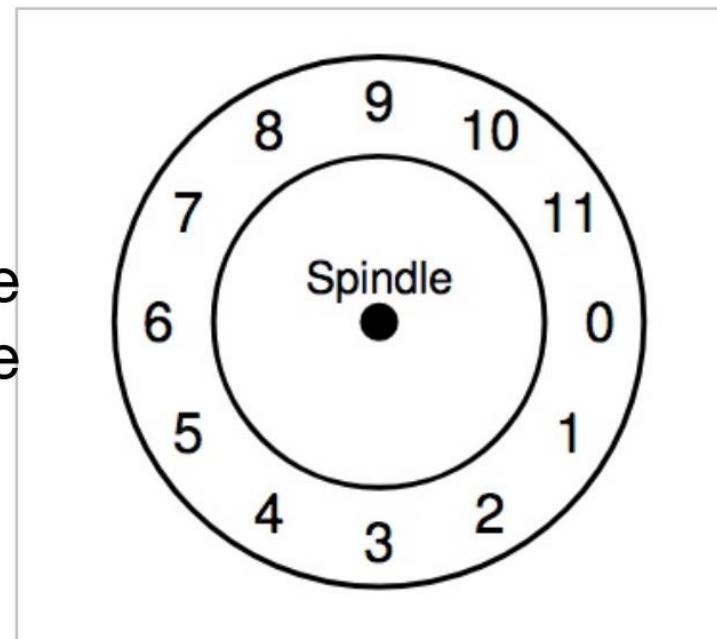
# Internal workings of a Hard Disk Drive (HDD)

A single track example



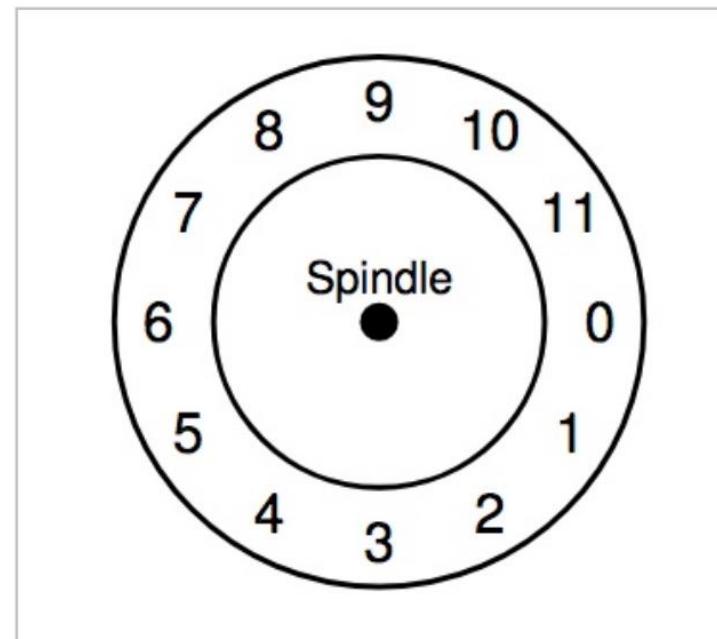
# Internal workings of a Hard Disk Drive (HDD)

Spindle in the center of the surface



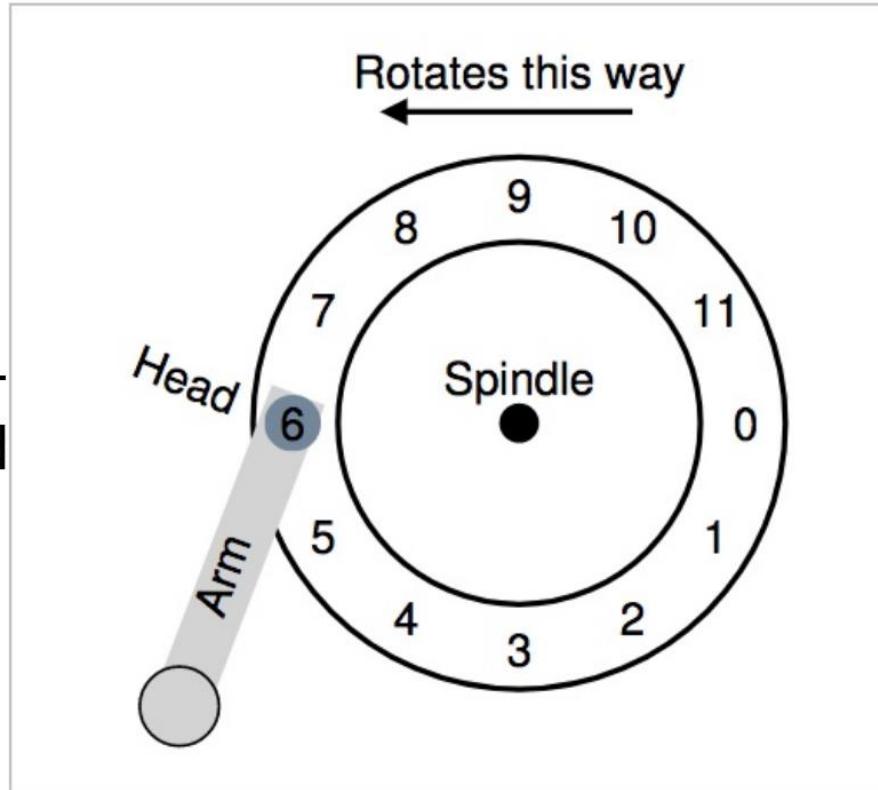
# Internal workings of a Hard Disk Drive (HDD)

The track is divided into numbered sectors

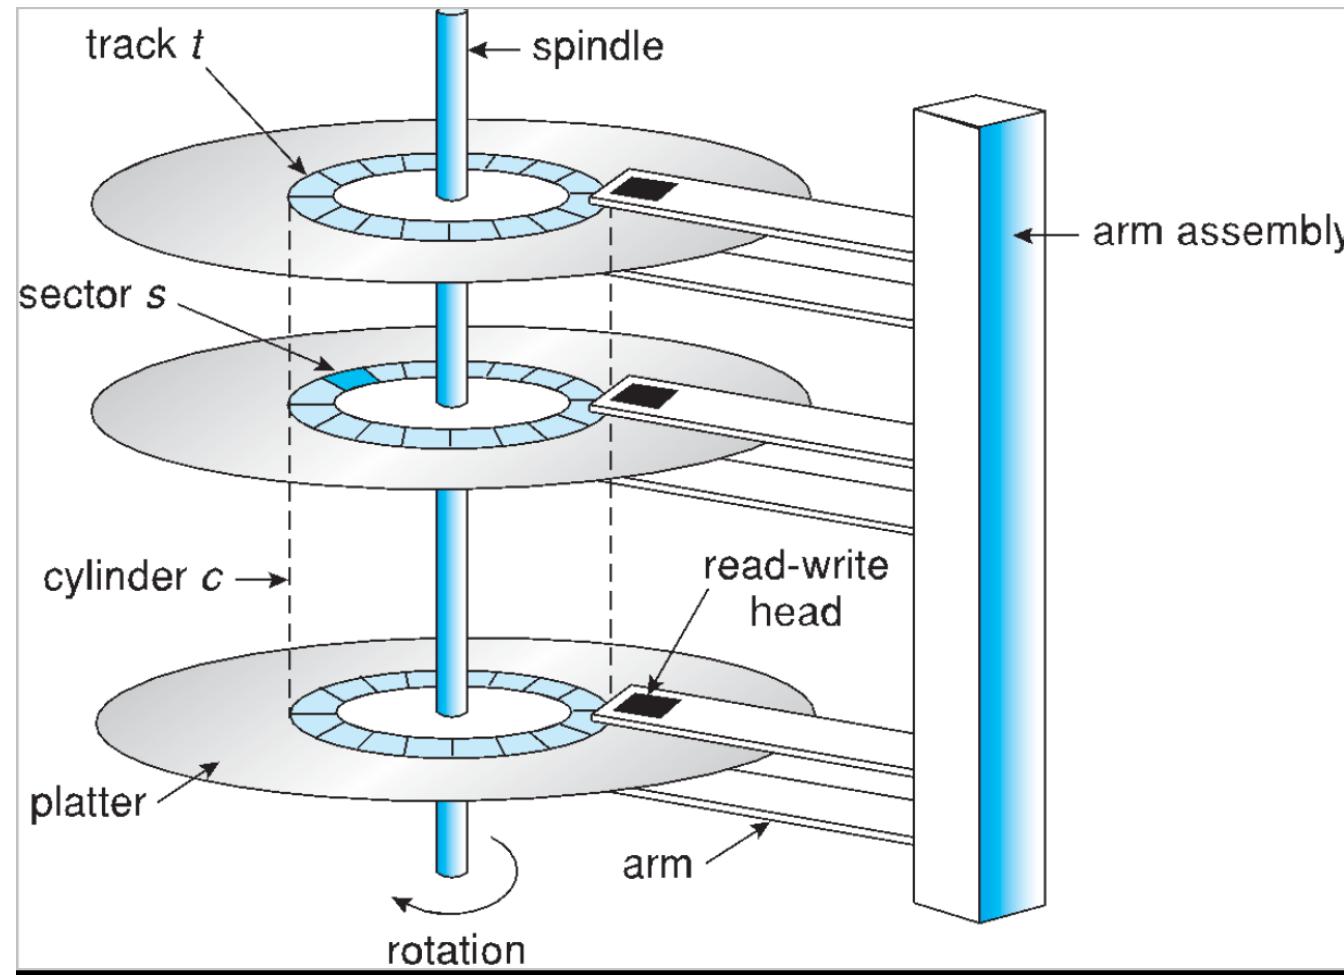


# Internal workings of a Hard Disk Drive (HDD)

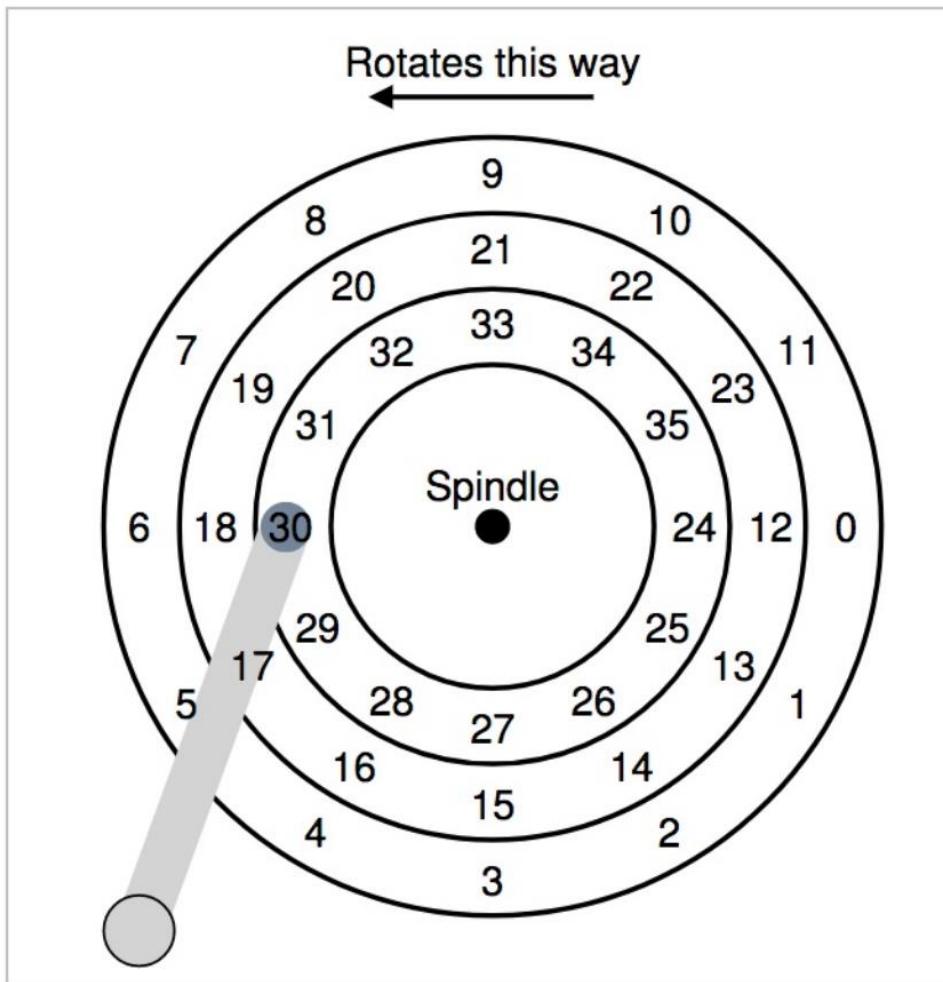
A single track + an arm + a head



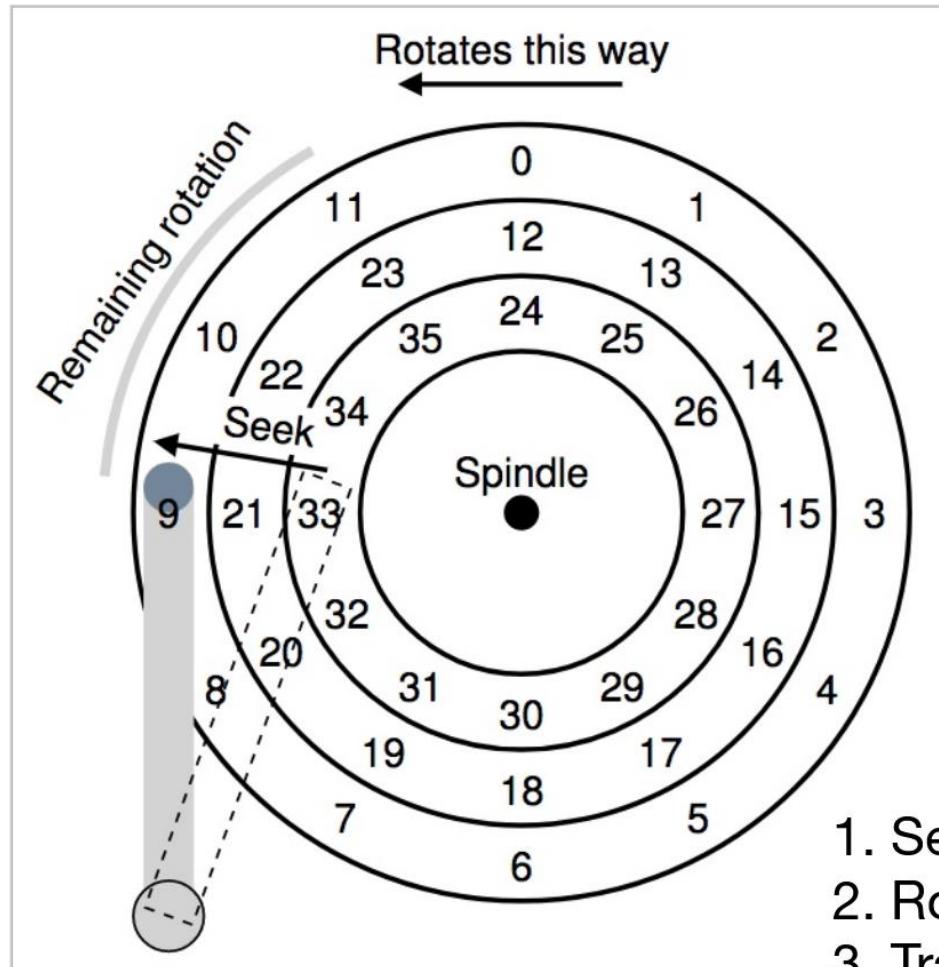
# Internal workings of a Hard Disk Drive (HDD) – 3D view



# Let's try to read sector 0. How would you do it?

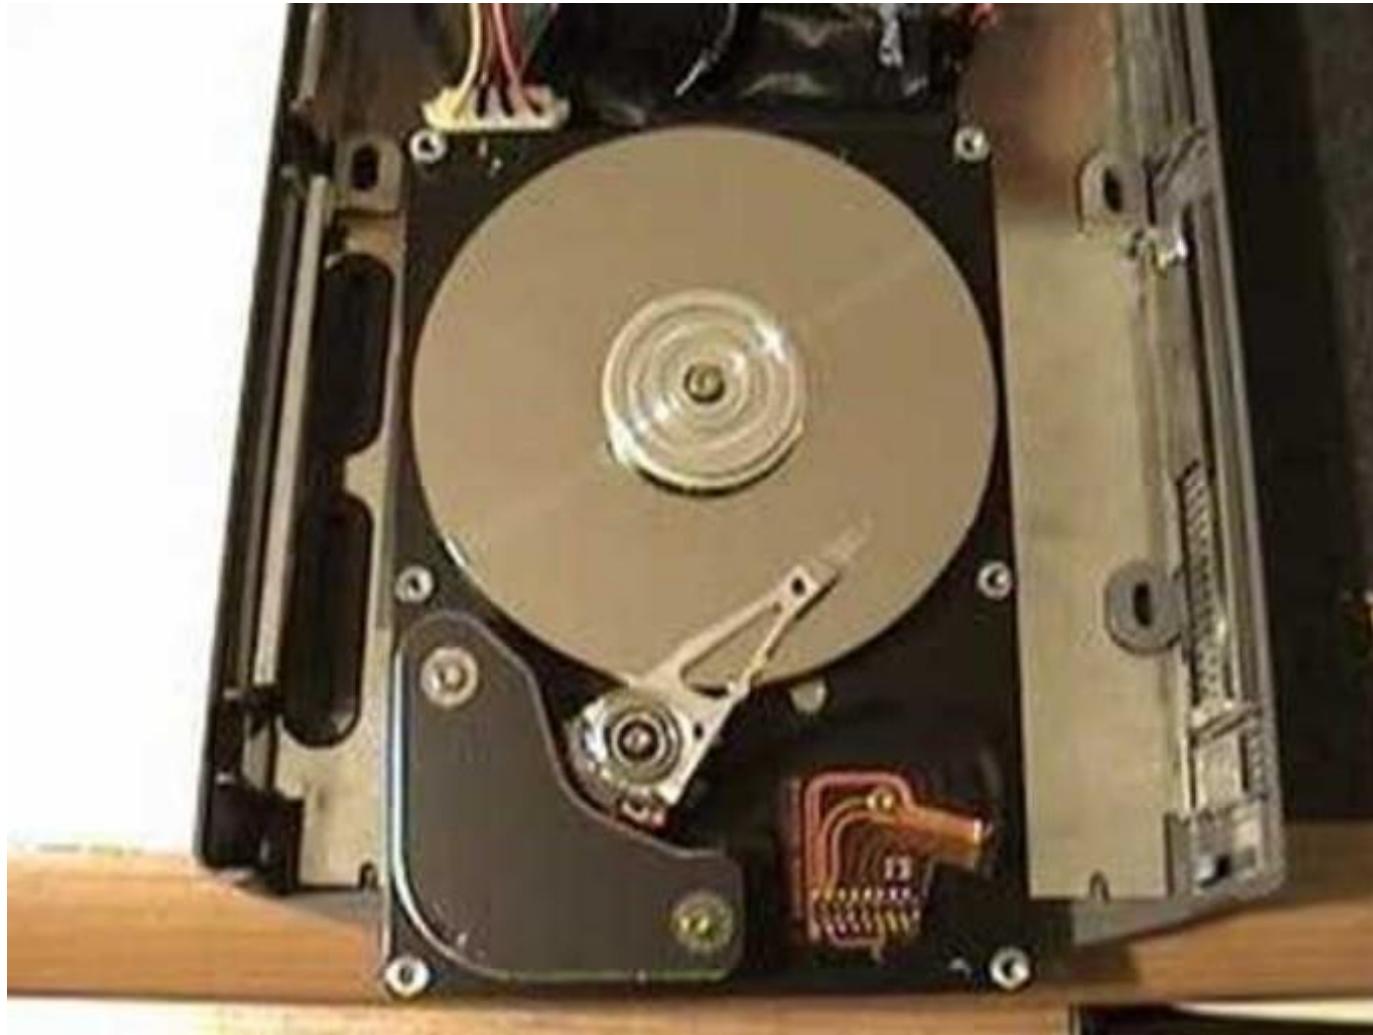


# Let's try to read sector 0. How would you do it?



1. Seek for right track
2. Rotate (sector 9 → 0)
3. Transfer data (sector 0)

Let's watch a video on the *internal HDD workings!*



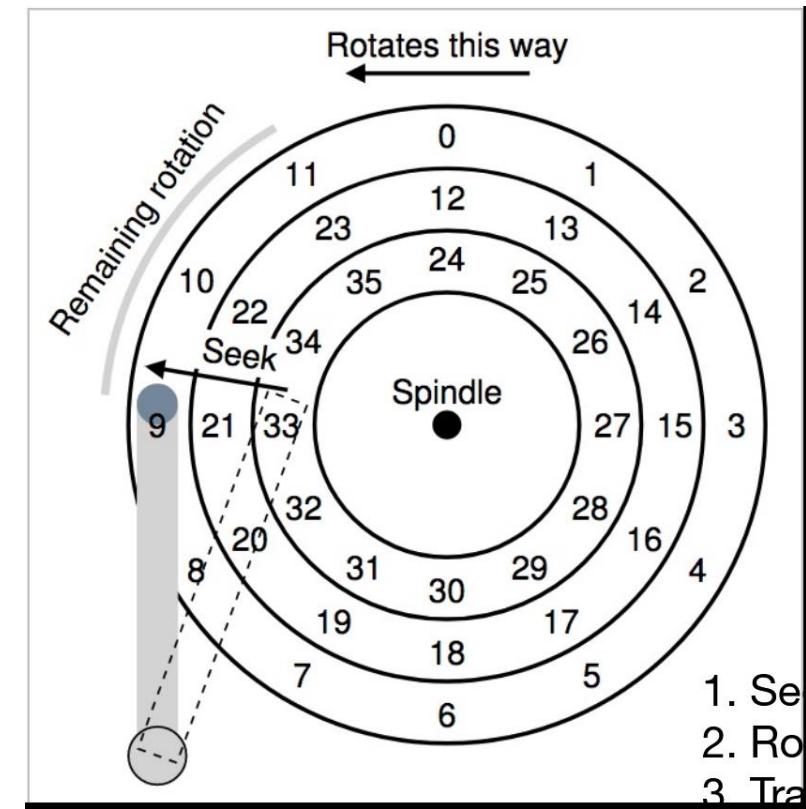
<https://www.youtube.com/watch?v=9eMWG3fwEU&t=30s>

# Understanding workings of HDD is important for the OS

- Disk **performance depends significantly** on the working mechanism
- The disk latency can be calculated by considering the three steps that the HDD must perform:
  - HDD I/O latency =  $\text{Latency}_{\text{seek}} + \text{Latency}_{\text{rotate}} + \text{Latency}_{\text{transfer}}$
  - Let's see how many milliseconds are taken by these (*seek, rotate, and transfer*) operations!

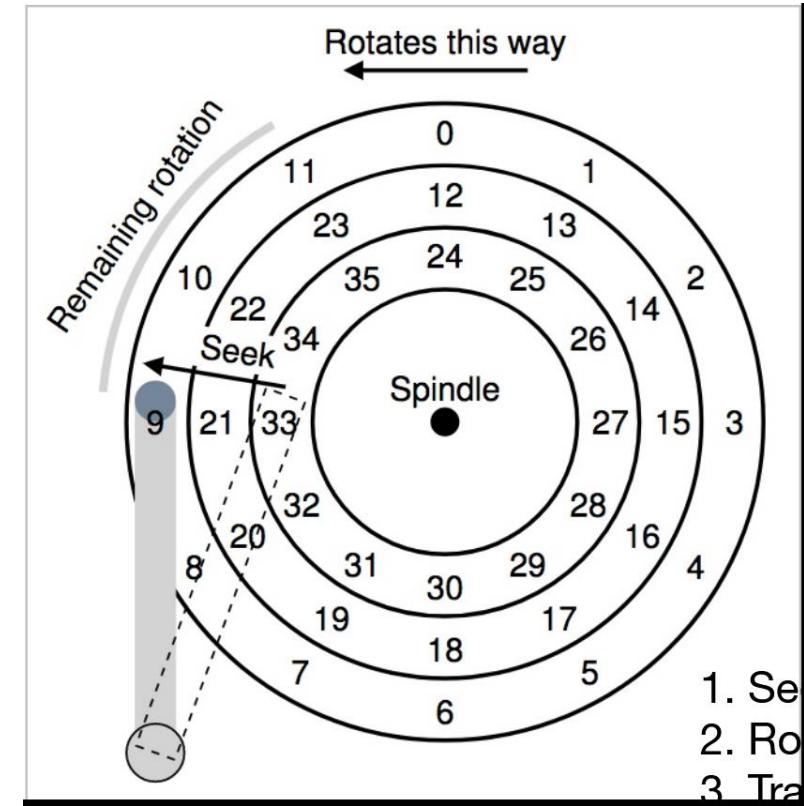
# Seek, Rotate, Transfer

- Seek may take several milliseconds (ms)
  - Depends on hardware
- Settling along can take 0.5 - 2ms
- Entire seek ( $L_{\text{seek}}$ ) takes 4 – 5 ms for a decent HDD, but it may take up to 10 ms for slow (past) HDDs



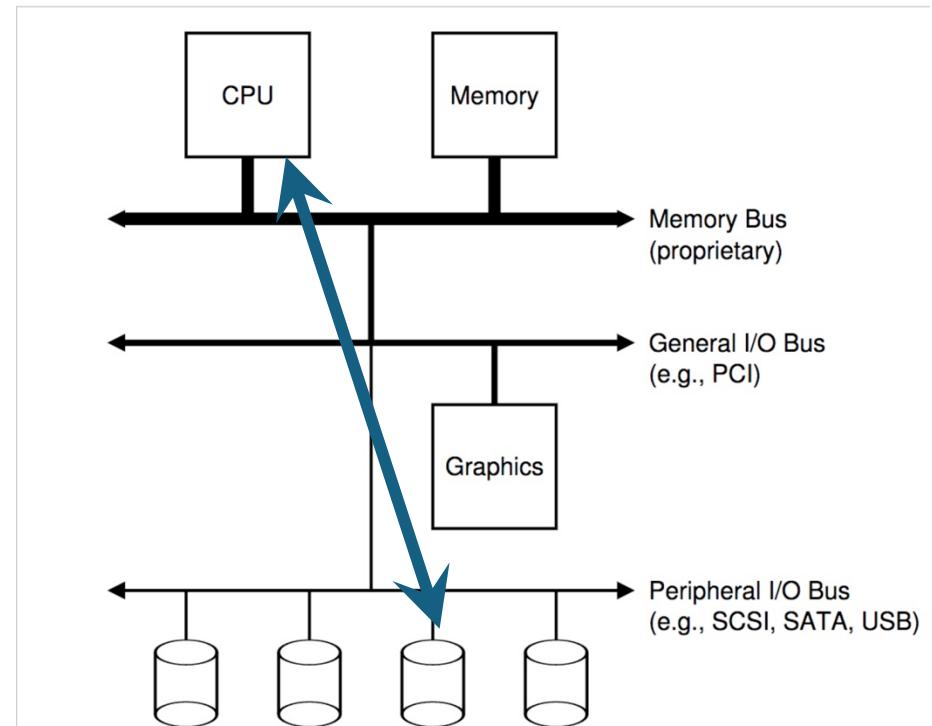
# Seek, **Rotate**, Transfer

- Depends on Rotation Per Minute (RPM) supported by the HDD
  - 7200 RPM is common nowadays
  - 15000 RPM is high end
  - Old computers may have 5400 RPM disks
- $L_{rotate}$  for 7200 RPM
  - $= 7200 \text{ rotations} / 60 \text{ s} = 120 \text{ rots / sec}$
  - $= 1 \text{ sec} / 120 \text{ rotations} = \mathbf{8.3 \text{ ms}} / \text{rotation}$
- It may take **4.2 ms on average** to rotate to target ( $0.5 * 8.3 \text{ ms}$ )



# Seek, Rotate, Transfer

- Relatively fast
  - Depends mostly on the link speed
  - SATA is the link for HDDs today
- 100+ MB/s is typical for SATA
  - Up to **600MB/s** for SATA III)
- Calculating  $L_{\text{transfer}}$ 
  - $1\text{s} / 100\text{MB} = 10\text{ms} / \text{MB} = 4.9\mu\text{s}$ /sector
    - Assuming 512-byte sector



# Examining workload suitability for HDDs

- Seek and rotations are slow while transfer is relatively fast
  - seek = ~4 – 10 ms, rotate = ~4 – 5 ms, transfer = ~5 us
- **What kind of workload is best suited for disks?**
  - **Sequential I/O:** access sectors in order (transfer dominated)
- **Random** workloads access sectors in a random order (seek+rotation dominated)
  - Typically, slow on disks
  - Never do random I/O unless you must!
  - For instance, **quicksort** is a terrible algorithm for disk

# Let's do some disk performance calculations together

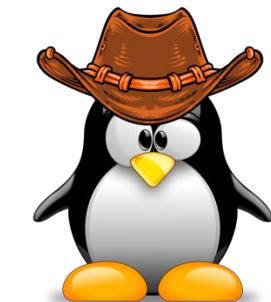
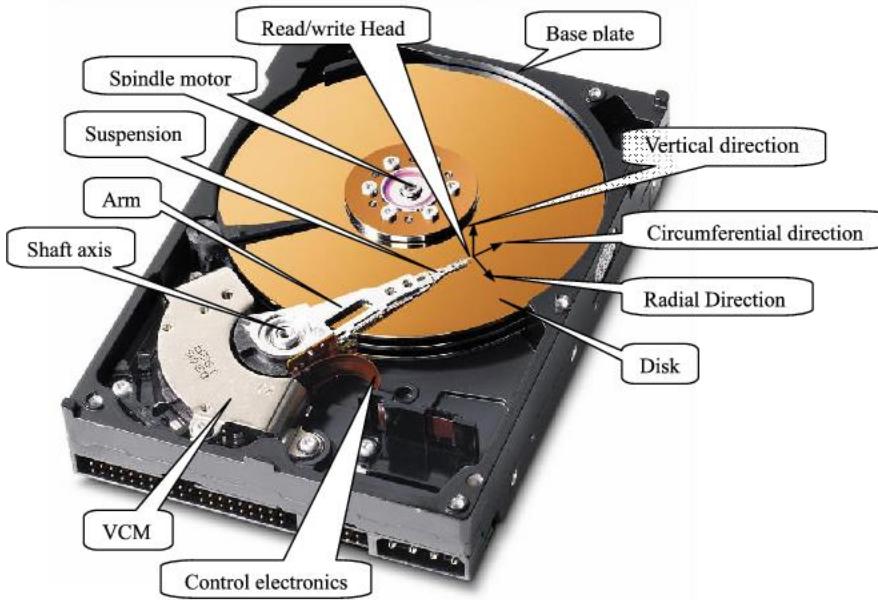
- Seagate Enterprise SATA III HDD
  - Rotations-per-minute = 7200
  - Average seek = 4.16 ms
  - Max transfer = 500 MB/s
- How long does an average 4KB read take?
  - $\text{Latency}_{\text{seek}} + \text{Latency}_{\text{rotate}} + \text{Latency}_{\text{transfer}}$
  - Transfer =  $\frac{1s}{500 * 1024 \text{ KB/s}} * 4 = 0.00000781248 \text{ seconds}$ 
    - = ~8 us
  - 4.16ms + 4.2 ms + 8 us = 8.368 ms



# A fun look at the oldest commercial HDD

- 1956 IBM RAMDAC computer
  - 5M (7-bit) characters
  - 50 x 24" platters
  - Access time = ~1 sec





# Disk (HDD) scheduling algorithms

# Disk scheduling introduction and requirements

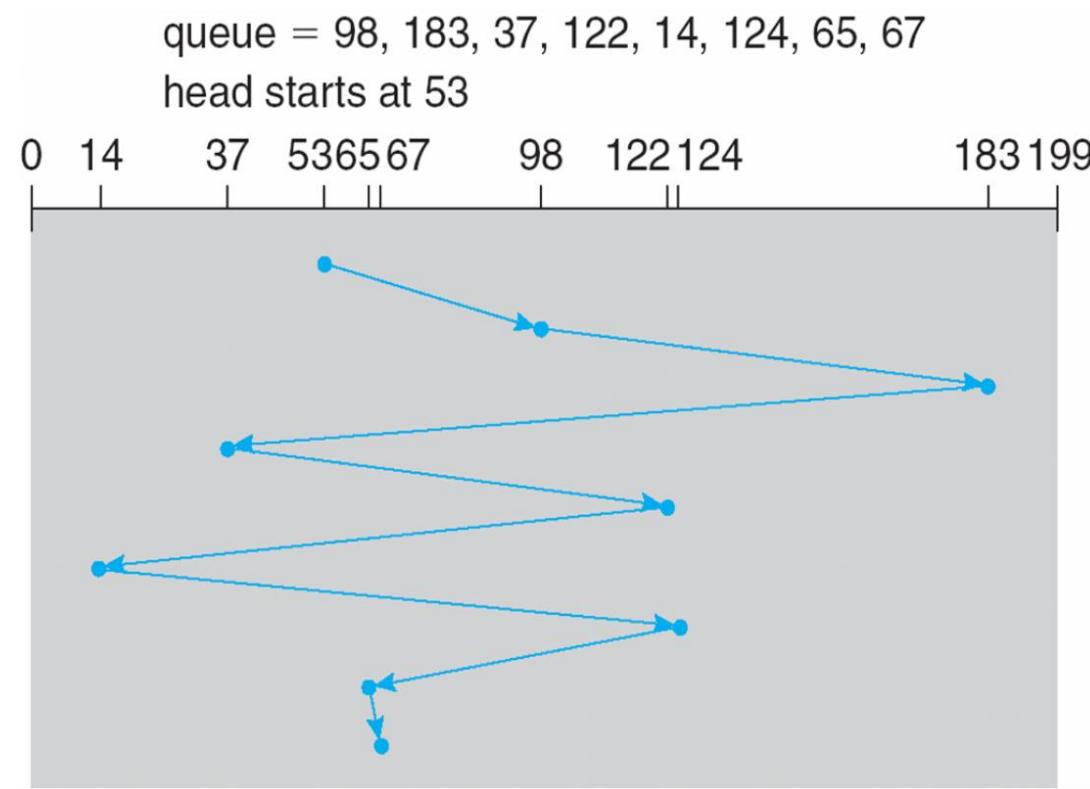
- The OS must answer: “Given a stream of I/O requests, in what order should they be served?”
- Strategy: **reorder** requests to meet certain goals
  - Performance (e.g., by making I/O sequential)
  - Fairness
  - Consistent latency
- **Performance objective:** minimize seek + rotation time
  - Minimize the distance the head needs to go

# Queuing and disk scheduling algorithms

- Note that drive controllers have small buffers and can manage a queue of I/O requests (of varying “depth”)
- OSs also maintain their own buffers in-memory before sending a request to the disk
- *Disk scheduling algorithms:*
  - Algorithms that schedule the orders of disk I/O requests

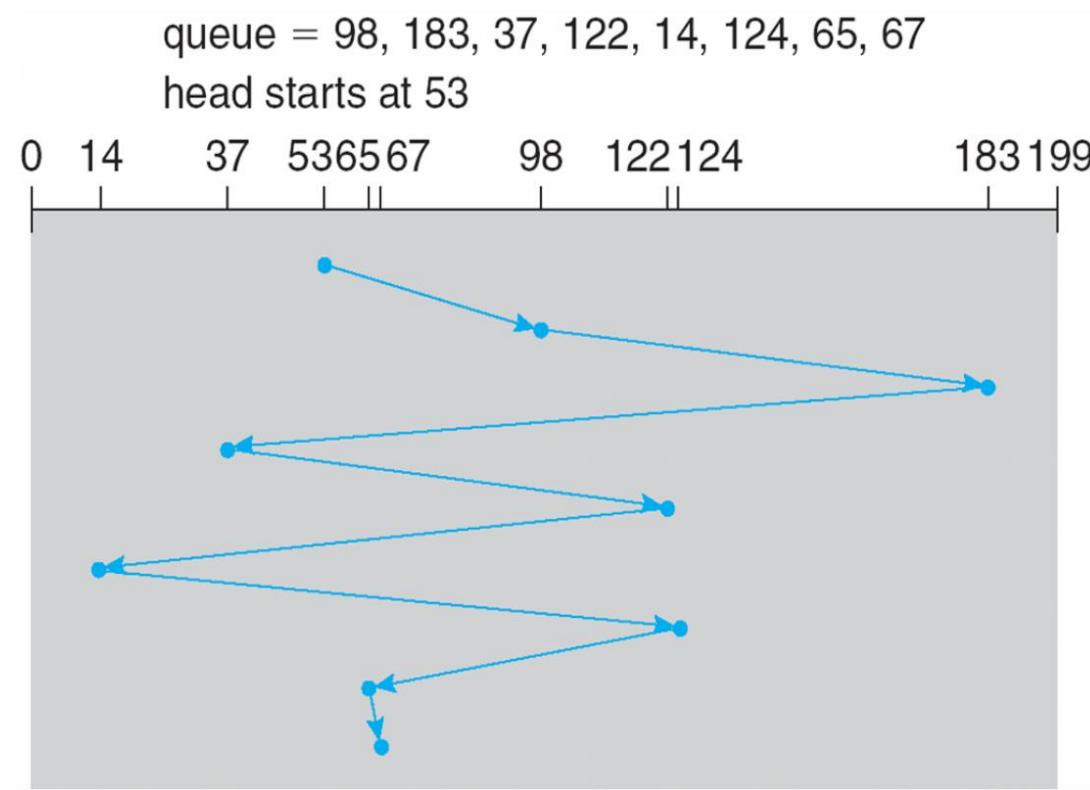
# The simplest approach: First-In-First-Out (FIFO)

- **Idea:** Serve the I/O request in the order they arrive



# Is there any problem with FIFO?

- **Idea:** Serve the I/O request in the order they arrive



**Problem: many total head movement of 640 cylinders**

# Calculating the number of head movements

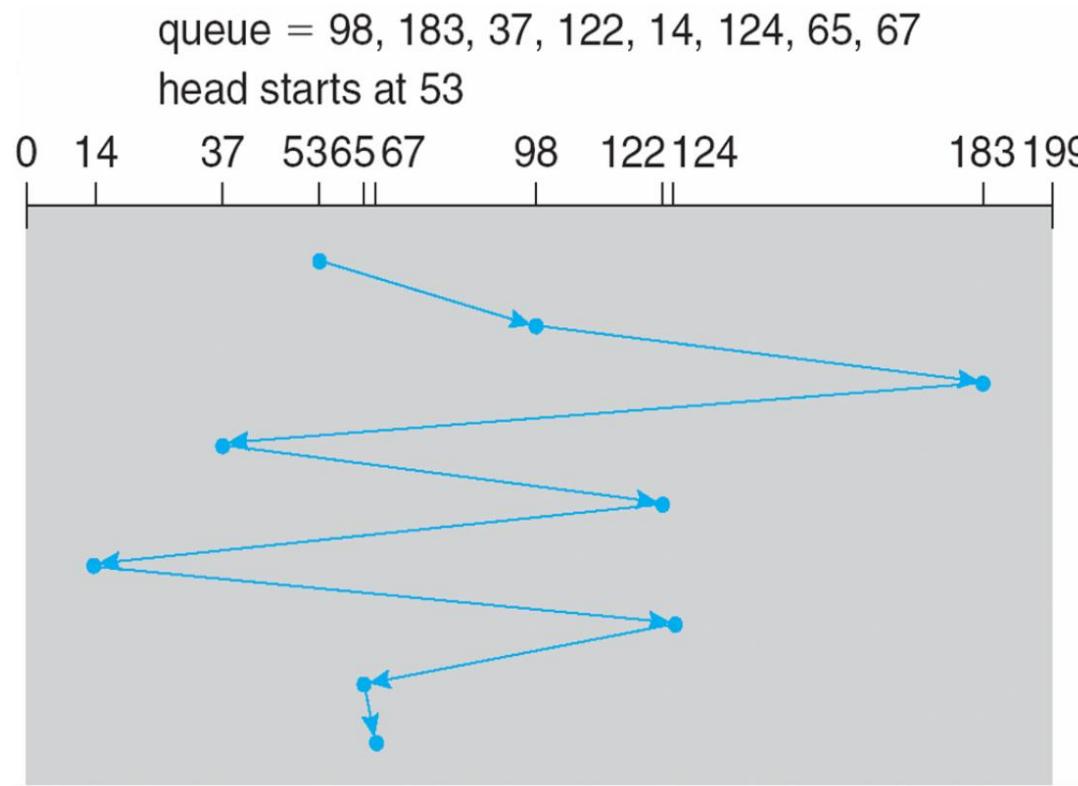
- **Idea:** Serve the I/O request in the order they arrive

**640 cylinders:**

- (98 – 53)
- (183 – 98)
- (183 – 37)
- (122 – 37)
- (122 – 14)
- (124 – 14)
- (124 – 65)
- (67 – 65)

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53



# How can we minimize head movement?

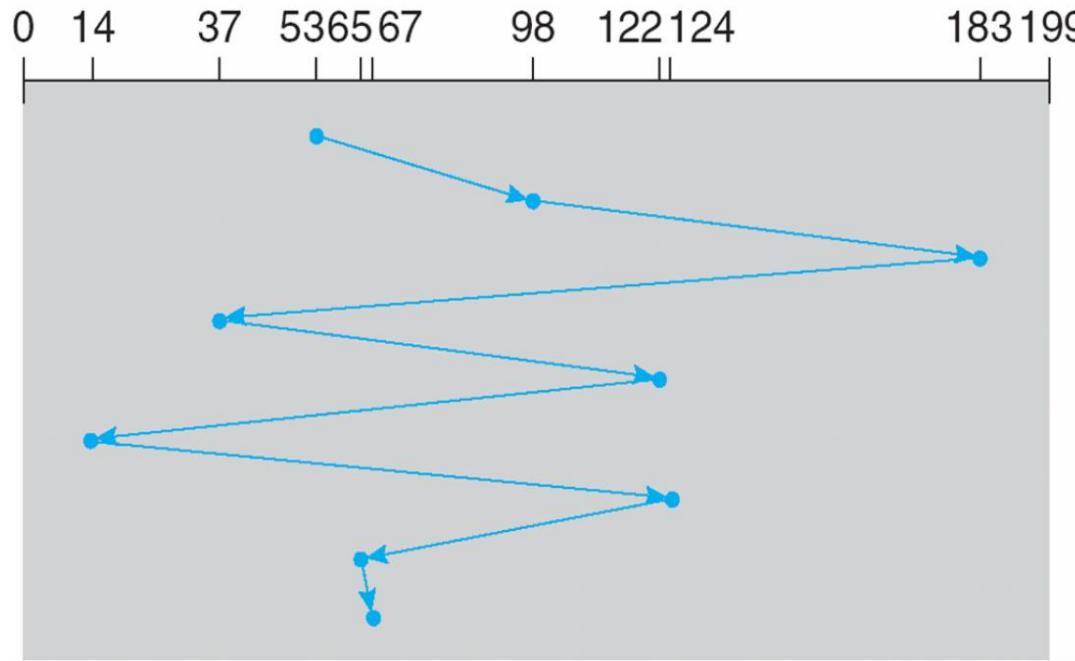
- **Idea:** Serve the I/O request in the order they arrive

## 640 cylinders:

- (98 – 53)
- (183 – 98)
- (183 – 37)
- (122 – 37)
- (122 – 14)
- (124 – 14)
- (124 – 65)
- (67 – 65)

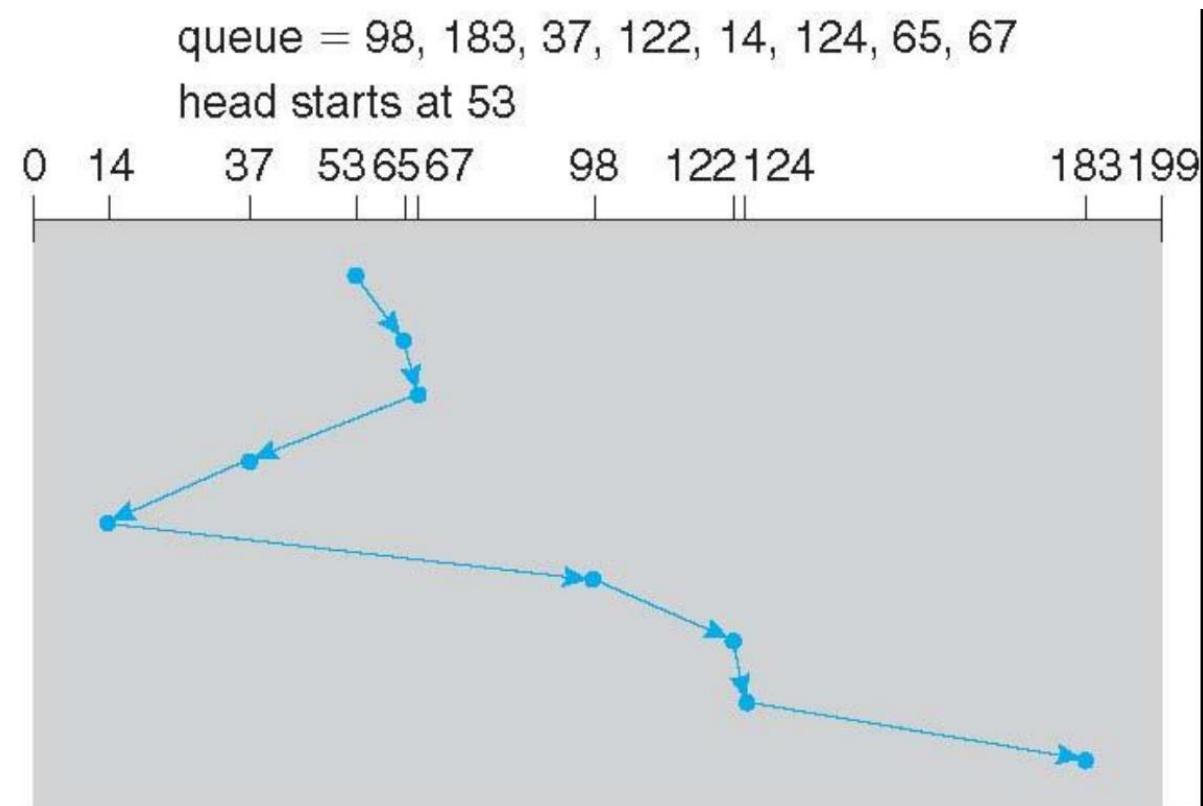
queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53



# *Shortest Positioning Time First (SPTF)*

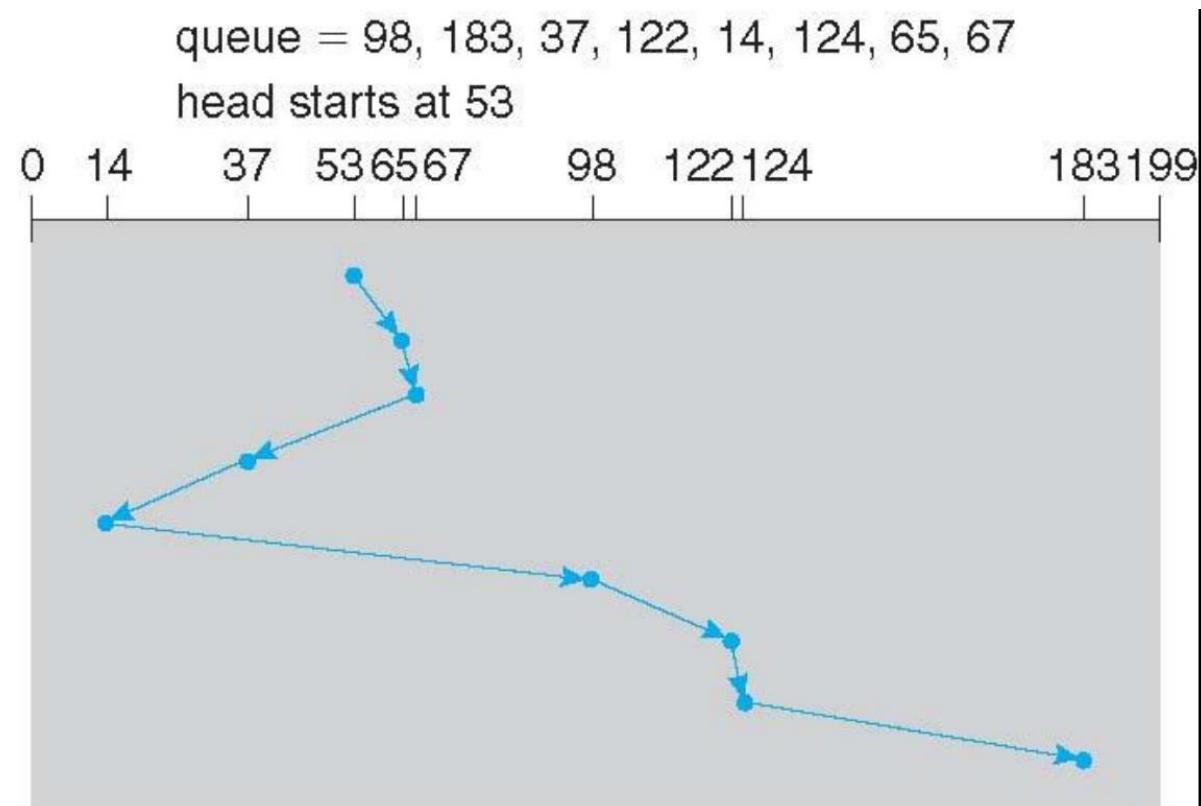
- **Idea:** Select the request with the least time for seeking and rotating



- Illustration shows total head movement of **236** cylinders.

# Is there any problem with SPTF?

- **Idea:** Select the request with the least time for seeking and rotating

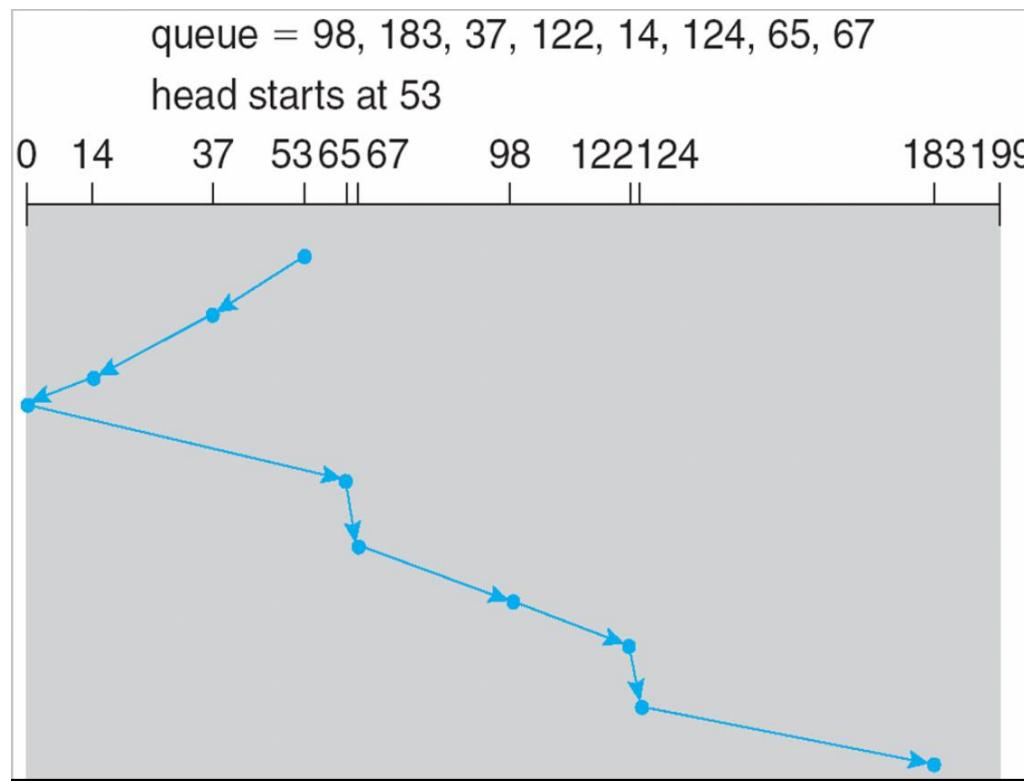


- Illustration shows total head movement of **236** cylinders.
- **Problem:** Like Shortest-Job-First (SJF) scheduling: may cause starvation of some requests!

# SCAN

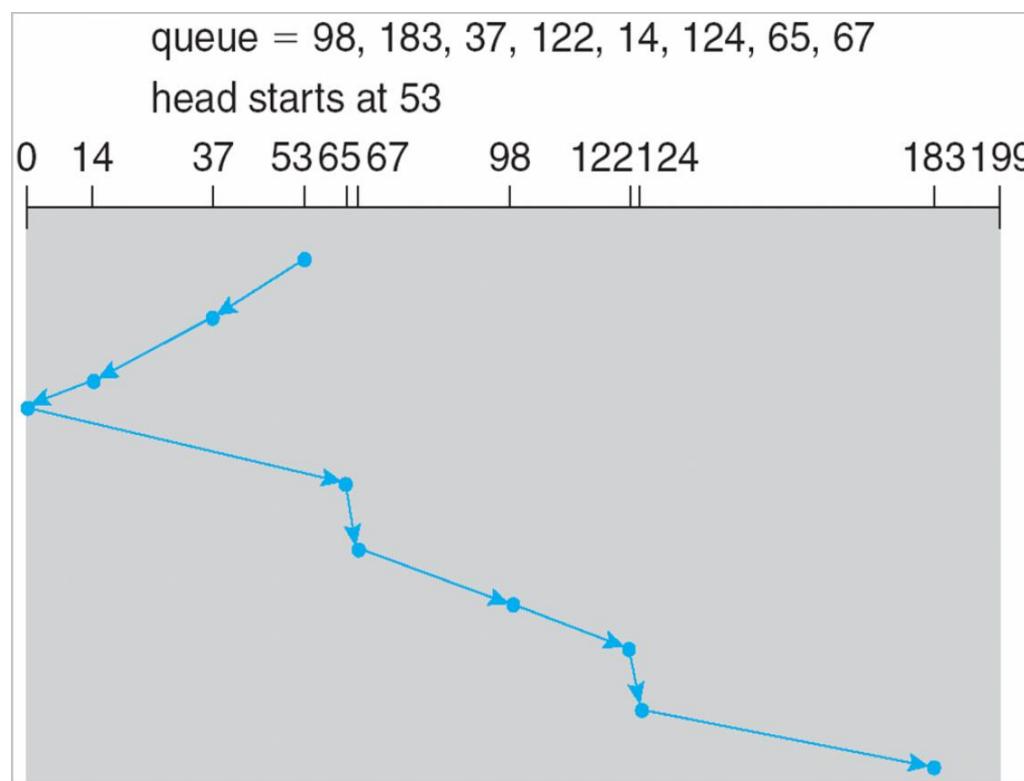
- **Idea:** Sweep back and forth, from one end of disk to the other, serving requests as you go (also called an “elevator” algorithm)

- Illustration shows total head movement of **236** cylinders.



# Is there any problem with SCAN?

- **Idea:** Sweep back and forth, from one end of disk to the other, serving requests as you go (also called an “elevator” algorithm)
- Illustration shows total head movement of **236** cylinders.
- **It is a bit unfair!**
  - Cylinders in the middle get better service, while at edges might have to wait for long periods of time.

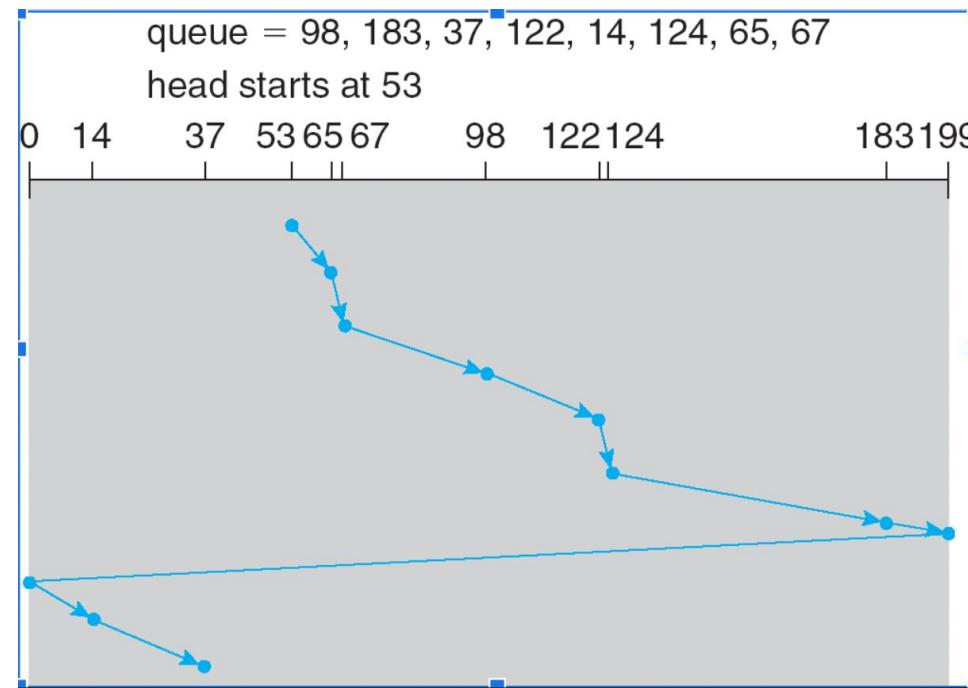


# C-SCAN

**Idea:** Only sweep in **ONE** direction.

- When it reaches the other end, however, it immediately returns to the beginning of the disk, without servicing any requests on the return trip

- Illustration shows total head movement of **382** cylinders.
- A bit slower than SCAN, but fairer to different computations

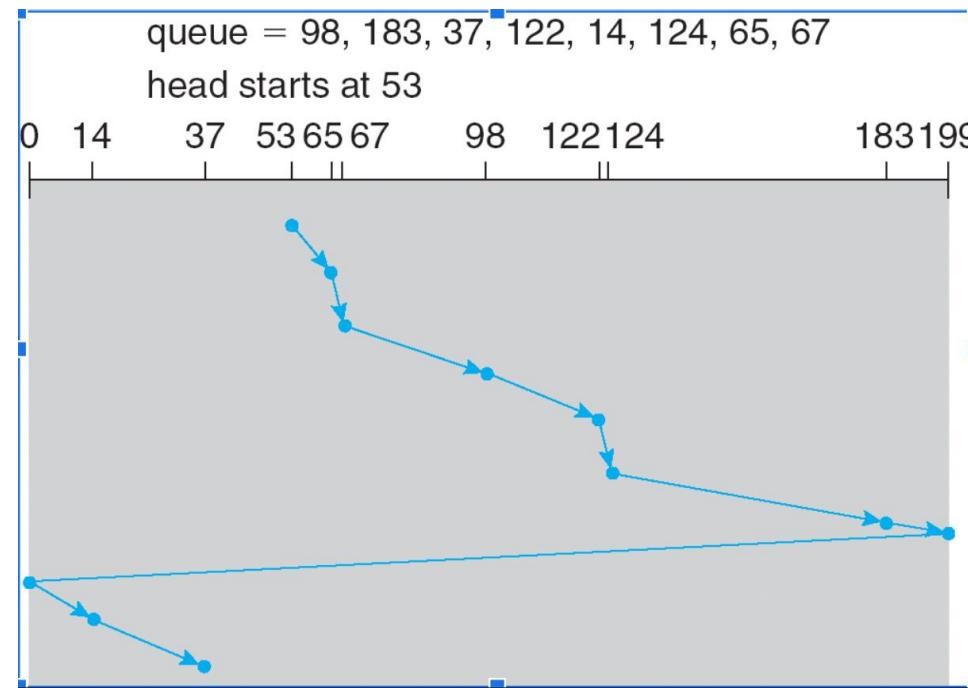


# Is there any problem with naive C-SCAN?

**Idea:** Only sweep in **ONE** direction.

- When it reaches the other end, however, it immediately returns to the beginning of the disk, without servicing any requests on the return trip

- **We are moving to the end of each direction (199 – 0), even though there is no such task**

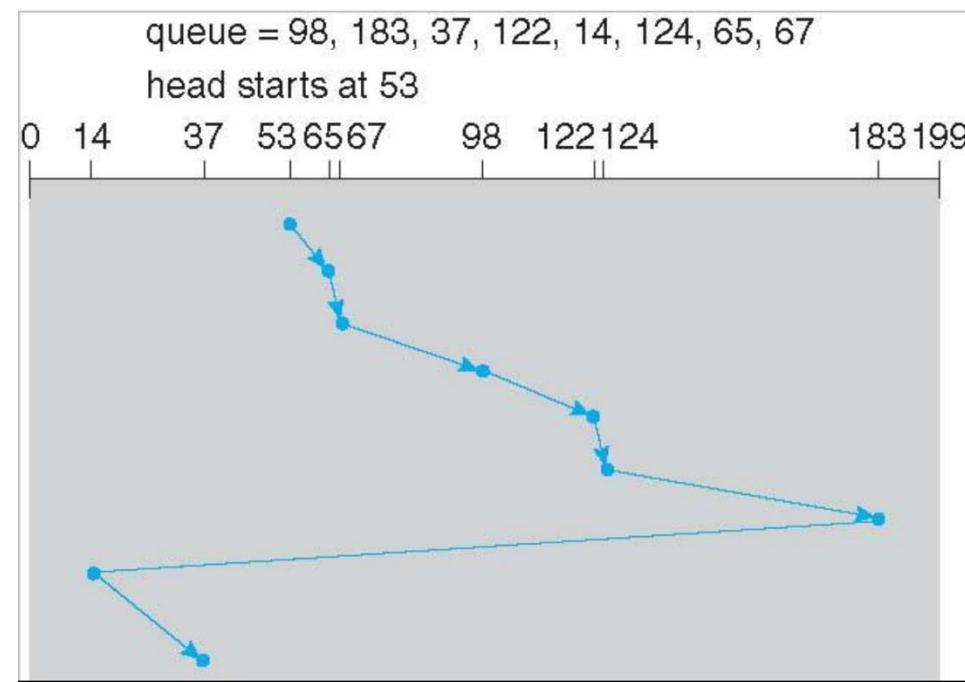


# C-LOOK

**Idea:** Only sweep in **ONE** direction, but also **look ahead**

- Arm only goes as far as the last request in each direction, then reverses direction immediately, without first going all the way to the end of the disk

- Illustration shows total head movement of **322** cylinders.
  - $(65-53)+(67-65)+(98-67)+(122-98)+(124-122)+(183-124)+(183-14)+(37-14) = 322$



# Work conservation versus anticipation in scheduling

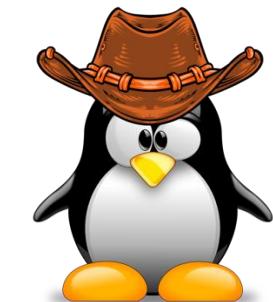
- **Work conserving** scheduling will always try to perform an I/O operation without any delay if
    - (a) there's an I/O request waiting to be handled
    - (b) the disk is ready to handle an I/O request
  - Sometimes, it is better to **wait (delay) instead if you anticipate** another request will appear nearby
    - E.g., consider the current queue: 0, 199
    - After a small time, there might be a request (e.g., 65)
- Such non-work-conserving schedulers are called *anticipatory*

# Default disk scheduling in Linux OSs: CFQ

- CFQ stands for “Completely Fair Queueing”
- OS keeps a I/O request queue for each process
- Do weighted round-robin among each process queue, with slice time proportional to process priority
- Optimize order within queue (e.g., using C-SCAN)
- Typically, disk scheduling in Linux is **anticipatory**

# Summarizing different HDD scheduling

- SPTF has a natural appeal towards minimizing average movements
  - **Starvation is a real problem**
- SCAN and C-SCAN perform better for systems that place a heavy load on the disk
  - **Better fairness since all I/O requests will be served after fixed time**
- Performance depends on the workload (i.e., number and types of requests)
- Requests for disk service can be impacted by the file-allocation method/pattern and metadata layout – **topic of file systems**



Questions? Otherwise, see you next class!