




CSE 330: Operating Systems


Adil Ahmad

Lecture #16: SSDs versus HDDs, higher-level abstractions



Project #4: brief introduction

An end-to-end memory access graph



The mmap system call in Linux

- **What is mmap?**
 - Used by processes to request the OS to allocate new memory
- `mmap (address, ..., permissions, size, ...)`
- Let's take a look at an example:
 - `mmap (0x1000, ..., READ | WRITE, 4096, ...)`
 - **What does this mean?**
 - Allocate 4096 bytes (1 page) of memory at virtual address 0x1000

Let's design the functionality of mmap together!


What steps would you follow to build mmap implementation?

```
mmap (0x1000, ..., READ | WRITE, 4096, ...)
```

- Step #0: check if the requested “addr” is already mapped
- Step #1: find a free physical page (or *frame*) to allocate
- Step #2: allocate (and map) the page tables
- Step #3: map the page frame

Step #0: Check if memory is already mapped at 0x1000


- **How do we know if memory is mapped at the address?**
 - One simple way is to traverse or *walk* the process' page tables!



How would you check that a PTE is “present”?

Step #0: Check if memory is already mapped at 0x1000


- **How do we know if memory is mapped at the address?**
 - One simple way is to traverse or *walk* the process' page tables!



How would you check that a PTE is “present”?

Step #0a: Checking if a page table entry is present

- Recall the layout of page table entries (PTEs):



If the PTE is present and mapped, we say page is **resident** in memory


Step #1: find free frames

- When a computer boots up, all frames not used by the OS are **free**
 - OS initializes an internal allocator that **keeps track of the location** (i.e., physical address) of all free frames
- In Linux, if you want to get a free frame from the internal allocator:
 - `va_frame = get_zeroed_page(. .)`
 - this will return the virtual address of a frame
- If you need the physical address, you can execute:
 - `__pa(va_frame)`

Step #2: Allocate a page table (if not allocated)

- Consider the following possible scenario:


Leve1-1 PT says
not present



Step #3: Allocate a page frame (if not allocated)


- Consider the following possible scenario:

Leve1-0 PT says
not present



Step #3a: Implement permissions in the bottom PTE

- Recall mmap (0x1000, ..., READ | WRITE, 4096, ...)
- In the bottom-level page table, you can set these permissions





Step #4: Keep track of pages allocated to process

- Required for later housekeeping (e.g., reclaiming/freeing certain pages from the process' address space)
 - `munmap` system call
- Future allocation requests:
 - Recall step #1 had us walking the page tables
 - If we keep track of what pages are allocated, no need to walk anymore!

Project #4: Build a custom `mmap` implementation!

- At a high-level, follow steps #0 – step #4 inside a kernel module
 - Please follow the ordering of steps mentioned in the document!!!!
- Use the helper functions/code provided to you:
 - Traversing (or walking) page tables
 - Checking if a page table entry is “present”
 - Allocating new frames using internal allocators
 - Setting different read/write/execute bits within PTEs
- **Note:** this is a hard assignment. Please start early.



Quick recap of HDDs and their scheduling

Storage performance overview

- Disk performance depends significantly on the working mechanism
- The disk latency can be calculated by considering the three steps that the HDD must perform:
 - HDD I/O latency = $\text{Latency}_{\text{seek}} + \text{Latency}_{\text{rotate}} + \text{Latency}_{\text{transfer}}$
- Let's see how many milliseconds are taken by these (*seek*, *rotate*, and *transfer*) operations!

Examining workload suitability for HDDs



- Seek and rotations are slow while transfer is relatively fast
 - seek = ~4 – 10 ms, rotate = ~4 – 5 ms, transfer = ~5 us
- **What kind of workload is best suited for disks?**
 - **Sequential I/O:** access sectors in order (transfer dominated)
- **Random** workloads access sectors in a random order (seek+rotation dominated); thus, slow on disks

Disk scheduling introduction and requirements

- The OS must answer: “Given a stream of I/O requests, in what order should they be served?”
- Strategy: **reorder** requests to meet certain goals
 - Performance (e.g., by making I/O sequential)
 - Fairness
 - Consistent latency
- **Performance objective:** minimize seek + rotation time
 - Minimize the distance the head needs to go


Summary of HDD scheduling

- Shortest Positioning Time First (SPTF) reduces seek and rotation
 - But, it leads to starvation
- SCAN, C-SCAN, and C-LOOK perform better for systems that place a heavy load on the disk
 - Also, provide better fairness
- Performance depends on the workload (i.e., number and types of requests)
- Linux leverages Completely Fair Queueing (CFQ)




Comparing HDDs and SSDs

Solid State Drive (SSD) overview




Solid State Drive (SSD) Internals



SSD Internals




Die Internals



4KB usually




Blocks and pages



One block

- Each flash array is made up of many blocks and this one block


Blocks and pages



Writing to a flash region


- All “1” in a block → block is **clean** and ready to be used
- Data is **written** by changing the “1” into a “0” (called **program**)
 - Operation at “page-level”
- To remove data, you must write back “1” (called **erase**)
 - Operation at “block-level” (block → a set of pages)
- Hence, **program** is much faster than **erase**

Writing to a flash region (illustration)




All pages are clean
("programmable")

Writing to a flash region (illustration)



Writing to a flash region (illustration)




Writing to a flash region (illustration)

1011	1111	1111	1111
1000	1111	1111	1111
1111	0110	1111	1111
1111	1010	1111	1111

Two pages hold data
(cannot be overwritten)


Writing to a flash region (illustration)

still want to write data into this page???




Two pages hold data
(cannot be overwritten)


Writing to a flash region (illustration)



Writing to a flash region (illustration)



Writing to a flash region (illustration)



After erase, again, **free state**
(can write new data in any page)


Writing to a flash region (illustration)

1011 0001	1111 1111	1111 1111	1111 1111
1111 1111	1111 1111	1111 1111	1111 1111


This blue page holds data

Comparison of APIs OS must support for HDDs and SSDs

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)



Device drivers



Device driver introduction

- Components of the kernel that support communication with different devices (e.g., HDDs, SSDs, Bluetooth, etc.)
- Typically, these can be installed on-demand during execution
- In Windows, device drivers are given a lower privilege (i.e., *ring-1/2*) than the kernel
 - Can access kernel memory but cannot execute privileged instructions

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)

Device driver functionality broken into **two halves**

- *Top half:* executed when the device is called by the process (e.g., open)
 - Process executes system call (e.g., open)
 - OS starts I/O request and waits for response from the device; sleeps
 - Executed during the context of I/O requesting process
- *Bottom half:* executed when the device signals an interrupt
 - Device has completed request
 - Could happen during the context of any process
 - Device interrupt handler is called, and it wakes-up the I/O requesting process
- Let's check **virtio-disk.c** for in the **xv6 OS** as an example

The xv6 operating system

- We've seen some code from the xv6 OS before, and we will see some more today related to it's block layer
- Let's briefly introduce the OS first:
 - xv6 is a “toy” operating system designed for teaching OSs
 - designed by professors at MIT in early 2005
- Great OS if you're interested in tinkering/hacking kernels

xv6 code: virtio_disk_rw (in kernel/virtio-disk.c)

Top-half execution of the virtio device driver during process' context

```
void  
virtio_disk_rw(struct buf *b, int write)  
{  
    uint64 sector = b->blockno * (BSIZE / 512);  
  
    acquire(&disk.vdisk_lock);  
    ...  
}
```

Executed on FS
read/write call

```
...  
  
_sync_synchronize();  
  
*R(VIRTIO_MMIO_QUEUE_NOTIFY) = 0; // value is queue number  
  
// Wait for virtio_disk_intr() to say request has finished.  
while(b->disk == 1) {  
    sleep(b, &disk.vdisk_lock);  
}
```

Notify the device using
MMIO register

Go to sleep while
waiting for response

xv6 code: virtio_disk_intr (in kernel/virtio-disk.c)

Bottom-half execution of the virtio device driver during any context

```
void  
virtio_disk_intr()  
{  
    acquire(&disk.vdisk_lock);
```


```
    // the device won't raise another interrupt until we tell it  
    // we've seen this interrupt, which the following line does.  
    // this may race with the device writing new entries to  
    // the "used" ring, in which case we may process the new  
    // completion entries in this interrupt, and have nothing to do  
    // in the next interrupt, which is harmless.  
    *R(VIRTIO_MMIO_INTERRUPT_ACK) = *R(VIRTIO_MMIO_INTERRUPT_STATUS) & 0x3;
```

Executed in response
to interrupt

Acknowledge interrupt
but nothing else

```
    struct buf *b = disk.info[id].b;  
    b->disk = 0; // disk is done with buf  
    wakeup(b);
```

Wake-up requesting process
and ask it to handle response



Moving to higher level abstractions

Can anyone tell me a problem with having different interfaces?

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)

- Each new storage disk could have its own R/W mechanism
- OS developers would have to cater individually to the specific APIs of a device


This adds **complexity** to using disks for any general task (e.g., writing a file system)

What's the solution to this “many interface” problem?

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)


Create a **unified interface** and translate all operations from different interfaces to the unified one

Visualizing a unified interface for storage access



Once again, OS does the plumbing in computers


Provides an abstraction for the messy hardware



- Manipulating hardware requires understanding of the hardware internals
- Abstractions simplify using the hardware and hide away all these messy details

The block abstraction layer

- We want storage devices to allow **random** access of sectors or **blocks** (typically 512 or 4096 bytes)
- Simplify accessing all these devices with a **block-level** abstraction
 - Higher-level OS code can directly interact with blocks, and the lower-level device driver can support that interaction
- Generally, two operations:
 - `bread(device, x, buffer)` → read from device at sector X into buffer
 - `bwrite(device, y, buffer)` → write buffer into sector Y of device



Questions? Otherwise, see you next class!