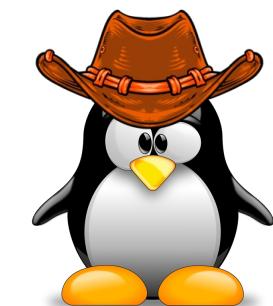




CSE 330: Operating Systems

Adil Ahmad

Lecture #17: The block abstraction and the file system

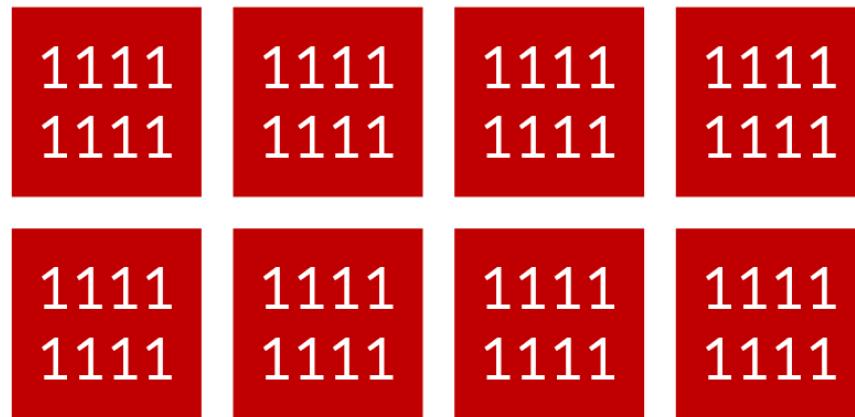


Recap of last week

Writing to a flash region

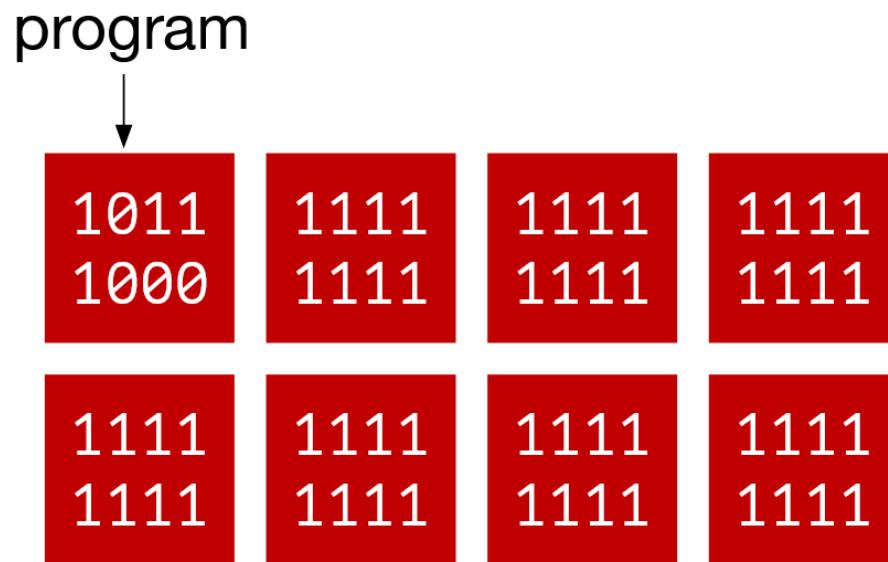
- All “1” in a block → block is **clean** and ready to be used
- Data is **written** by changing the “1” into a “0” (called **program**)
 - Operation at “page-level”
- To remove data, you must write back “1” (called **erase**)
 - Operation at “block-level” (block → a set of pages)
- Hence, **program** is much faster than **erase**

Writing to a flash region (illustration)

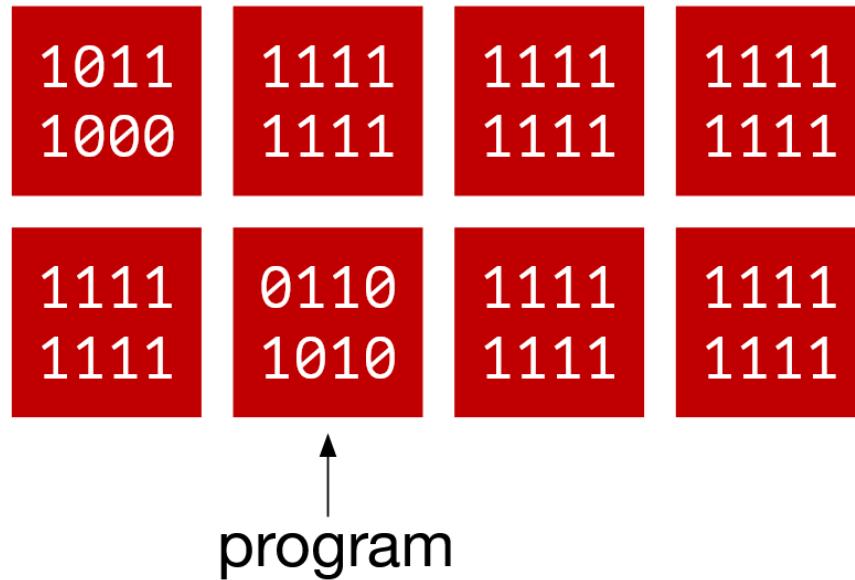


All pages are clean
("programmable")

Writing to a flash region (illustration)



Writing to a flash region (illustration)



Writing to a flash region (illustration)

1011	1111	1111	1111
1000	1111	1111	1111
1111	0110	1111	1111
1111	1010	1111	1111

Two pages hold data
(cannot be overwritten)

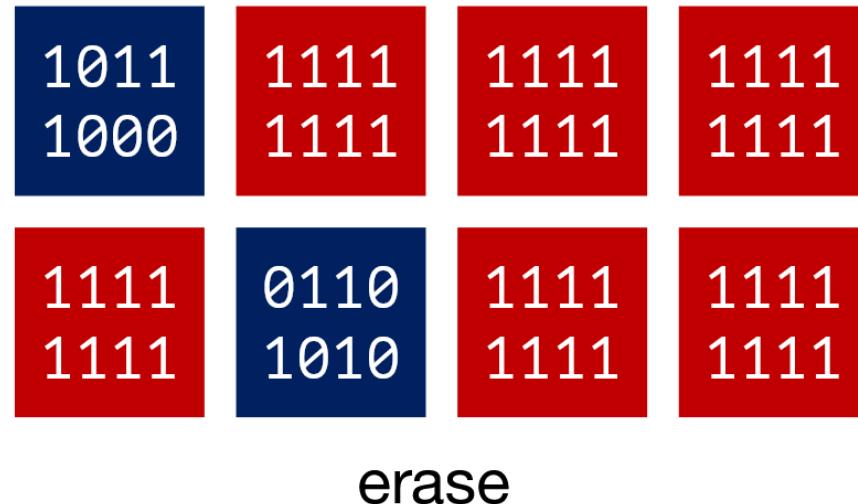
Writing to a flash region (illustration)

still want to write data into this page???

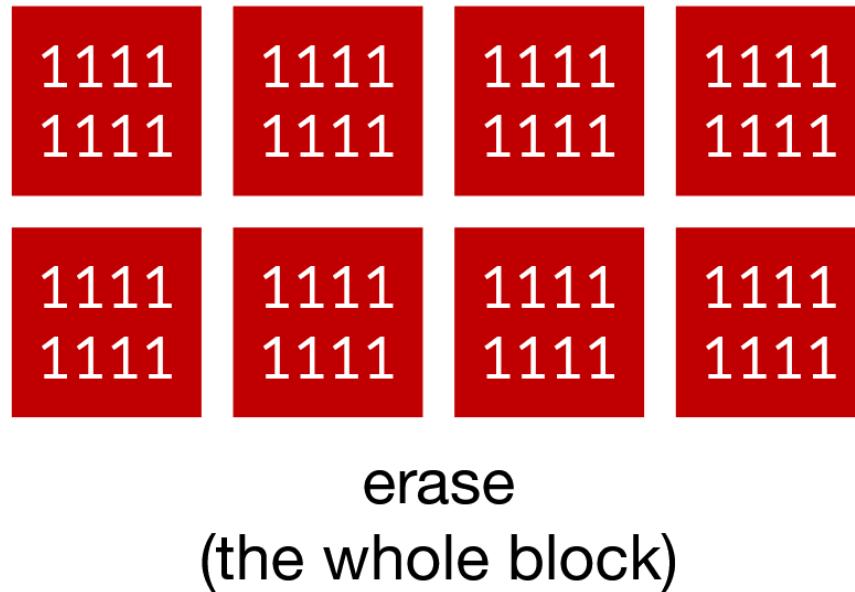
1011	1111	1111	1111
1000	1111	1111	1111
1111	0110	1111	1111
1111	1010	1111	1111

Two pages hold data
(cannot be overwritten)

Writing to a flash region (illustration)



Writing to a flash region (illustration)



Writing to a flash region (illustration)

1111	1111	1111	1111
1111	1111	1111	1111
1111	1111	1111	1111
1111	1111	1111	1111

After erase, again, **free state**
(can write new data in any page)

Writing to a flash region (illustration)

1011 0001	1111 1111	1111 1111	1111 1111
1111 1111	1111 1111	1111 1111	1111 1111

This blue page holds data

Comparison of APIs OS must support for HDDs and SSDs

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)

Can anyone tell me a problem with having different interfaces?

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)

- Each new storage disk could have its own R/W mechanism
- OS developers would have to cater individually to the specific APIs of a device

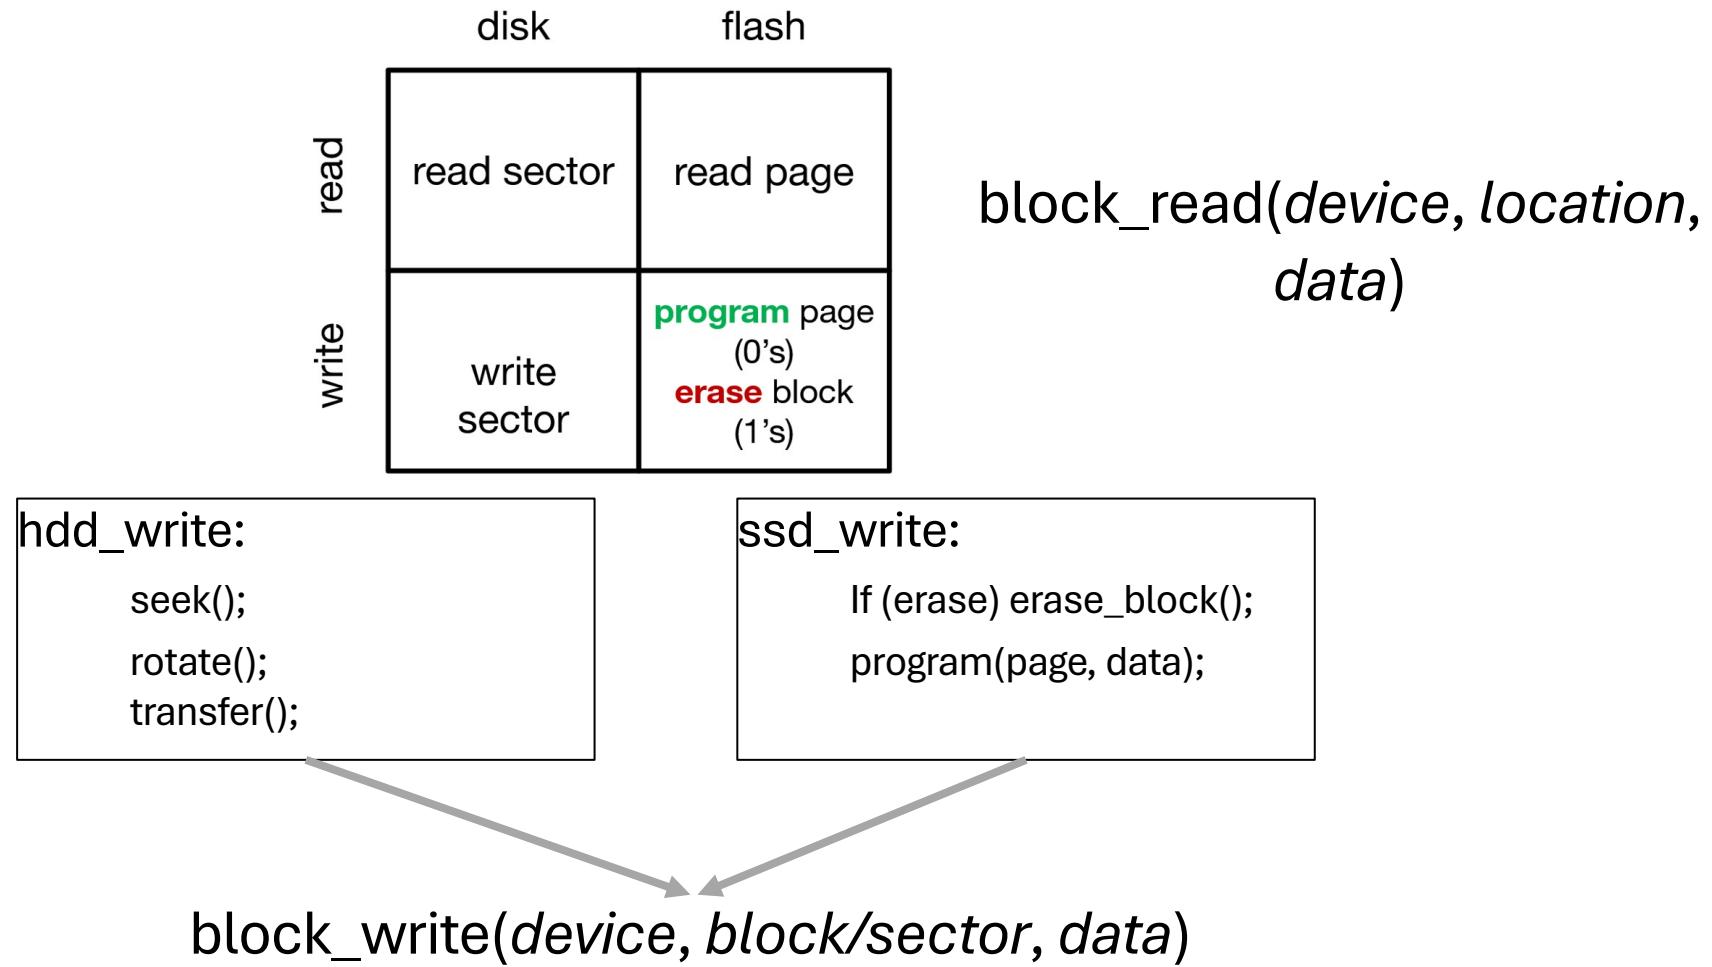
This adds **complexity** to using disks for any general task (e.g., writing a file system)

What's the solution to this “many interface” problem?

	disk	flash
read	read sector	read page
write	write sector	program page (0's) erase block (1's)

Create a **unified interface** and translate all operations from different interfaces to the unified one

Visualizing a unified interface for storage access

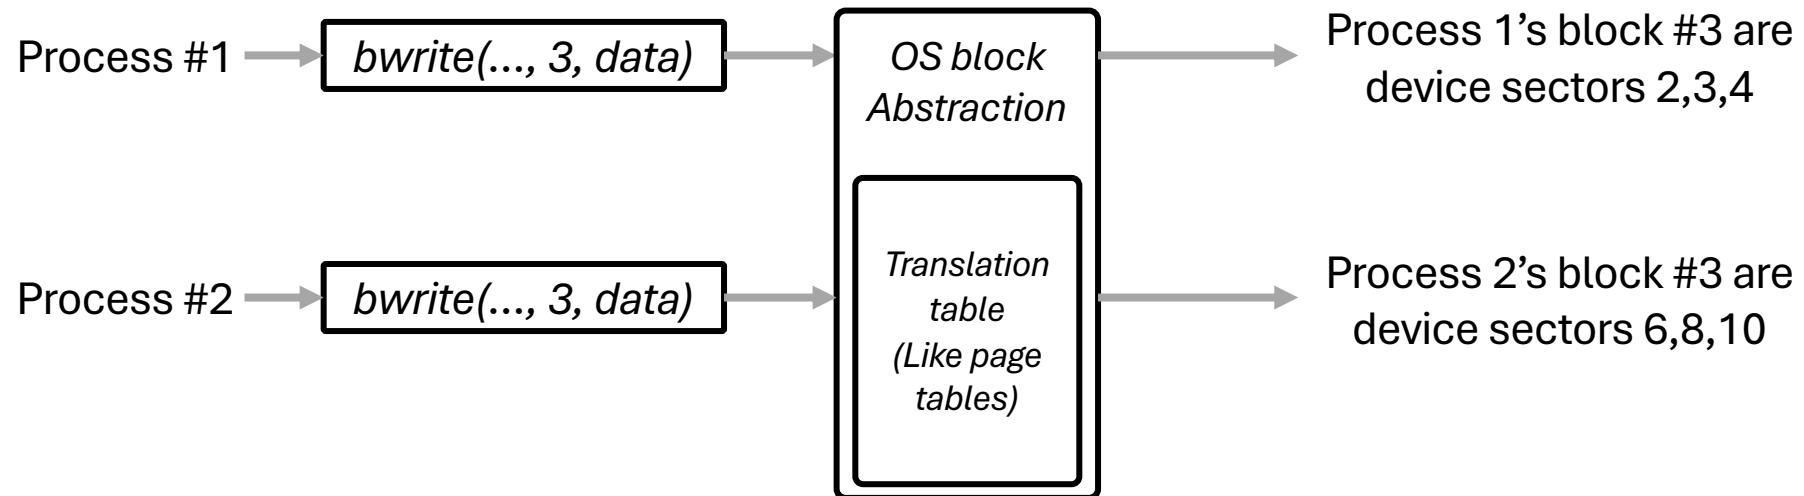


The block abstraction layer

- We want storage devices to allow **random** access of sectors or **blocks** (typically 512 or 4096 bytes)
- Simplify accessing all these devices with a **block-level** abstraction
 - Higher-level OS code can directly interact with blocks, and the lower-level device driver can support that interaction
- Generally, two operations:
 - `bread(device, x, buffer)` → read from device at sector X into buffer
 - `bwrite(device, y, buffer)` → write buffer into sector Y of device

Apart from simplicity, another benefit of block abstraction?

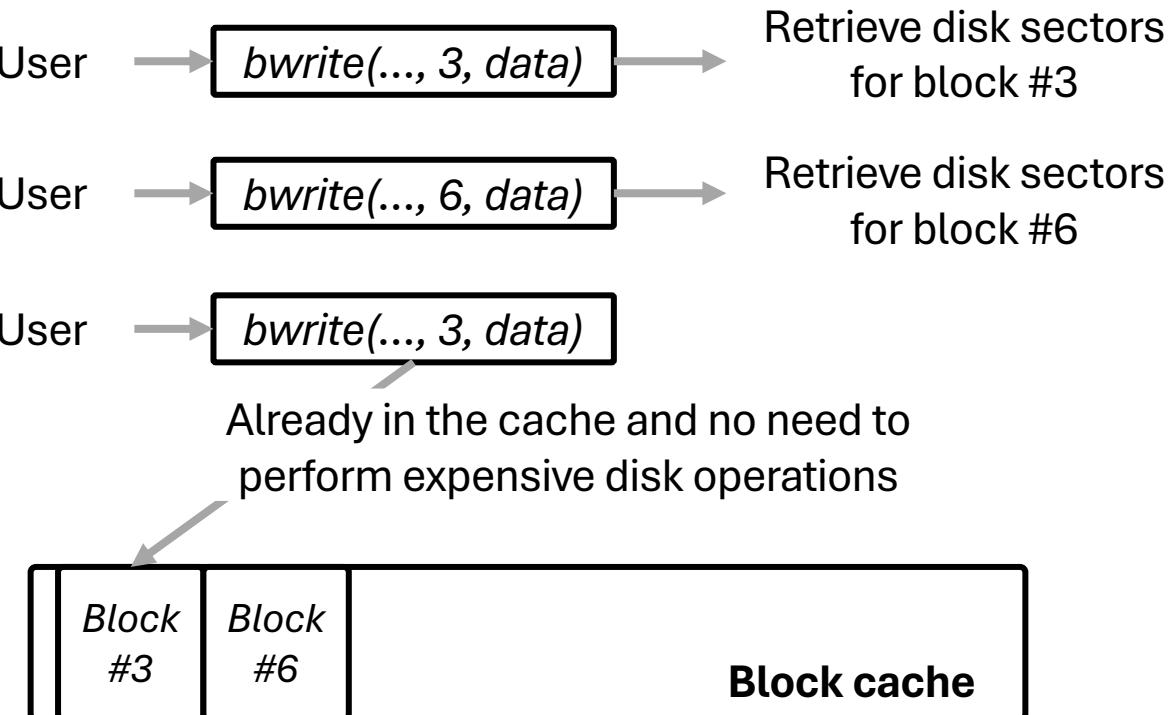
- **Logical disk view:** OS can virtualize how a user thinks of disk regions, and even combine disk “sectors” into larger “blocks”



- Multiple users/processes can simply operate on block #1 → #N

Apart from simplicity, another benefit of block abstraction?

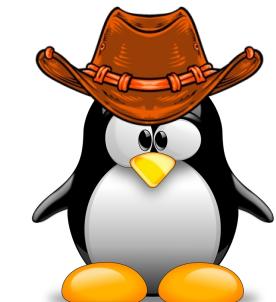
- **Caching efficiency:** Conveniently keep track of frequently-used disk regions by users and keep them in-memory for later use



Are there any disadvantages of the block abstractions?

➤ **Sacrifice raw device I/O performance**

- Each layer of abstraction requires translation → many more function calls and memory accesses
- Much faster to directly access block-level storage for special scenarios (e.g., where disk writes are needed)



Let's review a block layer implementation

High-level working of the **xv6 OS** block layer

- **Step #1:** Initializes a list of in-memory block cache
- **Step #2:** Reads data from disk into an in-memory block cache, while synchronizing block access with other threads
- **Step #3:** Performs a write (if needed) on an in-memory block
 - Similar synchronization as step #2

Block layer initialization (xv6/kernel/bio.c)

```
struct {
    struct spinlock lock;
    struct buf buf[NBUF];
} bcache;

// Linked list of all buffers, through prev/next.
// Sorted by how recently the buffer was used.
// head.next is most recent, head.prev is least.
struct buf head;
} bcache;

void
binit(void)
{
    struct buf *b;

    initlock(&bcache.lock, "bcache");

    // Create linked list of buffers
    bcache.head.prev = &bcache.head;
    bcache.head.next = &bcache.head;
    for(b = bcache.buf; b < bcache.buf+NBUF; b++){
        b->next = bcache.head.next;
        b->prev = &bcache.head;
        initsleeplock(&b->lock, "buffer");
        bcache.head.next->prev = b;
        bcache.head.next = b;
    }
}
```

Synchronize using a lock

Why use spinlock here?

Linked list of blocks to cache, and keep *head*

Initialize the spinlock

Why use sleeplock here?

Initialize the sleeplock

Read a block from disk (xv6/kernel/bio.c)

```
// Return a locked buf with the contents of the indicated block.  
struct buf*  
bread(uint dev, uint blockno)  
{  
    struct buf *b;  
  
    b = bget(dev, blockno);  
    if(!b->valid) {  
        virtio_disk_rw(b, 0);  
        b->valid = 1;  
    }  
    return b;  
}
```

Get a free block

If block is empty,
fetch from disk

```
// Look through buffer cache for block on device dev.  
// If not found, allocate a buffer.  
// In either case, return locked buffer.  
static struct buf*  
bget(uint dev, uint blockno)  
{  
    struct buf *b;
```

```
    acquire(&bcache.lock);
```

Acquire spinlock

```
    // Is the block already cached?  
    for(b = bcache.head.next; b != &bcache.head; b = b->next){  
        if(b->dev == dev && b->blockno == blockno){  
            b->refcnt++;  
            release(&bcache.lock);  
            acquiresleep(&b->lock);  
            return b;  
        }  
    }
```

Go through list

```
    // Not cached.  
    // Recycle the least recently used (LRU) unused buffer.  
    for(b = bcache.head.prev; b != &bcache.head; b = b->prev){  
        if(b->refcnt == 0) {  
            b->dev = dev;  
            b->blockno = blockno;  
            b->valid = 0;  
            b->refcnt = 1;  
            release(&bcache.lock);  
            acquiresleep(&b->lock);  
            return b;  
        }  
    }  
    panic("bget: no buffers");
```

Release spl and
get sleeplock

Go through list again,
but find free block

Write a block from disk (xv6/kernel/bio.c)

bread(..) should always happen first, before writing according to xv6

Why is this the case?

All block operations happen in the “cache” before being written to disk

```
// Write b's contents to disk. Must be locked.  
void  
bwrite(struct buf *b)  
{  
    if(!holdingsleep(&b->lock))  
        panic("bwrite");  
    virtio_disk_rw(b, 1);  
}
```

Make sure lock is being held

Perform the actual write

Finalizing the block-level abstraction

Provides us a way to simplify “direct” writes to different storage devices

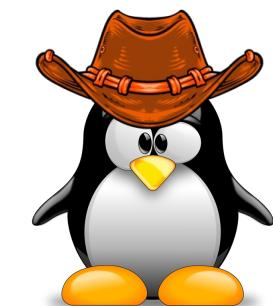
HDD → seek, rotate, transfer → bread(dev, block, data);

SSD → program, erase → bread(dev, block, data);

The xv6 block abstraction also implements the following:

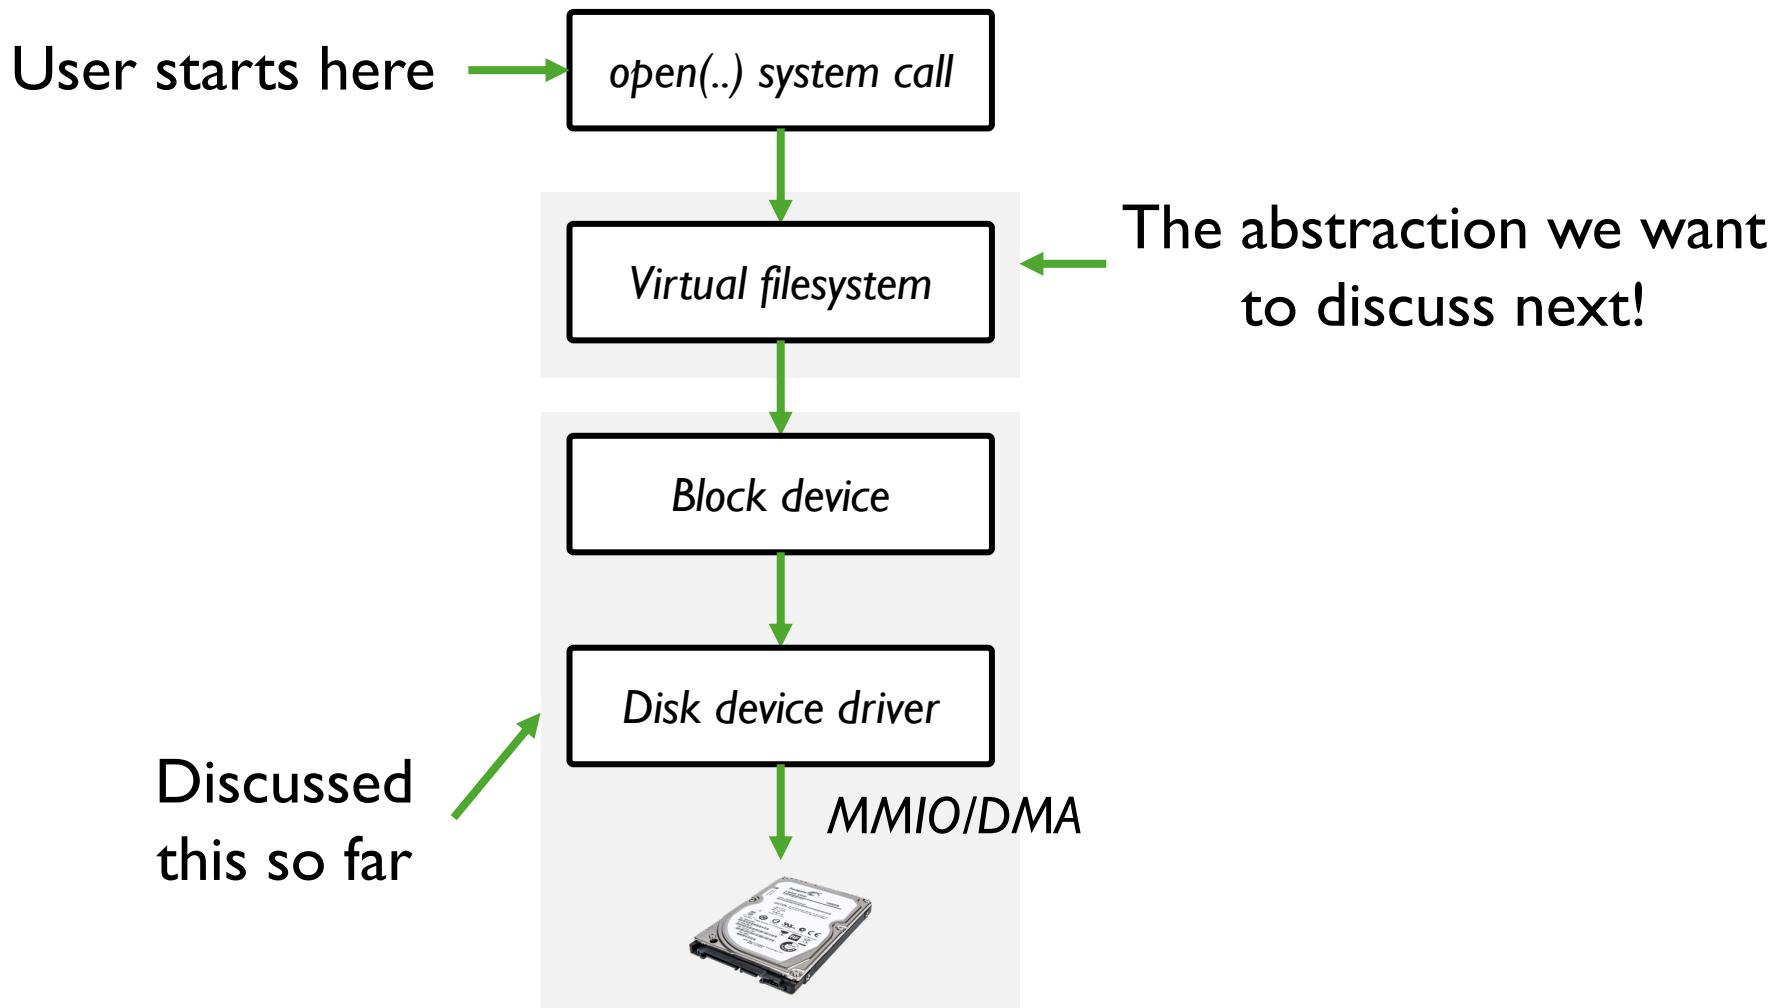
- **Synchronization** b/w block-level access of a device to avoid corruptions
- **Caching** of recently-used blocks for improved performance

Is the block abstraction enough for development?



The (virtual) file system

Disk interactions starting from a system call



The (virtual) filesystem abstraction

- High-level **intuitive** view of how we look at data stored on disks
- Not just disk-related; UNIX philosophy → “**everything is a file**”

Can anyone tell me of other ways in which you have used files?

- High-level **intuitive** view of how we look at data stored on disks
- Not just disk-related; UNIX philosophy → “**everything is a file**”
 - `/dev/memalloc` → virtual file to communicate with modules
 - **CPU features** can be enabled or disabled using files.
 - E.g., entire CPUs can be disabled as follows:
`echo 0 | sudo tee /sys/devices/system/cpu/cpu1/online`
 - **Perform console R/W** → `write(1, "hello", 5)`
 - `1` → file descriptor for console (terminal) output

Let's get back to (disk-related) file

- File → a set of blocks that the OS has combined and operates on together
 - Recall that the filesystem is composed of the the block layer
- Files are given identifiers (e.g., “hello.txt” is a human-readable version) so programs/users can distinguish between them
- File system (FS) → an intricate hierarchical collection of files built using the blocks in your storage disk

Different “names” for a file

- Three different names typically
 - ✓ **inode** (low-level names)
 - Internal name (number) given to a file by the OS
 - ✓ **path** (human readable)
 - The version that we see when we open the file browser (e.g., Windows explorer or MacOS finder)
 - ✓ **file descriptor** (runtime state)
 - Represents the runtime status of a certain file

The inode (OS-level representation)

- Each file has exactly one inode number
- Inodes are unique (at a given time) within a FS
- **File names can be the same, why can't inodes?**
 - Something must be unique for the OS to track

PROMPT>: stat test.dat

File: 'test.dat'

Size: 5

Blocks: 8

IO Block: 4096 regular file

Device: 803h/2051d

Inode: 119341128

Links: 1

Access: (0664/-rw-rw-r--)

Uid: (1001/ yue)

Gid: (1001/ yue)

Context: unconfined_u:object_r:user_home_t:s0

Access: 2015-12-17 04:12:47.935716294 -0500

Modify: 2014-12-12 19:25:32.669625220 -0500

Change: 2014-12-12 19:25:32.669625220 -0500

Birth: -

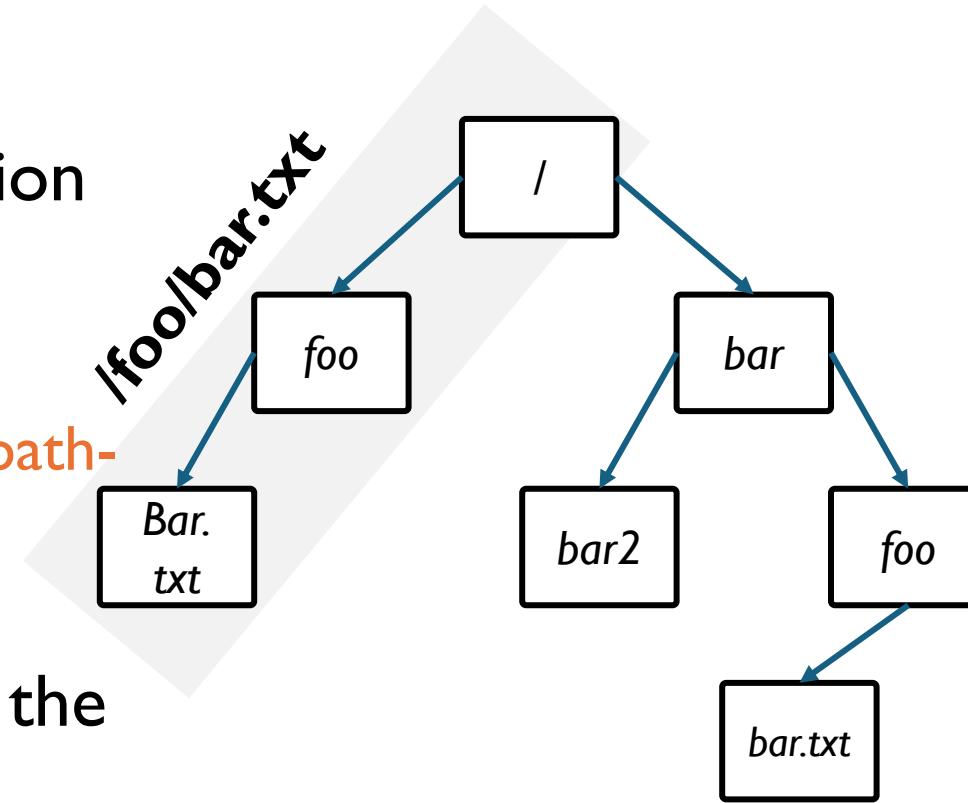
stat example

```
PROMPT>: stat test.dat
```

```
  File: 'test.dat'  Size: 5      Blocks: 8      IO Block:  
4096   regular file  
Device: 803h/2051d Inode: 119341128  Links: 1  
Access: (0664/-rw-rw-r--) Uid: ( 1001/      yue)  Gid: (  
1001/      yue)  
Context: unconfined_u:object_r:user_home_t:s0  
Access: 2015-12-17 04:12:47.935716294 -0500  
Modify: 2014-12-12 19:25:32.669625220 -0500  
Change: 2014-12-12 19:25:32.669625220 -0500  
Birth: -
```

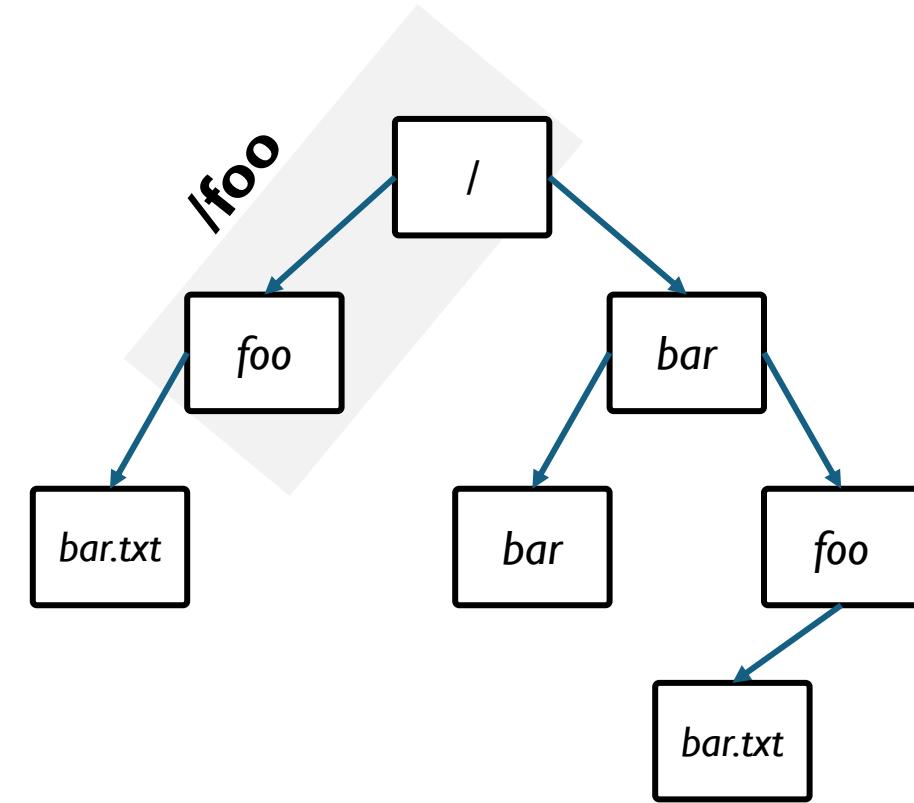
The path (human-readable)

- Human-readable interpretation of every inode
- Typically, represented as – <path-to-directory , filename>
- **Traversing** a tree – getting the final *inode* from a location
 - E.g., ls /foo/bar.txt
 - Gets the inode and prints details



The path (human-readable)

- Directories and files can have the same name if they are in different tree locations
- Cannot have the same name if they're in the same directory
 - Not really a problem for OSs, but a problem for humans to distinguish between ☺



`ls` example

```
prompt> ls -al
```

```
total 216
```

```
drwxr-xr-x 19 yue staff 646 Nov 23 16:28 .
drwxr-xr-x+ 40 yue staff 1360 Nov 15 01:41 ..
-rw-r--r--@ 1 yue staff 1064 Aug 29 21:48 common.h
-rwxr-xr-x 1 yue staff 9356 Aug 30 14:03 cpu
-rw-r--r--@ 1 yue staff 258 Aug 29 21:48 cpu.c
-rwxr-xr-x 1 yue staff 9348 Sep 6 12:12 cpu_bound
-rw-r--r-- 1 yue staff 245 Sep 5 13:10 cpu_bound.c
...
```

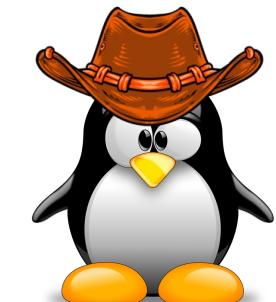
- **Two additional files in every folder. Can anyone tell me what these are?**

The file descriptor

- “Everything is a file”
 - File descriptor tracks what each ‘file’ really does in the system

```
struct file {  
    enum { FD_NONE, FD_PIPE, FD_INODE, FD_DEVICE } type;  
    int ref; // reference count  
    char readable;  
    char writable;  
    struct pipe *pipe; // FD_PIPE  
    struct inode *ip; // FD_INODE and FD_DEVICE  
    uint off; // FD_INODE  
    short major; // FD_DEVICE  
};
```

If file belongs to the disk, it has an inode



Questions? Otherwise, see you next class!