



CSE 330: Operating Systems

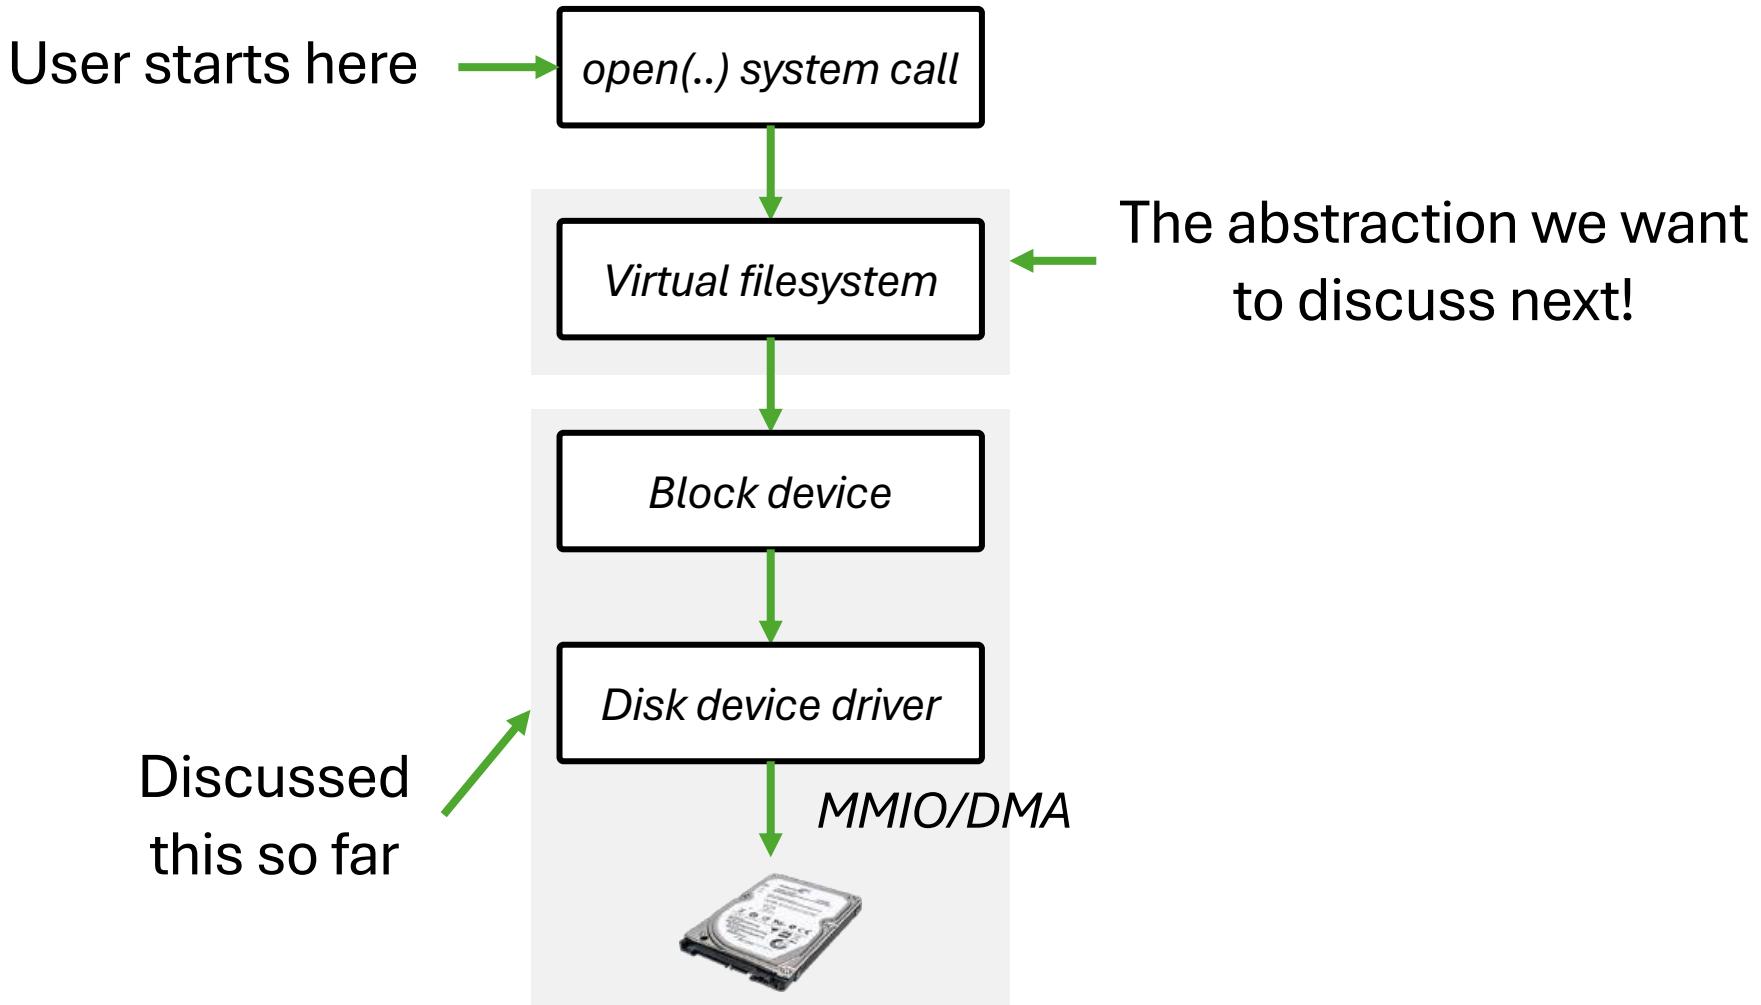
Adil Ahmad

Lecture #18: File system overview and internals



The (virtual) file system recap

Disk interactions starting from a system call



The (virtual) filesystem abstraction

- High-level **intuitive** view of how we look at data stored on disks
- Not just disk-related; UNIX philosophy → “**everything is a file**”

Can anyone tell me of other ways in which you have used files?

- High-level **intuitive** view of how we look at data stored on disks
- Not just disk-related; UNIX philosophy → “**everything is a file**”
 - `/dev/memalloc` → virtual file to communicate with modules
 - **CPU features** can be enabled or disabled using files.
 - E.g., entire CPUs can be disabled as follows:
`echo 0 | sudo tee /sys/devices/system/cpu/cpu1/online`
 - **Perform console R/W** → `write(1, "hello", 5)`
 - 1 → file descriptor for console (terminal) output

Let's get back to (disk-related) file

- File → a set of blocks that the OS has combined and operates on together
 - Recall that the filesystem is composed of the the block layer
- Files are given identifiers (e.g., “hello.txt” is a human-readable version) so programs/users can distinguish between them
- File system (FS) → an intricate hierarchical collection of files built using the blocks in your storage disk

Different “names” for a file

- Three different names typically
 - ✓ **inode** (low-level names)
 - Internal name (number) given to a file by the OS
 - ✓ **path** (human readable)
 - The version that we see when we open the file browser (e.g., Windows explorer or MacOS finder)
 - ✓ **file descriptor** (runtime state)
 - Represents the runtime status of a certain file

The inode (OS-level representation)

- Each file has exactly one inode number
- Inodes are unique (at a given time) within a FS
- **File names can be the same, why can't inodes?**
 - Something must be unique for the OS to track

PROMPT>: stat test.dat

File: 'test.dat'

Size: 5

Blocks: 8

IO Block: 4096 regular file

Device: 803h/2051d

Inode: 119341128

Links: 1

Access: (0664/-rw-rw-r--)

Uid: (1001/ yue)

Gid: (1001/ yue)

Context: unconfined_u:object_r:user_home_t:s0

Access: 2015-12-17 04:12:47.935716294 -0500

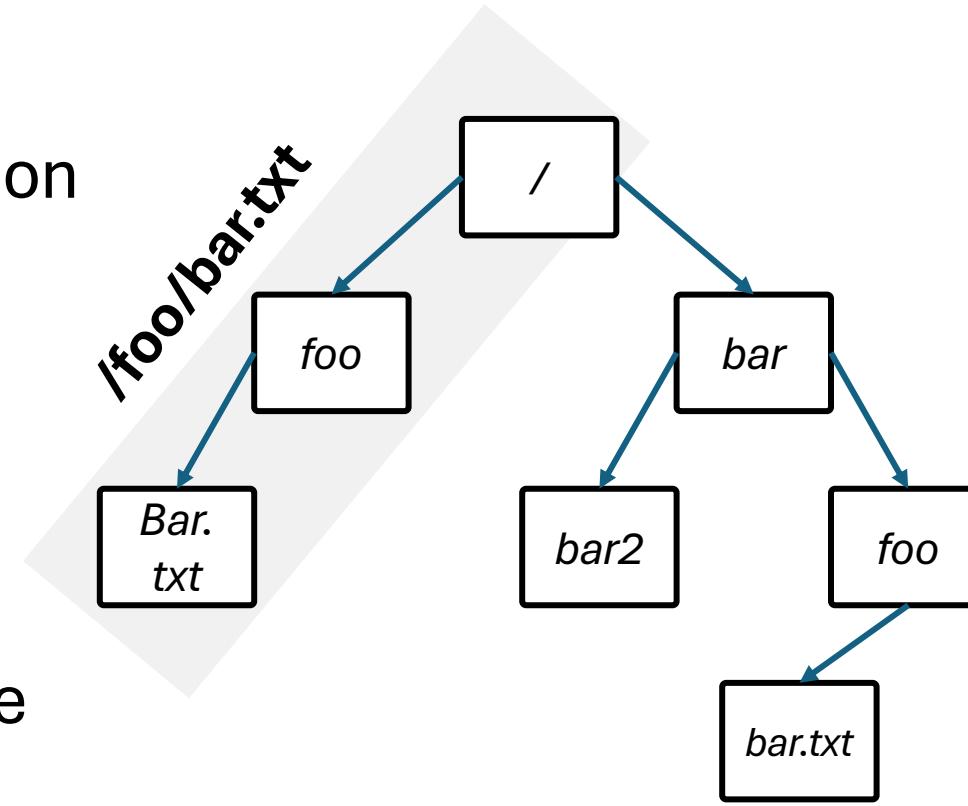
Modify: 2014-12-12 19:25:32.669625220 -0500

Change: 2014-12-12 19:25:32.669625220 -0500

Birth: -

The path (human-readable)

- Human-readable interpretation of every inode
- Typically, represented as –
<path-to-directory , filename>
- **Traversing** a tree – getting the final *inode* from a location
 - E.g., ls /foo/bar.txt
 - Gets the inode and prints details

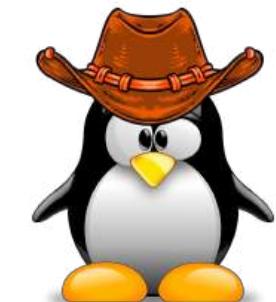


The file descriptor

- “Everything is a file”
 - File descriptor tracks what each ‘file’ really does in the system

```
struct file {  
    enum { FD_NONE, FD_PIPE, FD_INODE, FD_DEVICE } type;  
    int ref; // reference count  
    char readable;  
    char writable;  
    struct pipe *pipe; // FD_PIPE  
    struct inode *ip; // FD_INODE and FD_DEVICE  
    uint off; // FD_INODE  
    short major; // FD_DEVICE  
};
```

If file belongs to the disk, it has an inode



Commonly-used file system interfaces

Creating a new file or opening an existing file

- UNIX system call: open()

```
int fd = open(char *path, int flag, mode_t mode);
```

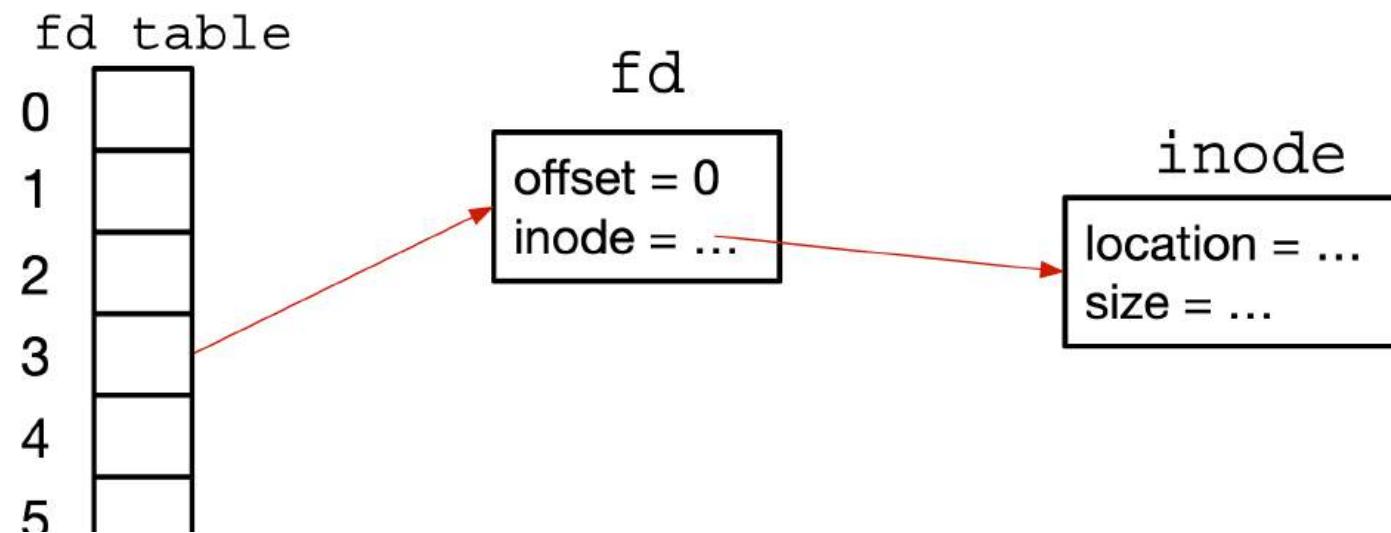
-OR-

```
int fd = open(char *path, int flag);
```

- open() returns a file descriptor (fd)
 - A fd is an integer
 - Private per process
- fd is a handle that gives caller the power to perform certain operations
 - You can think of a “fd” as a pointer to an object of the file

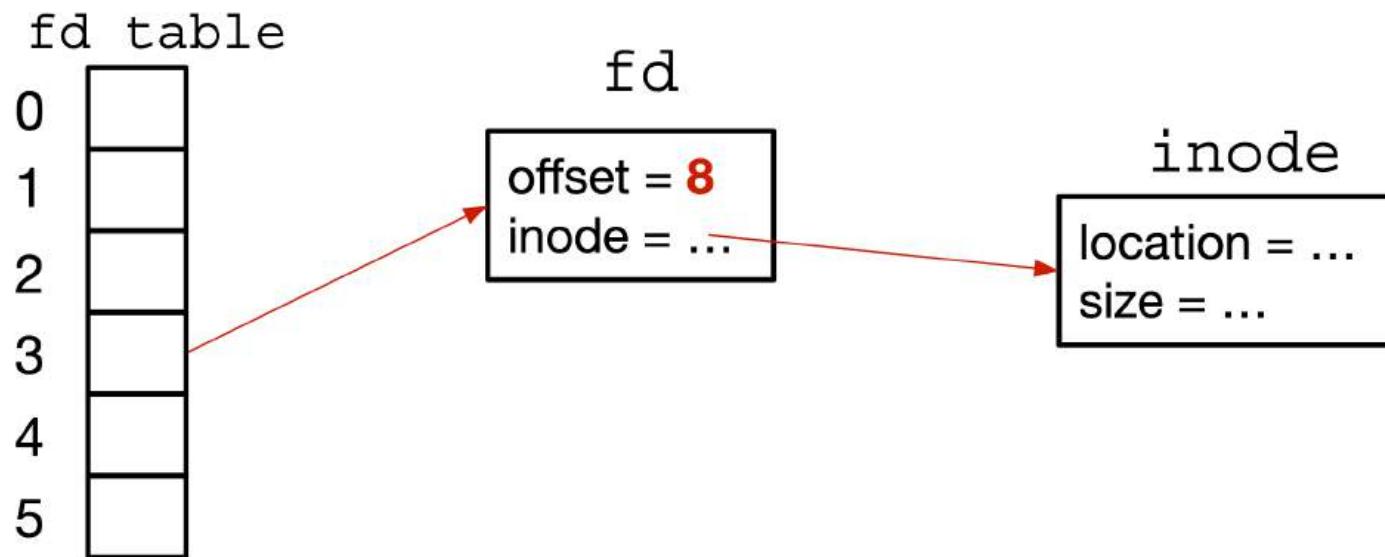
open () example

```
int fd1 = open("file.txt", O_CREAT); // return 3
```



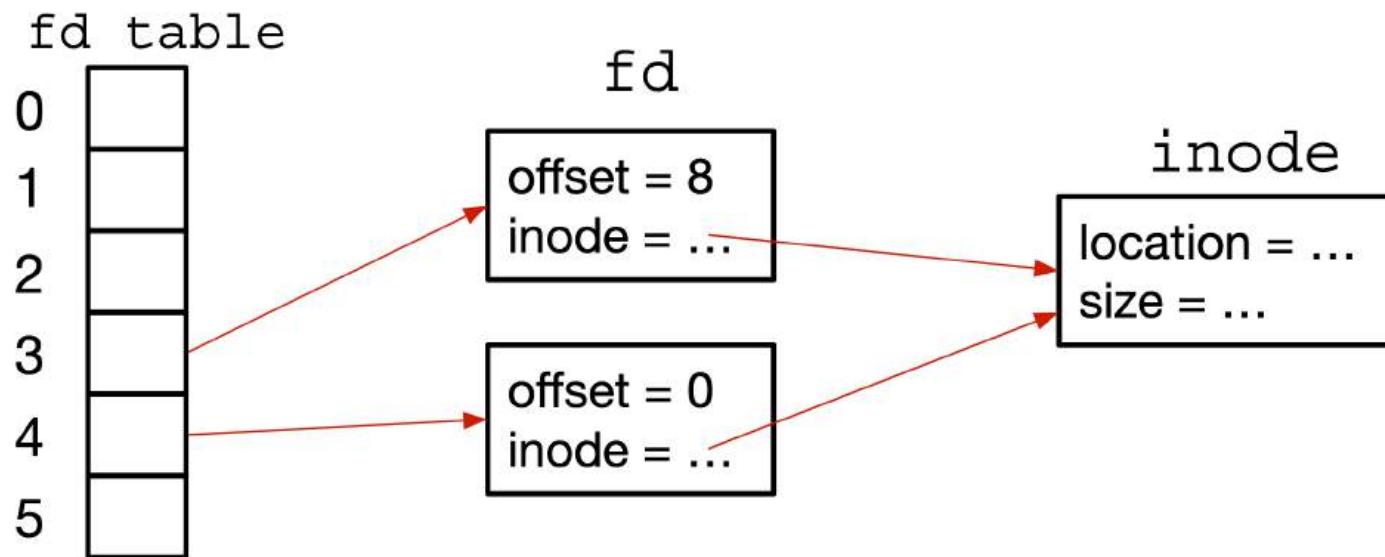
open () example

```
int fd1 = open("file.txt", O_CREAT); // return 3  
read(fd1, buf, 8);
```



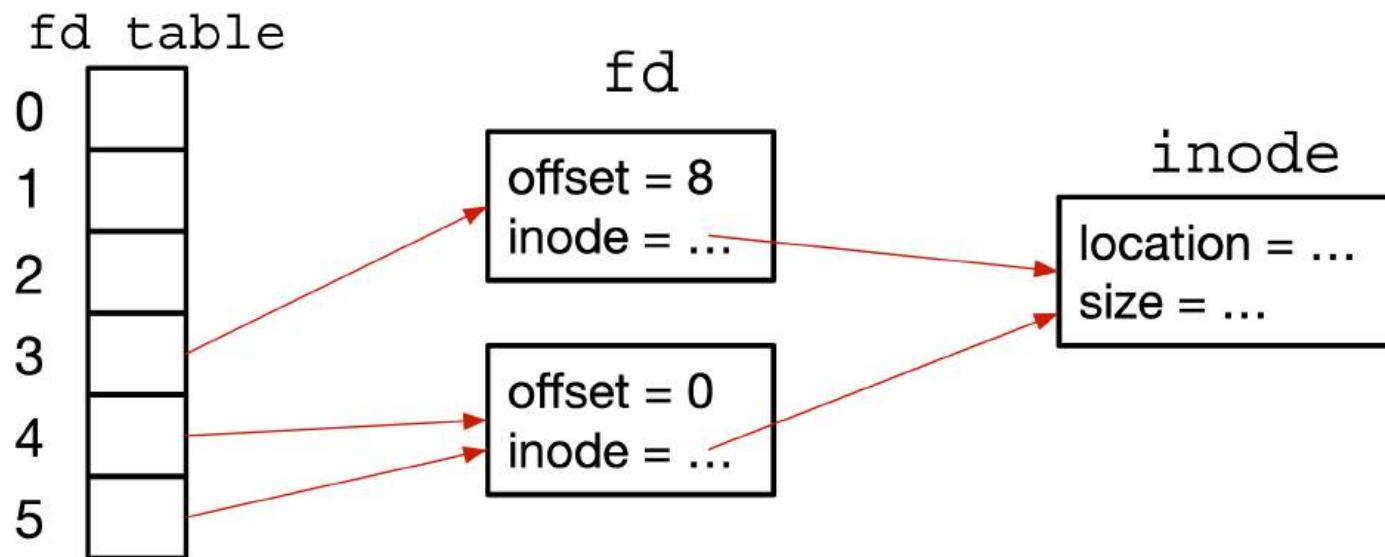
open () example

```
int fd1 = open("file.txt", O_CREAT); // return 3  
read(fd1, buf, 8);  
int fd2 = open("file.txt", O_WRONLY); // return 4
```



open () example

```
int fd1 = open("file.txt", O_CREAT); // return 3  
read(fd1, buf, 8);  
int fd2 = open("file.txt", O_WRONLY); // return 4  
int fd3 = dup(fd2); // return 5
```



Other common file system interfaces

- UNIX system call: read() and write()

```
int ret = read(int fd, char *buffer, size_t size);  
int ret = write(int fd, char *buffer, size_t size);
```

- **Can anyone tell me why they have the same arguments?**

Can anyone tell me what FS calls this command invokes?

- `cat test.txt`

Can anyone tell me what FS calls this command invokes?

- `cat test.txt`

```
prompt> strace cat file.txt
...
open("file.txt", O_RDONLY)      = 3
read(3, "hello\n", 65536)       = 6
write(1, "hello\n", 6)          = 6
read(3, "", 65536)             = 0
close(3)                        = 0
...
prompt>
```

Can anyone tell me what FS calls this command invokes?

Open the file with read
only mode

```
prompt> strace cat file.txt
```

```
...
open("file.txt", O_RDONLY)      = 3
read(3, "hello\n", 65536)       = 6
write(1, "hello\n", 6)          = 6
read(3, "", 65536)             = 0
close(3)                        = 0
```

```
...
```

```
prompt>
```

Can anyone tell me what FS calls this command invokes?

Open the file with read
only mode

Read content from file

```
prompt> strace cat file.txt
```

```
...
open("file.txt", O_RDONLY)      = 3
read(3, "hello\n", 65536)       = 6
write(1, "hello\n", 6)          = 6
read(3, "", 65536)             = 0
close(3)                        = 0
...
prompt>
```

Can anyone tell me what FS calls this command invokes?

Open the file with read
only mode

Read content from file

Write string to std
output fd 1

```
prompt> strace cat file.txt
```

...

```
open("file.txt", O_RDONLY) = 3
```

```
read(3, "hello\n", 65536) = 6
```

```
write(1, "hello\n", 6) = 6
```

```
read(3, "", 65536) = 0
```

```
close(3) = 0
```

...

```
prompt>
```

Can anyone tell me what FS calls this command invokes?

Open the file with read
only mode

Read content from file

Write string to std
output fd 1

cat tries to read more
but reaches EOF

prompt> strace cat file.txt

```
...
open("file.txt", O_RDONLY)      = 3
read(3, "hello\n", 65536)       = 6
write(1, "hello\n", 6)          = 6
read(3, "", 65536)             = 0
close(3)
...
prompt>
```

Can anyone tell me what FS calls this command invokes?

Open the file with read
only mode

Read content from file

Write string to std
output fd 1

cat tries to read more
but reaches EOF

cat done with file ops
and closes the file

prompt> strace cat file.txt

...	
open("file.txt", O_RDONLY)	= 3
read(3, "hello\n", 65536)	= 6
write(1, "hello\n", 6)	= 6
read(3, "", 65536)	= 0
close(3)	= 0
...	

prompt>

Are writes to disk performed exactly when you call write()?

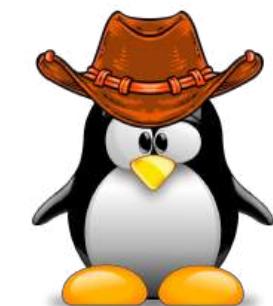
- No, that would be very **slow** since more writes could be needed later, and they can be **batched together**
 - Hence, your FS will buffer writes *in-memory*
- Sometimes though, you may want to force writes, e.g., if a later operation depends on disk write or to ensure persistence.
 - For such cases, FS provides the `fsync()` system call

Example of how “vim” uses the filesystem

```
prompt> vim file.txt
... vim editing session
...
prompt> .....▶ :W
q
```

The diagram illustrates the interaction between a terminal session and a corresponding C program. A terminal window at the top shows a user entering commands: 'vim file.txt', followed by an ellipsis indicating an editing session, another ellipsis, and then ':W' (to write changes) and 'q' (to quit). A downward arrow points from the terminal window to a block of C code at the bottom. The C code implements the logic for these commands. It starts by opening a temporary swap file ('.file.txt.swp') in write-only mode with creation and truncation flags, and permissions matching the original file. It then writes the new content to this file. After writing, it performs an `fsync(fd);` operation to ensure data durability. Finally, it closes the temporary file and renames it back to 'file.txt', effectively replacing the original file.

```
int fd = open(".file.txt.swp", O_WRONLY|O_CREAT|O_TRUNC, S_IRUSR|S_IWUSR);
write(fd, buffer, size); // write out new version of file
fsync(fd); // make data durable
close(fd); // close tmp file
rename(".file.txt.swp", "file.txt"); // change name and replacing old file
```



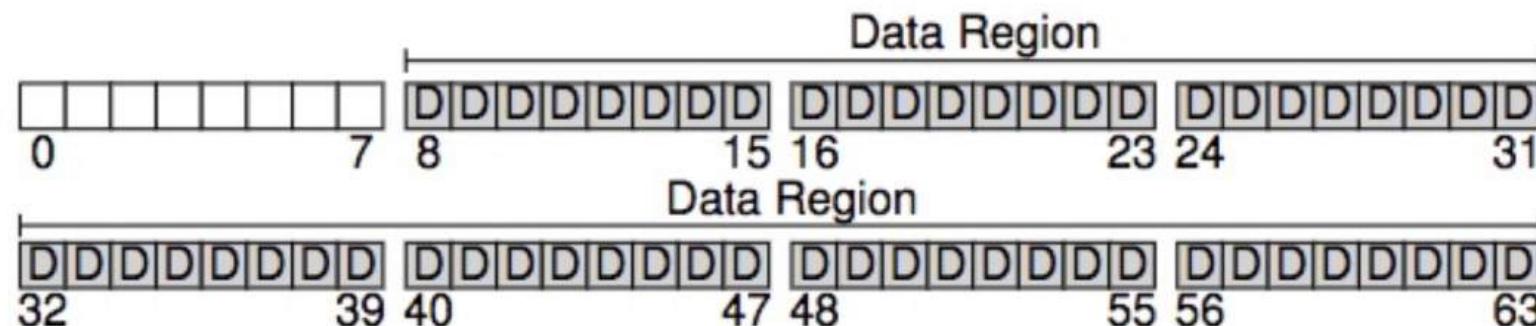
Implementation of a filesystem

What are the questions we may need to answer to build an FS?

- Where to store the files in the disk?
- Where/how to store the information about existing files (e.g., inodes, path, blocks, etc.) in the disk?
- How to quickly find free locations in the disk for new files?

Storing file “data blocks” within the disk

- Initially, the entire disk is empty
- We will save some space to store information
- Use the rest to store our “data blocks” (figure below)



Now, we must answer:

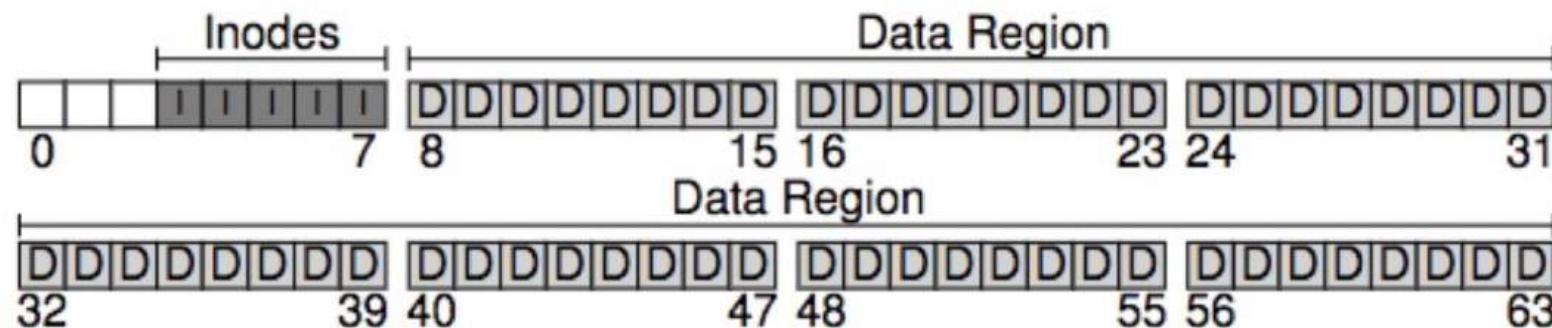
- Where to store the files in the disk?
- Where/how to store the information about existing files (e.g., inodes, path, blocks, etc.) in the disk?
- How to quickly find free locations in the disk for new files?

An **inode** data structure

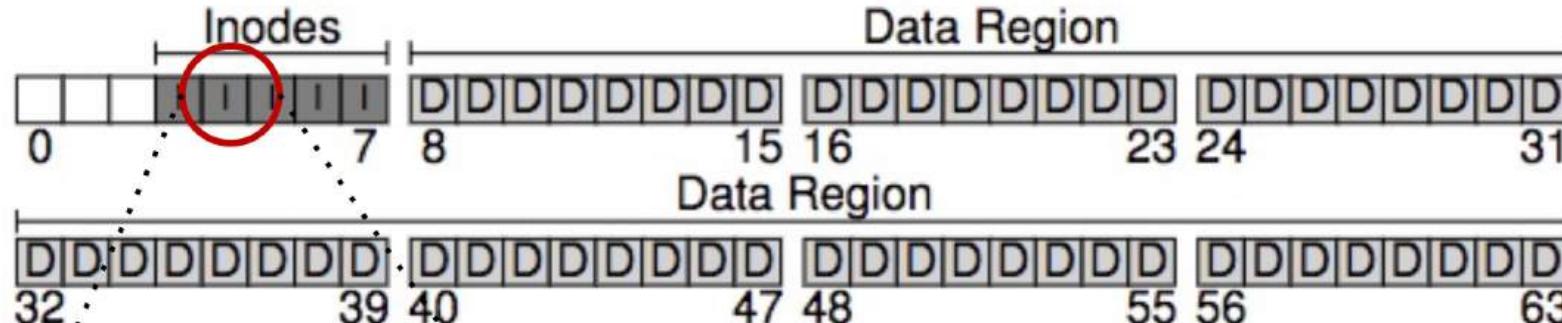
- Answers the “*how to store file information*” question
- Inode → The number corresponds to a data structure consisting of :
 - Type (e.g., file or directory)
 - Size
 - Number of data blocks
 - Address of data blocks
 - User permissions (e.g., root, etc.)
 - ...
- Inodes are 128/256 bytes, depending on the FS implementation

An inode table

- Answers the “*where to store file information*” question
- In a part of the reserved disk space, as a table/array format



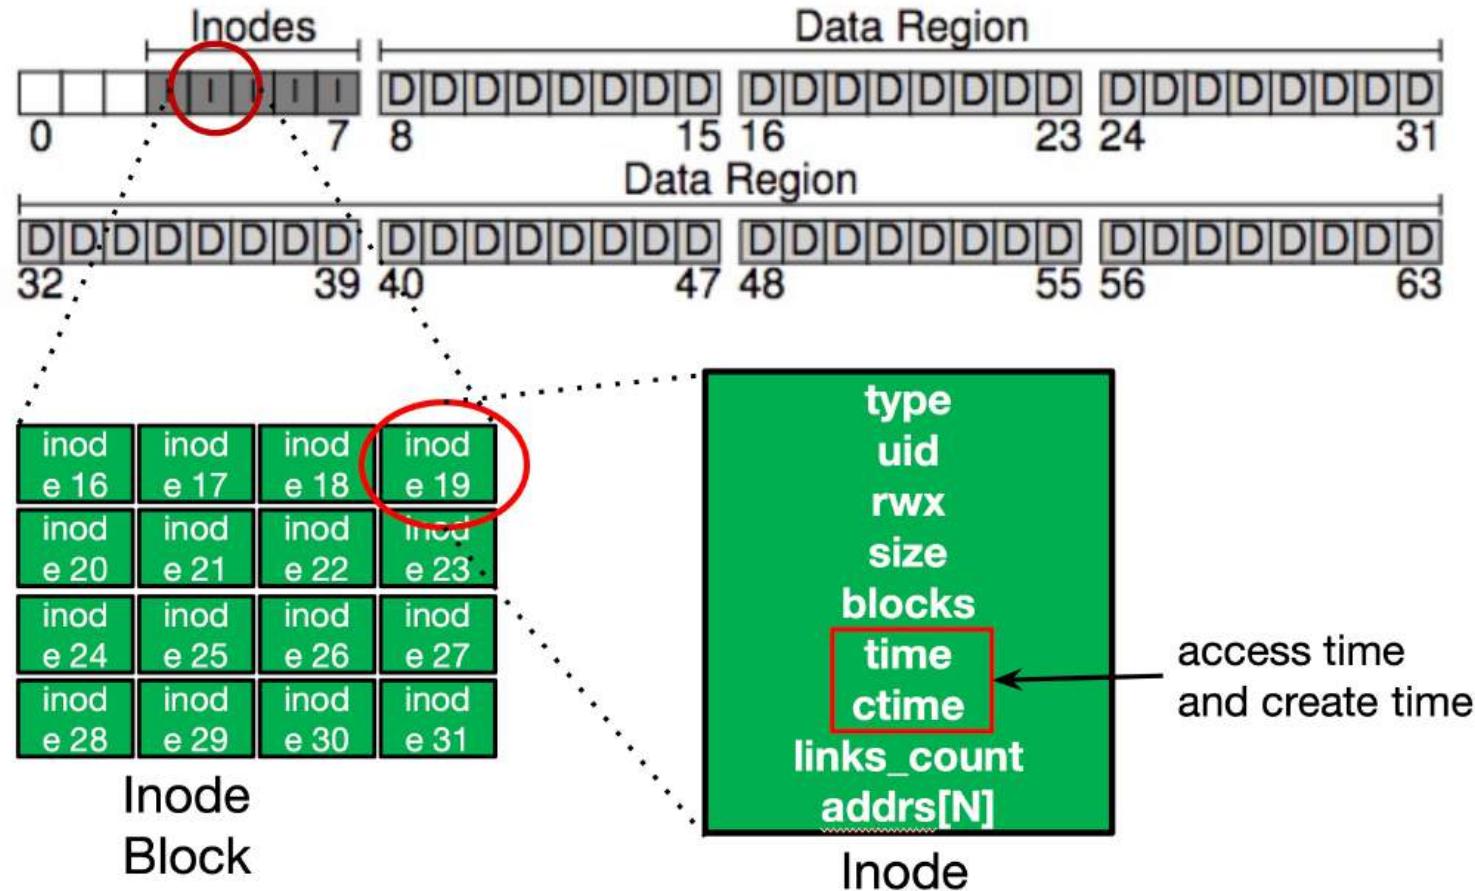
Visualizing the **inode** and its **table** in the disk



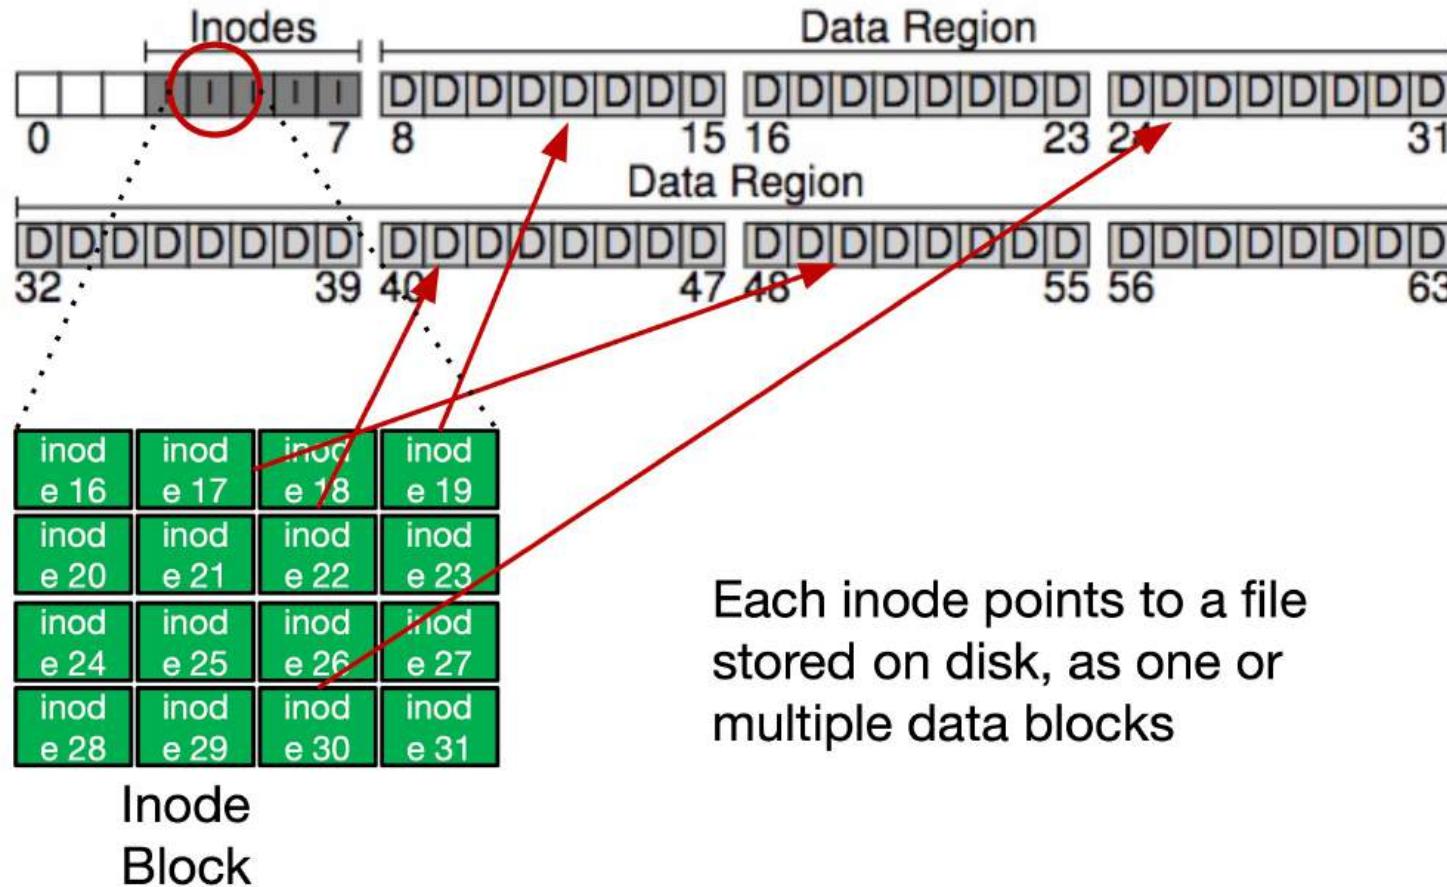
inod e 16	inod e 17	inod e 18	inod e 19
inod e 20	inod e 21	inod e 22	inod e 23
inod e 24	inod e 25	inod e 26	inod e 27
inod e 28	inod e 29	inod e 30	inod e 31

- Inodes are typically 128 or 256 bytes (depends on the file system)
 - 16—32 inodes per inode block

Visualizing the **inode** and its **table** in the disk



Visualizing the **inode** and its **table** in the disk



Now, we must answer:

- Where to store the files in the disk?
- Where/how to store the information about existing files (e.g., inodes, path, blocks, etc.) in the disk?
- How to quickly find free locations in the disk for new files?

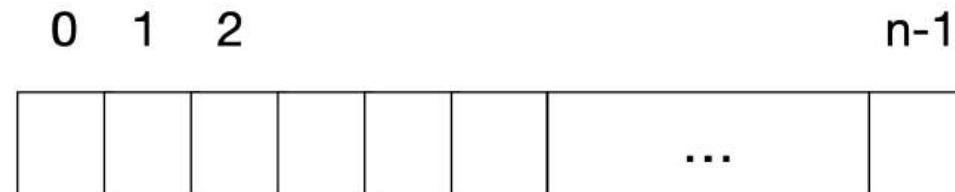
This question has two sub-questions. What are they?

- How to quickly find free locations in the disk for new files?
 - 1) How to quickly find free “data blocks”?
 - 2) How to quickly find free “inodes”?

Using **bitmaps** to store free data and inode blocks

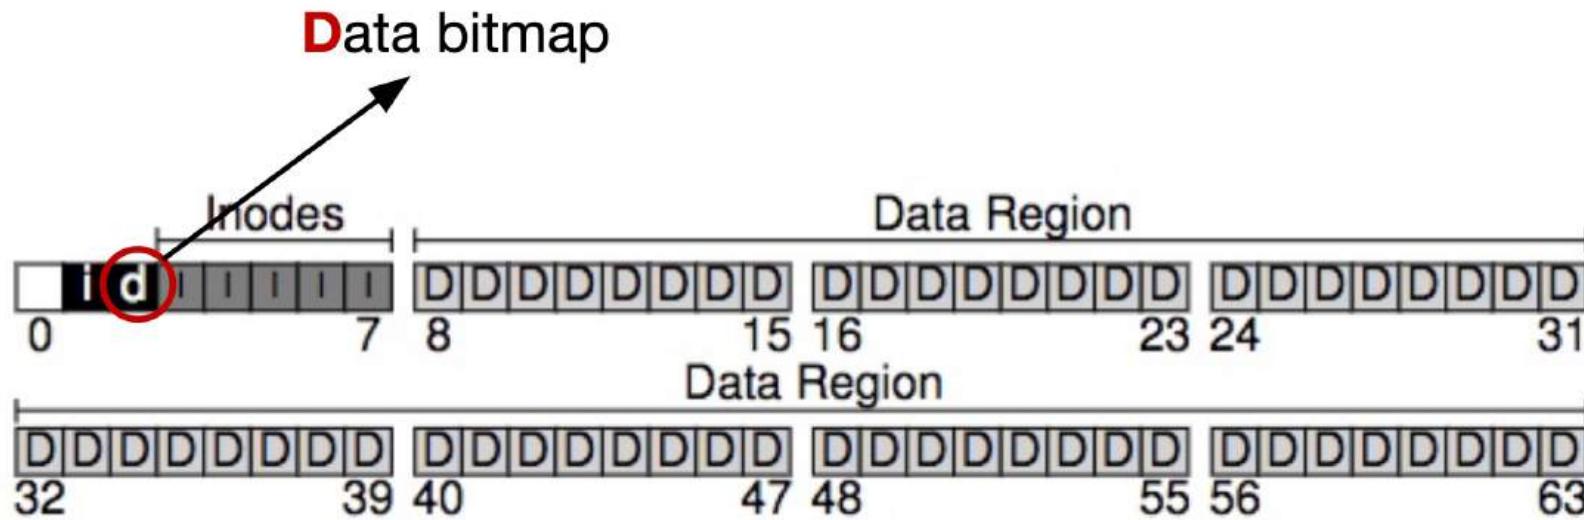
- Can anyone tell me what's a bitmap?

Each bit of the bitmap is used to indicate whether the corresponding object/block is **free** (0) or **in-use** (1)

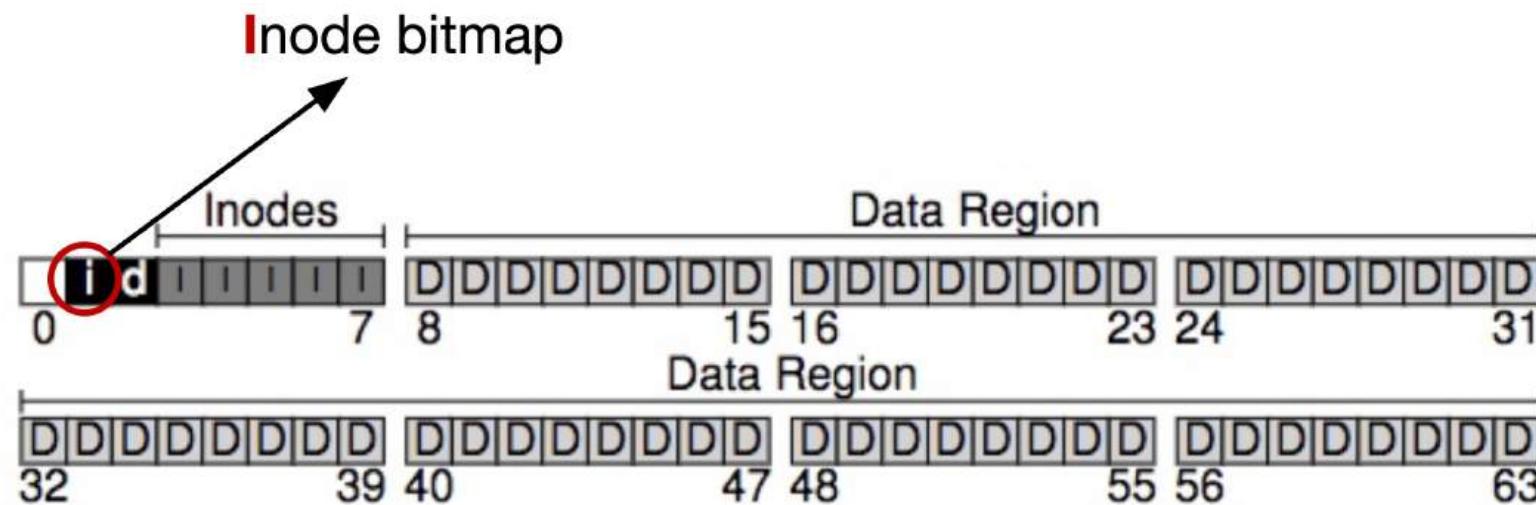


bit[i] 1 \Rightarrow object[i] in use
= 0 \Rightarrow object[i] free

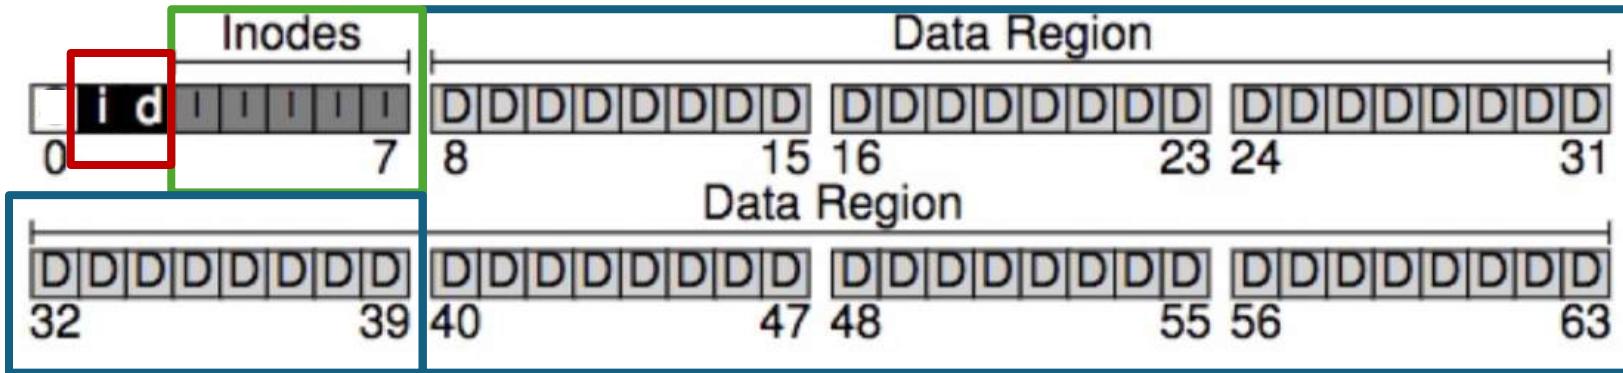
Bitmap to find free data blocks: data bitmap



Bitmap to find free data blocks: inode bitmap



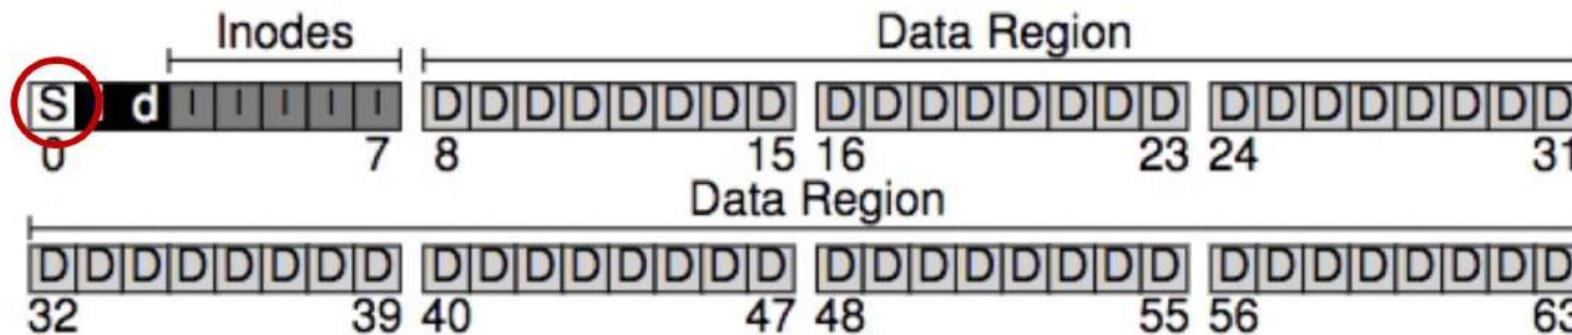
Is there another question that we may have missed?



- How to persistently keep track of which regions contain all these disk structures (e.g., data blocks, inode table, bitmaps, etc.)?
- Their size may be variable (e.g., depending on disk)

The “superblock”

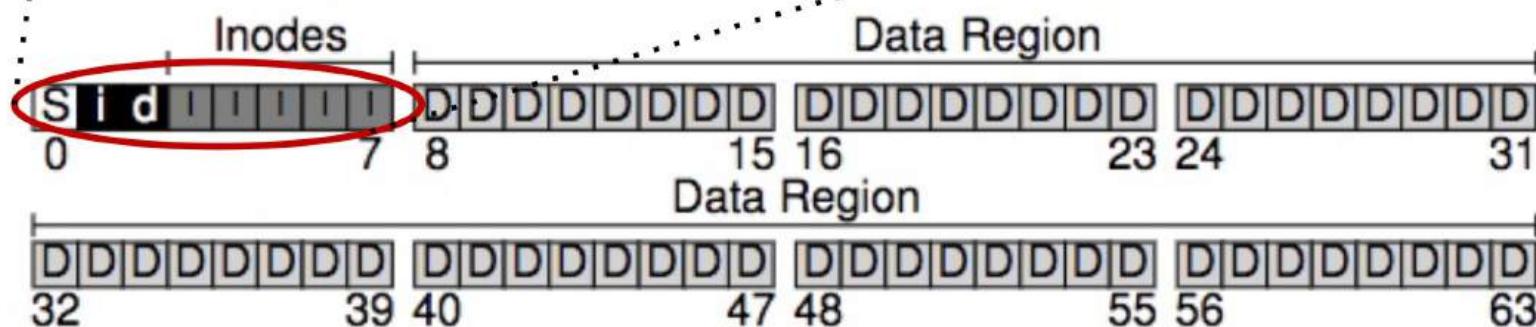
- Keeps track of basic filesystem information, like
 - Block size
 - How many inodes are there?
 - How much space is free?

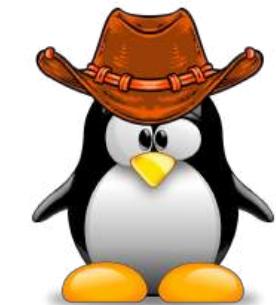


Putting it all together

The Inode Table (Closeup)

					iblock 0	iblock 1	iblock 2	iblock 3	iblock 4													
Super	i-bmap	d-bmap	0KB	4KB	8KB	12KB	16KB	20KB	24KB	28KB	32KB											
			0	1	2	3	16	17	18	19	32	33	34	35	48	49	50	51	64	65	66	67
			4	5	6	7	20	21	22	23	36	37	38	39	52	53	54	55	68	69	70	71
			8	9	10	11	24	25	26	27	40	41	42	43	56	57	58	59	72	73	74	75
			12	13	14	15	28	29	30	31	44	45	46	47	60	61	62	63	76	77	78	79





Questions? Otherwise, see you next class!