




# CSE 330: Operating Systems


**Adil Ahmad**

**Lecture #20:** Filesystem logging (crash recovery and consistency)




# FS optimization(s) recap

# What is the difference between the two workloads?





- **Spatial** locality: addresses are close to each other in “space”




- **Temporal** locality: addresses are close to each other in “time”

# Which of the following is actually better for disks?




- **Spatial** locality: addresses are close to each other in “space”
- **Temporal** locality: addresses are close to each other in “time”

# Imagine that we must choose inode and data blocks




- **Is this a good or a bad policy to choose blocks?**
- Bad policy; will require slow “random” access


# A better policy would be as follows



- **Is there something else that we can optimize?**
- Location of the inode itself, in this particular case. ☺

# An even better policy







# FS optimization: buffering in memory

# Avoiding excessive I/O operations is important

- Even with all possible disk optimizations, I/O remains slow and unpredictable depending on workloads
- File system will cache “files” in memory (recall how the xv6 block layer caches “blocks”) to improve performance
- To do so, the file system will:
  - Reserve a region for the “file cache”
  - Find out which files are frequently read/written
  - Keep those files within the cache

# Modern OSs take a “unified” cache approach

- Keep a single cache for (a) file pages and (b) program memory pages
- Avoids having multiple caches, and separately dealing with them



# Buffering creates a tiny headache for the file system..

- Imagine the following set of events from an FS:
  1. Program writes to a new 1G file (e.g., bar)
  2. FS caches all writes to the file in-memory
  3. FS updates the inodes within the disk
  4. FS writes all changes to data blocks within disk (e.g., block-by-block using Linux's block I/O)
- **Can you spot potential problems?**
  - If power goes off after 3, your system has already reserved the data block → FS lost a few free data blocks

# Worse problems could also occur on crashes

- Imagine the following set of events from an FS:
  1. Imagine that the FS is deleting a file X:
  2. Frees block for X and changes its data bitmap
  3. Before FS can delete the *inode* for the block, power goes off.
  4. OS restarts; thinks block is free and allocates it to a new file Y.
- **Can you spot potential problems?**
  - Two inodes are pointing to the same data blocks → write to two different files will overwrite each other
- **Can this be a security issue?**
  - Yes, imagine two different users get allocated the blocks

# Crash recovery is a significant problem for FSs


- Many filesystem operations require **multiple writes** to disk
  - `creat(file)` → 5 disk blocks to be written
    - inode, directory, and actual file data
- A crash (e.g., power failure) after a subset of these writes might put the filesystem in an **inconsistent state**
- **Note this is only a problem for write operations. Why is that?**
  - Data/Inode blocks only get changed on writes, not reads.

# How would you solve this problem of FS corruption on crashes?

- Imagine the following set of events from an FS:
  1. Program writes to a new 1G file (e.g., bar)
  2. FS caches all writes to the file in-memory
  3. FS updates the inodes within the disk
  4. FS writes all changes to data blocks within disk (e.g., block-by-block using Linux's block I/O)
- **Can you spot potential problems?**
  - If power goes off after 3, your system has already reserved the data block → FS lost a few free data blocks

# How would you solve this problem of FS corruption on crashes?

- Imagine the following set of events from an FS:
  1. Imagine that the FS is deleting a file X:
  2. Frees block for X and changes its data bitmap
  3. Before FS can delete the *inode* for the block, power goes off.
  4. OS restarts; thinks block is free and allocates it to a new file Y.
- **Can you spot potential problems?**
  - Two inodes are pointing to the same data blocks → write to two different files will overwrite each other



Ensuring correctness after system crash

# Filesystem **logging** addresses crash problems

*Let's look at the high-level idea of logging:*

- Never write directly to the on-disk data structures
- Place all writes to a “log” on the disk
- Once all log writes in one operation are complete, issue a “commit”
- FS writes the data to actual disk locations using “log”
- FS removes “log” entry


# Does logging prevent corruption on crashes?

*Let's look at the high-level idea of logging:*

- Never write directly to the on-disk data structures
  - Place all writes to a “log” on the disk
  - Once all log writes in one operation are complete, issue a “commit”
    - FS writes the data to actual disk locations using “log”
    - FS removes “log” entry
- 
- The diagram illustrates the high-level idea of logging through a list of steps. Step 3, 'Once all log writes in one operation are complete, issue a “commit”', is highlighted with a red border. Three callout boxes provide context for this step:
  - A box to the right of step 3 contains the text: 'Discard all writes on reboot after crash here'. An arrow points from this box to the word 'commit'.
  - A box below step 3 contains the text: 'All needed data is in disk to complete write'. An arrow points from this box to the word 'commit'.
  - A box to the right of step 4 contains the text: 'Can recover the data using log on crash here'. An arrow points from this box to the word 'commit'.

Ensures atomicity: all writes in a transaction are logged or none

# Where is the log in this figure BTW?



- We must reserve it as well in some fixed part of the disk, and note its information (e.g., starting block, size) in the superblock

# *Code review: xv6's crash recovery through logging*

*Step #1: Signaling start of log at the write system call*

*Step #2: Writing data to in-memory buffers*

*Step #3: Tracking the buffers that must be flushed to disk*

*Step #4: Commit cache to log disk sectors*

*Step #5: Moving data from log to final disk locations*

# Step #1: Signaling the start of a log at write syscall

write(..) system call

```
int  
filewrite(struct file *f, uint64 addr, int n)  
{  
...  
begin_op();  
ilock(f->ip);  
if ((r = writei(f->ip, 1, addr + i, f->off, n1)) > 0)  
    f->off += r;  
iunlock(f->ip);  
end_op();  
}
```

Signalling starts here

Perform the writes now

Called at the end

```
// called at the start of each FS system call.  
void  
begin_op(void)  
{  
    acquire(&log.lock);  
    while(1){  
        if(log.committing){  
            sleep(&log, &log.lock);  
        } else if(log.lh.n + (log.outstanding+1)*MAXOPBLOCKS > LOGSIZE){  
            // this op might exhaust log space; wait for commit.  
            sleep(&log, &log.lock);  
        } else {  
            log.outstanding += 1;  
            release(&log.lock);  
            break;  
        }  
    }  
}
```

Is log being committed?

Is there enough space?

# Step #2: Writing data to in-memory buffers

```
int  
writei(struct inode *ip, int user_src, uint64 src, uint off, uint n)
```

Called from filewrite()

...

```
for(tot=0; tot<n; tot+=m, off+=m, src+=m){  
    uint addr = bmap(ip, off/BSIZE);  
    if(addr == 0)  
        break;  
    bp = bread(ip->dev, addr);  
    m = min(n - tot, BSIZE - off%BSIZE);  
    if(either_copyin(bp->data + (off % BSIZE), user_src, src, m) == -1) {  
        brelse(bp);  
        break;  
    }  
    log_write(bp);  
    brelse(bp);  
}
```

Try to find a cache to keep  
block in-memory

Read block from disk

Copy data from userspace

Similar to the `copy_from_user(..)`  
you have been using for kernel modules

# Step #3: Tracking that a buffer must be written later

```
int
writei(struct inode *ip, int user_src, uint64 src, uint off, uint n)

for(tot=0; tot<n; tot+=m, off+=m, src+=m){
    uint addr = bmap(ip, off/BSIZE);
    if(addr == 0)
        break;
    bp = bread(ip->dev, addr);
    m = min(n - tot, BSIZE - off%BSIZE);
    if(either_copyin(bp->data + (off % BSIZE), user_src, src, m) == -1) {
        brelse(bp);
        break;
    }
    log_write(bp);
    brelse(bp);
}
```

Must keep metadata for log

Log metadata structures

```
struct log {
    struct logheader;
    ...
}
```

```
struct logheader {
    int n;
    int block[LOGSIZE];
}
```

log\_write(..

Reuse if already allocated

```
) for (i = 0; i < log.lh.n; i++) {
    if (log.lh.block[i] == b->blockno) // log absorption
        break;
}
log.lh.block[i] = b->blockno;
```

Place each block into an empty log block

# Step #4: Commit cache to log sectors

*writei(..) → filewrite(..) → end\_op()*

```
static void  
commit()  
{  
    if (log.lh.n > 0) {  
        write_log(); // Write modified blocks from cache to log  
        write_head(); // Write header to disk -- the real commit  
        install_trans(0); // Now install writes to home locations  
        log.lh.n = 0;  
        write_head(); // Erase the transaction from the log  
    }  
}
```

Write from cache to log

Not actually ‘committed’  
until we save log header

```
static void  
write_head(void)  
{  
    struct buf *buf = bread(log.dev, log.start);  
    struct logheader *hb = (struct logheader *) (buf->data);  
    int i;  
    hb->n = log.lh.n;  
    for (i = 0; i < log.lh.n; i++) {  
        hb->block[i] = log.lh.block[i];  
    }  
    bwrite(buf);  
    brelse(buf);  
}
```

Read log header from disk

Is the log actually  
‘committed’ at this point?

Change, then write-back

# Step #5: Write cache to disk and delete log entry

*writei(..) → fwrite(..) → end\_op()*

```
static void  
commit()  
{  
    if (log.lh.n > 0) {  
        write_log(); // Write modified blocks from cache to log  
        write_head(); // Write header to disk -- the real commit  
        install_trans(0); // Now install writes to home locations  
        log.lh.n = 0;  
        write_head(); // Erase the transaction from the log  
    }  
}
```

Remove the log to avoid  
re-copy on crash

```
// Copy committed blocks from log to their home location  
static void  
install_trans(int recovering)  
{  
    int tail;  
  
    for (tail = 0; tail < log.lh.n; tail++) {  
        struct buf *lbuf = bread(log.dev, log.start+tail+1); // read log block  
        struct buf *dbuf = bread(log.dev, log.lh.block[tail]); // read dst  
        memmove(dbuf->data, lbuf->data, BSIZE); // copy block to dst  
        bwrite(dbuf); // write dst to disk  
        if(recovering == 0)  
            bunpin(dbuf);  
        brelse(lbuf);  
        brelse(dbuf);  
    }  
}
```

Read log src and  
disk dst blocks

Write-back to dst

Is the store to disk “actually  
committed” at this point?

# Aspects of FS logging can become challenging

- All data at system call is first written to in-memory block cache
- Then transferred to the log disk space inside your disk
- Also, FS system calls can happen concurrently

Can you spot the major challenges to ensure smoothness?

# Challenges of logging?

- **System call data must fit inside the log disk space**
  - Recall that each write(..) puts all blocks into log regions
  - xv6's solutions:
    1. Set an upper bound of all log entries
      - Set log size  $\geq$  upper bound
    2. Breakdown large writes into smaller transactions
      - Problem: large writes are thus not atomic
      - However, overall system remains in a consistent state

# Challenges of logging?

- **Concurrent system call handling**
  - 1. Must allow different FS system calls from different proc at the same time
  - 2. Block may be written multiple times while still in-memory
    - Creating a new log entry for each time is wasteful
- xv6's solutions:
  - 1. Check if a system call's data can fit in log before starting it
    - Else, sleep and wait for log to free up
  - 2. If the block is already assigned a log entry, return without creating new entry
    - Also called “write absorption” → absorb multiple writes into a single

# Pros of xv6's logging?

- FS internal invariants are maintained
  - Metadata remains consistent (e.g., no two inodes point to same blocks)
- Except for the last few operations, data is preserved on the disk
  - Every system call is written at the end before return
  - No surprises of losing data written *way back*
- Write order is preserved
  - \$ echo A > X ; echo B > Y
  - X will be written before Y

# Cons of xv6's crash recovery method?

- Several inefficiencies

*Design choice:* ensure either old or new copy is correct. No corruption


- Every block is written twice (once for log + once for final)
- Eagerly writes to disk after every write system call
- Writes each block one-by-one to disk

**Implementation choice:** How would you overcome?

Keep data in-memory for longer periods and batch disk writes

# Ordered mode in ext3/4 FS avoids twice data writes

- Log the FS metadata but write data directly to the final disk location
  - Crash **consistency**, but **not recovery**
- High-level overview of ordered mode:
  1. Metadata is committed to log
  2. Data is written to final disk location
  3. Metadata is moved from log to final disk location
- Limitation?
- Suggested reading(s):  
Journaling the Linux ext2fs: <https://pdos.csail.mit.edu/6.1810/2022/readings/journal-design.pdf>  
Analysis and Evolution of Journaling File Systems (Prabhakaran et. al)



Questions? Otherwise, see you next class!