



# CSE 330: Operating Systems

**Adil Ahmad**

**Lecture #21:** System virtualization overview

# Recap of FS logging

- All data at system call is first written to in-memory block cache
- Then transferred to the log disk space inside your disk
- Also, FS system calls can happen concurrently

Can you spot the major challenges to ensure smoothness?

# Challenges of logging?

- **System call data must fit inside the log disk space**
  - Recall that each write(..) puts all blocks into log regions
  - xv6's solutions:
    1. Set an upper bound of all log entries
      - Set log size  $\geq$  upper bound
    2. Breakdown large writes into smaller transactions
      - Problem: large writes are thus not atomic
      - However, overall system remains in a consistent state

# Challenges of logging?

- **Concurrent system call handling**
  - 1. Must allow different FS system calls from different proc at the same time
  - 2. Block may be written multiple times while still in-memory
    - Creating a new log entry for each time is wasteful
- xv6's solutions:
  - 1. Check if a system call's data can fit in log before starting it
    - Else, sleep and wait for log to free up
  - 2. If the block is already assigned a log entry, return without creating new entry
    - Also called “write absorption” → absorb multiple writes into a single

# Pros of xv6's logging?

- FS internal invariants are maintained
  - Metadata remains consistent (e.g., no two inodes point to same blocks)
- Except for the last few operations, data is preserved on the disk
  - Every system call is written at the end before return
  - No surprises of losing data written *way back*
- Write order is preserved
  - \$ echo A > X ; echo B > Y
  - X will be written before Y

# Cons of xv6's crash recovery method?


- Several inefficiencies

*Design choice:* ensure either old or new copy is correct. No corruption

- Every block is written twice (once for log + once for final)
- Eagerly writes to disk after every write system call
- Writes each block one-by-one to disk

**Implementation choice:** How would you overcome?

Keep data in-memory for longer periods and batch disk writes



Moving on to system virtualization

# What is the meaning of the term ‘virtualization’?


Provide an abstracted or ‘fake’ view of computer resources

## **How have we used the concept of ‘virtualization’ so far?**

- Memory virtualization for processes
  - Set-up page tables to prevent the process from accessing other regions
- Disk virtualization
  - We leverage the filesystem and block layers to implement access control, etc.

# What's a virtual machine (or a VM)?

A *simulated* instance of a computer inside a computer



## Desired VM properties:

1. Accurate simulation
2. Isolation from ‘host’
3. Fast!

# Why do we use VMs in today's world?

## A. Cloud computing:

- Each user has their own isolated ‘machine’
- Can run different OSs and remain isolated from other users
- Better resource management for cloud providers

## B. System development:

- Testing OSs without breaking your own system
- Testing programs built for certain OS/packages

## C. System security:


- Many important OS features can be secured by implementing in a higher privileged layer (remember nested kernel?)

# Different ways to design a VM today

1. Full software-based emulation
2. Trap-and-emulate
  - Also includes para-virtualized (PV) interfaces
3. Hardware-assisted virtualization

# 1. Full software emulation-based virtualization

Simulate the entire machine's execution in software-mode



Other state that must be simulated: privileged registers, UART, PMP, etc.

# An example of a RISC-V instruction interpretation

- Guest executes address 0xa: “csrr a1,mhartid”
- VMM reads 0xa and finds the binary code: 0xf14025f3
- VMM knows RISC-V instruction convention
  - Instruction is 4-bytes (or 32-bits)
  - Bits 0-6 is instruction opcode (0x73 → CSR instructions)
  - Bits 7-11 is destination register (0xb → a1 register)
  - Bits 15-19 is src1 register (0x0 → zero register/unused)
  - Bits 20-31 is src2 register (0xf14 → mhartid register)
- Based on convention, VMM can keep update simulated registers

# What are the advantages of full software emulation?

## A. Architecture and OS-agnostic

- Can run any architecture (e.g., RISC-V on x86 using QEMU)
- Do not actually even need support from the Operating System

## B. Great for prototyping new devices/hardware

- Testing new CPU features without having real hardware
- Some very famous examples:
  - Intel SGX emulation using QEMU
  - Arm CCA emulation

# What are the disadvantages of full software emulation?

## A. Pretty slow!

- Must translate *every* instruction as well as maintain internal register/system state (e.g., all registers)

## B. Complex VMM design

- Your VMM must have lots of emulation code (e.g., check QEMU's software emulation stack)
- Easy to make mistakes – emulated HW does not mimic actual machine (very common!!)

# How can we speed up virtual machine execution?

# Allowing VMs to execute instructions on CPU is faster!

- Does not incur the overhead of instruction interpreter  
*In fact, we allow a process to execute instructions directly!*
- **Why not allow the VM to directly execute its instructions?**  
OS inside the VM might execute a privileged instruction. E.g., load a new page table and defeat memory isolation
- **Is there a solution to this problem?**  
*Execute the VM in user-mode (U-mode)?!*

## 2. Trap-and-emulate (T&E) virtualization

- Run the VM as a user-mode process
- User instructions are directly executed by the VM
  - `ld a0, s1` → no trap
- Supervisor instructions by the VM OS are trapped to the VMM
  - VMM emulates these just like in full software emulation
- **Advantages of this approach?**
  - Helps avoid the cost of translating *all* instructions
  - Still can build this design in software entirely

# What system state must we trap and emulate?

*Short answer: any change to privileged CPU state must be T&E-ed*

- Page tables (registers)
- Privileged registers (e.g., model-specific registers)
- Interrupt controls (e.g., PLIC/CLINT)
- ..

# Understanding the trap semantics in different arch.

## How exactly does the VMM trap supervisor instructions?


- RISC-V → easy
  - Automatically handled by the architecture
  - Execute ‘csrr’ in u-mode – direct trap to the S-mode SW
- x86 (32-bit) → hard
  - Executing a privileged instruction → skip the instruction
  - *Solution: para-virtualization.*
    - Make the OS aware of virtualization
    - Change the OS’ privileged operations to directly call the VMM using hyper-calls (e.g., VMCALL instruction in x86)

# Let's look at a few important T&E scenarios


1. Executing a system call from U-mode inside the VM
2. Reading a privileged register by the guest OS
3. Making changes to page tables

# *Scenario #1: Executing a system call from the guest*

Executing a system call in  
a non-virtualized system



Executing a system call in  
trap+emulate system



1. syscall always traps
2. update emulated guest state

4. read virtual registers  
(next slide)

## *Scenario #2: Reading from a privileged register from guest*

OS tries to read CR3 register. Following happens:

- Trap into the VMM since CR3 is a privileged instruction
- VMM checks what the guest wants by examining guest instruction
- Performs the operation in SW; moves *instruction pointer* to a1
- Moves the instruction;  $\text{sepc} += 4$  bytes (4B is instruction size, e.g.,)
- Returns to the guest OS

## *Scenario #3: Making changes to page tables*

Guest OS creates page table and tries to write to CR3 register.

- VMM inspects page table and creates a “shadow”
  - Guest’s page table: guest VA → guest PA (host VA)
  - Current page table: guest PA (host VA) → host PA
  - Shadow page table: guest VA → host PA
- VMM installs shadow to real CR3 after making sure there are no malicious entries (e.g., VM trying to read/write outside)

# Disadvantages of trap-and-emulate?

## A. Still slow (lots of supervisor-level traps)

- Kernel's performance affects every software's performance
- Can be improved using several techniques:
  - Especially for an 'enlightened' guest; one that knows it's a VM and actively tries to reduce traps

## B. Complexity remains

- Privileged operations are the most important operations in a system and their emulation is no trivial matter

# How do we speed-up virtualization beyond T&E?

## A. Still slow (lots of supervisor-level traps)

- Kernel's performance affects every software's performance
- Can be improved using several techniques:
  - Especially for an 'enlightened' guest; one that knows it's a VM and actively tries to reduce traps

## B. Complexity remains

- Privileged operations are the most important operations in a system and their emulation is no trivial matter

# We should allow the VM to execute *all most* instructions!

- Does not incur the overhead of *any* traps  
**Even privileged instructions can be directly executed on the CPU**
- **How is this even possible, let alone secure?**
  - Think about a process – save kernel context in TRAPFRAME and move to process execution at ‘sret’ instruction
  - What if we can save the host context and jump to the guest VM’s execution and resume later?

# Today's standard: Hardware-assisted virtualization


- New mode of execution: root mode and non-root mode.
  - Root mode is for host/VMM, non-root mode for guest/VM
- CPUs provide support for isolating the execution of host and VM
  - VMENTER/VMRESUME – instruction to enter/resume a VM from host
  - VMEXIT – exit the VM back to the host
- On **VMENTER**, host should save all context that it needs for later
- On **VMEXIT**, CPU saves guest context in a “TRAPFRAME” called the “Virtual Machine Control Structure” or VMCS (x86)

# The context to be saved on exit/entry to VMs

**All** register contexts including the privileged/unprivileged registers

- Control registers (CRs)
- Segment registers
- Floating point registers, etc.

# x86 HW-assisted virtualization illustration



# There are three main **security problems**

We need to systematically address the following .

A. Guest could read/write outside its own memory regions

- E.g., using modified page tables

B. Guest could talk to devices that the host wants to isolate.

- Could also refuse to give up control to host at interrupts

C. Guest could adversely impact important system functionality

- Change voltage on a certain CPU core

# How would you address these problems? (one-by-one)

We need to systematically address the following .

A. Guest could read/write outside its own memory regions


- E.g., using modified page tables

B. Guest could talk to devices that the host wants to isolate.

- Could also refuse to give up control to host at interrupts

C. Guest could adversely impact important system functionality

- Change voltage on a certain CPU core



Questions? Otherwise, see you next class!